Introduction to deep learning 1/?

Esten Høyland Leonardsen

02.11.22

UiO:Life Science, University of Oslo

Introduction

1. Building an artificial neural network (ANN)
2. Training the ANN
3. Transformation to a Convolutional Neural Network (CNN)

Building a neural network: Linear regression

$$y = ax + b$$

Building a neural network: Linear regression

$$y = ax + b$$

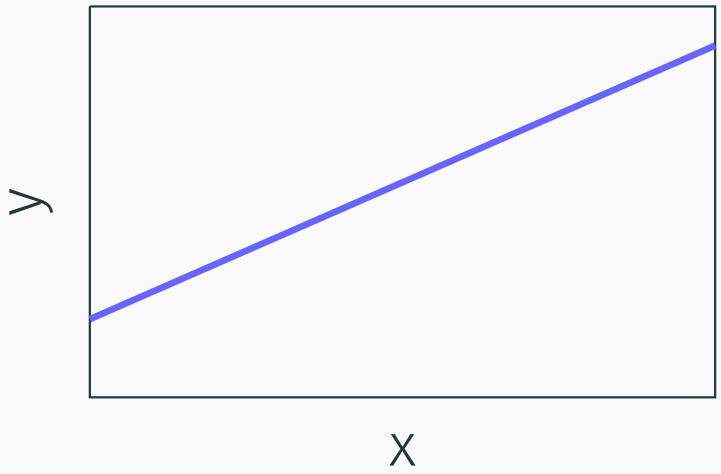

Building a neural network: Linear regression

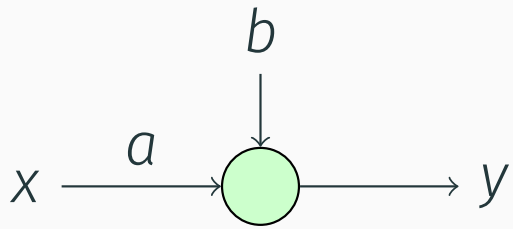

$$y = ax + b$$

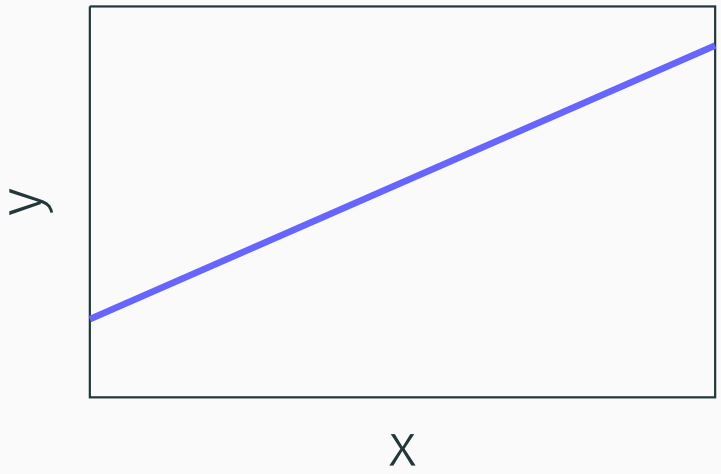

Building a neural network: Artificial neuron

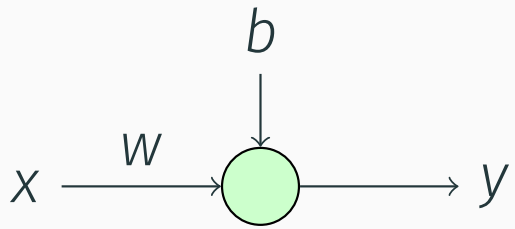

$$y = max(0, wx + b)$$

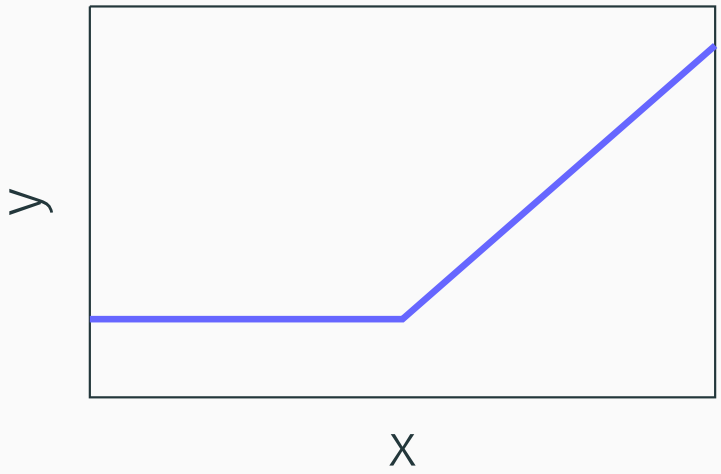

Building a neural network: Artificial neural network (ANN)

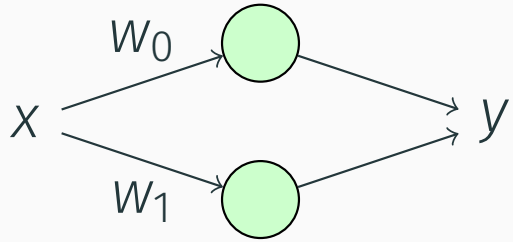

$$y = max(0, w_0 x + b_0) + max(0, w_1 x + b_1)$$

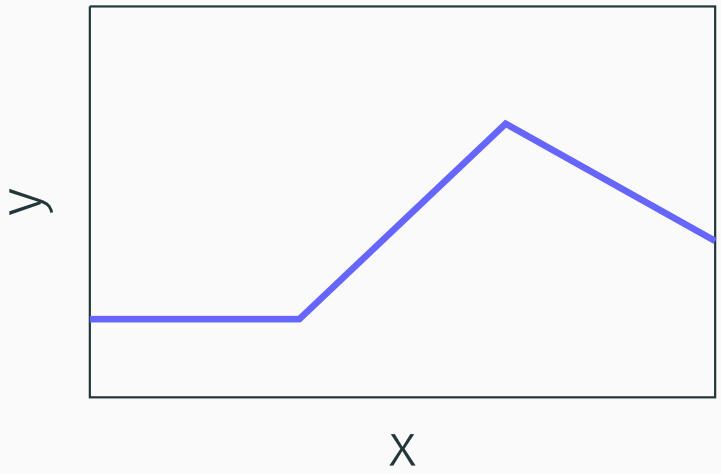

Building a neural network: Universal approximation theorem

”Any relationship that can be described with a polynomial function can be approximated by a neural network with a single hidden layer.”

- Some guy in the 80s, probably

Building a neural network: Universal approximation theorem

”Any relationship that can be described with a polynomial function can be approximated by a neural network with a single hidden layer.”

- Some guy in the 80s, probably

Building a neural network: Increasing dimensionality

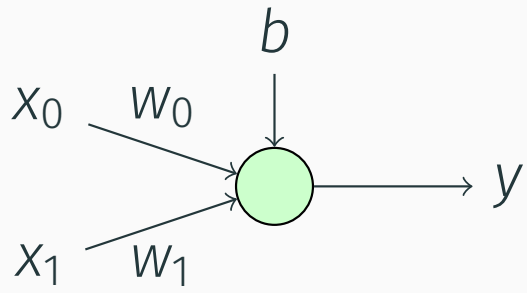

$$y = max(0, w_0 x_0 + w_1 x_1 + b)$$

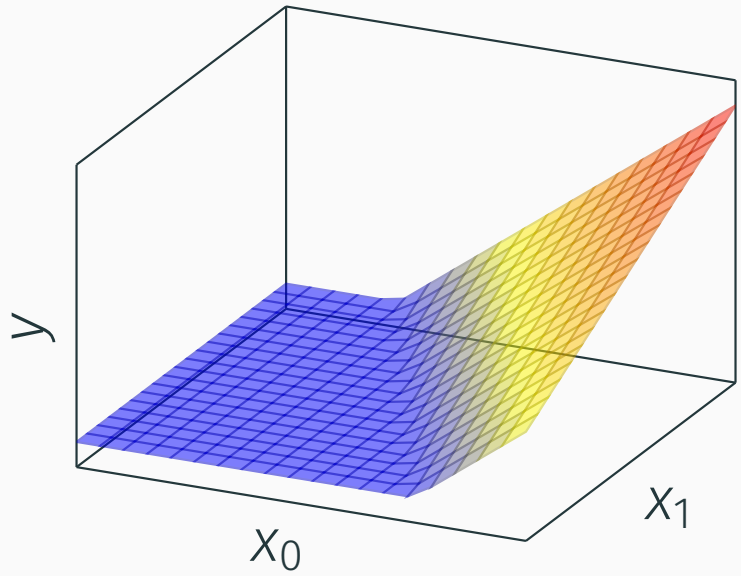

Building a neural network

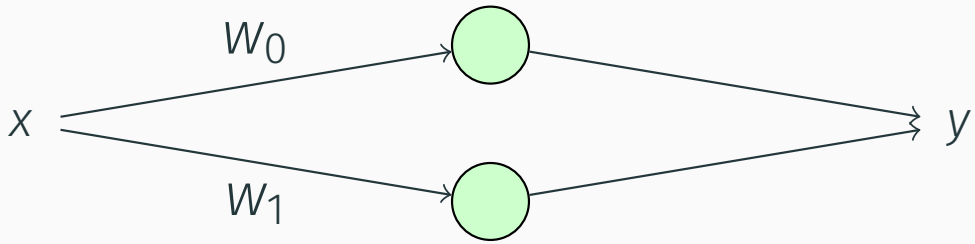

$$y = max(0, w_0 x + b_0) + max(0, w_1 x + b_1)$$

Building a neural network

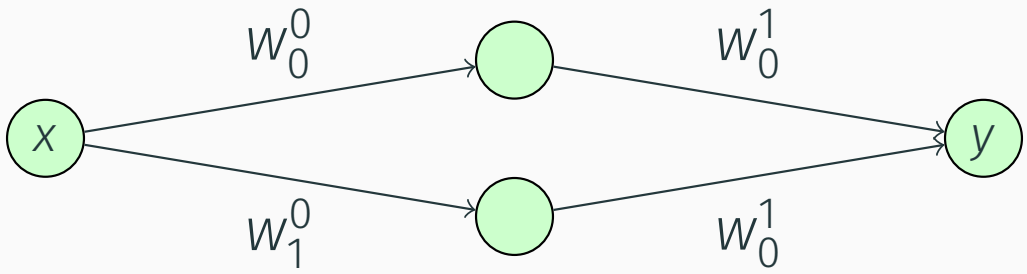

$$y = max(0, w^1_{0,0} * max(0, w^0_{0,0} * x + b_{0,0}) + w^1_{1,0} * max(0, w^0_{0,1} * x + b_{1,0}) + b_1)$$

Building a neural network

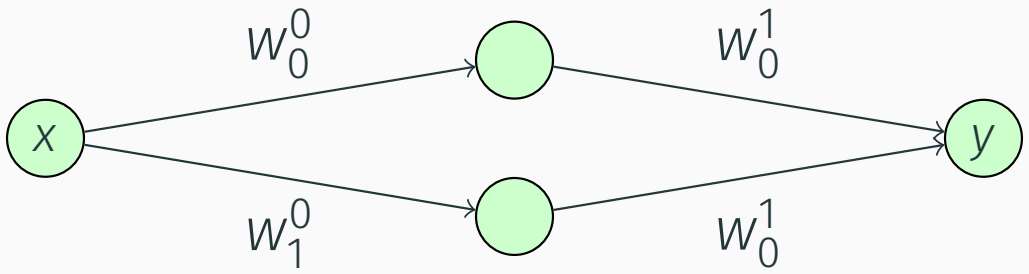

$$y = max(0, w^1_{0,0} * max(0, w^0_{0,0} * x + b_{0,0}) + w^1_{1,0} * max(0, w^0_{0,1} * x + b_{1,0}) + b_1)$$

Building a neural network

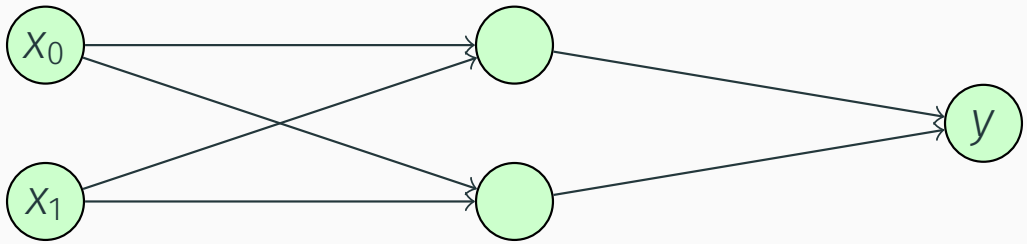

$$y = max(0, w^1_{0,0} * max(0, w^0_{0,0} * x_0 + w^0_{1,0} * x_1 + b_{0,0}) + w^1_{1,0} * max(0, w^0_{0,1} * x_0 + w^0_{1,1} * x_1 + b_{0,1}) + b_1)$$

Building a neural network

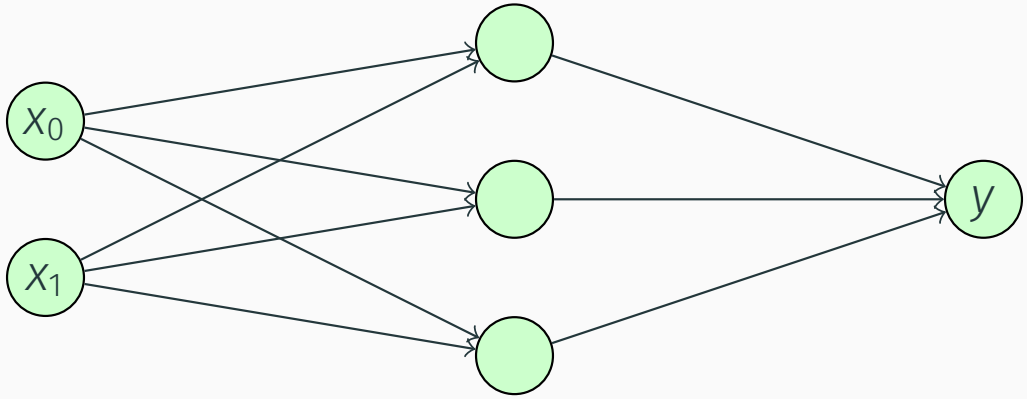

$$y = max(0, w^1_{0,0} * max(0, w^0_{0,0} * x_0 + w^0_{1,0} * x_1 + b_{0,0}) + w^1_{1,0} * max(0, w^0_{0,1} * x_0 + w^0_{1,1} * x_1 + b_{0,1}) + w^1_{2,0} * max(0, w^0_{0,2} * x_0 + w^0_{1,2} * x_1 + b_{0,2}) + b_1)$$

Building a neural network

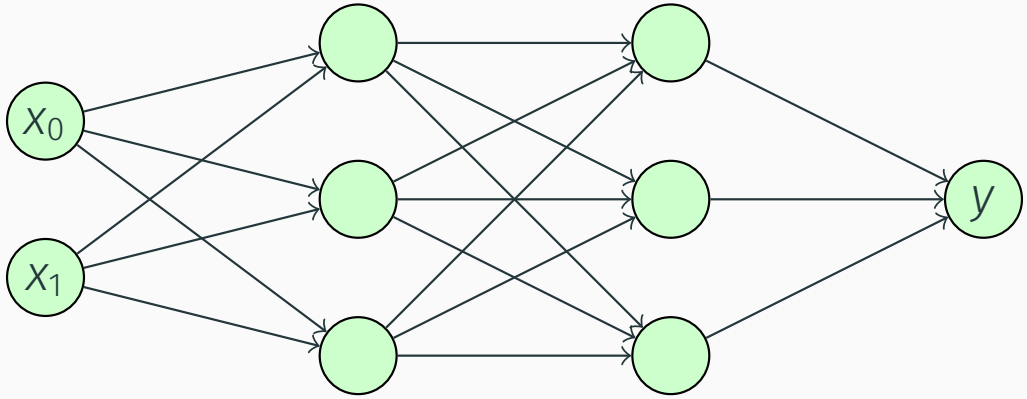

$$
\begin{aligned}
y = max(0, & w^2_{0,0} * max(0, w^1_{0,0} * max(0, w^0_{0,0} * x_0 + w^0_{1,0} * x_1 + b_{0,0}) + \\
& w^1_{1,0} * max(0, w^0_{0,1} * x_0 + w^+_{1,1} * w_1 + b_{0,1}) + \\
& w^1_{2,0} * max(0, w^0_{0,2} * x_0 + w^+_{1,2} * w_1 + b_{0,2}) + \\
& b_{1,0}) + \\
& w^2_{1,0} * max(0, w^1_{0,1} * max(0, w^0_{0,0} * x_0 + w^0_{1,0} * x_1 + b_{0,0}) + \\
& w^1_{1,1} * max(0, w^0_{0,1} * x_0 + w^+_{1,1} * w_1 + b_{0,1}) + \\
& w^1_{2,1} * max(0, w^0_{0,2} * x_0 + w^+_{1,2} * w_1 + b_{0,2}) + \\
& b_{1,1}) + \\
& w^2_{2,0} * max(0, w^1_{0,2} * max(0, w^0_{0,0} * x_0 + w^0_{1,0} * x_1 + b_{0,0}) + \\
& w^1_{1,2} * max(0, w^0_{0,1} * x_0 + w^+_{1,1} * w_1 + b_{0,1}) + \\
& w^1_{2,2} * max(0, w^0_{0,2} * x_0 + w^+_{1,2} * w_1 + b_{0,2}) + \\
& b_{1,2}) + \\
& b_2)
\end{aligned}
$$

Building a neural network: Summary

- Artificial neurons are (linear) weighted sums wrapped in non-linear activation functions
- Multiple artificial neurons stacked together in a layerwise fashion comprise an artificial neural network
- Artificial neural networks allow us to model arbitrarily complex relationships between inputs and outputs

Training a neural network: Loss functions

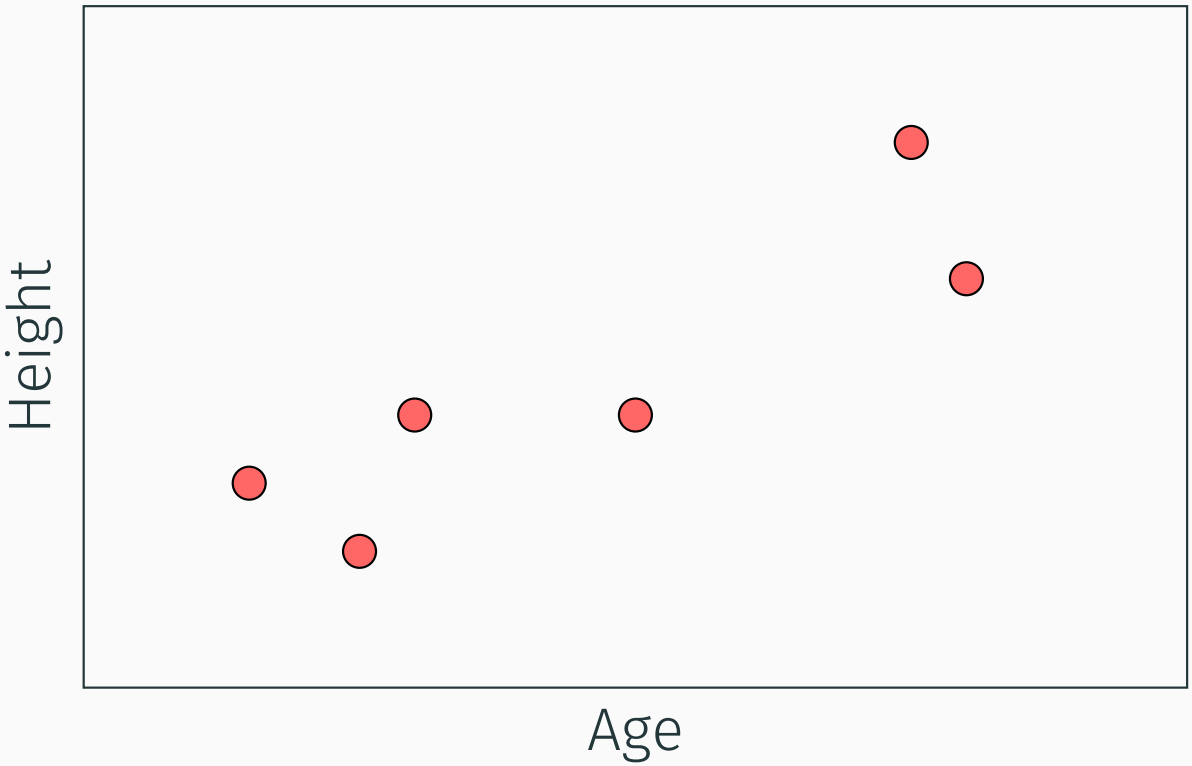

Training a neural network: Loss functions

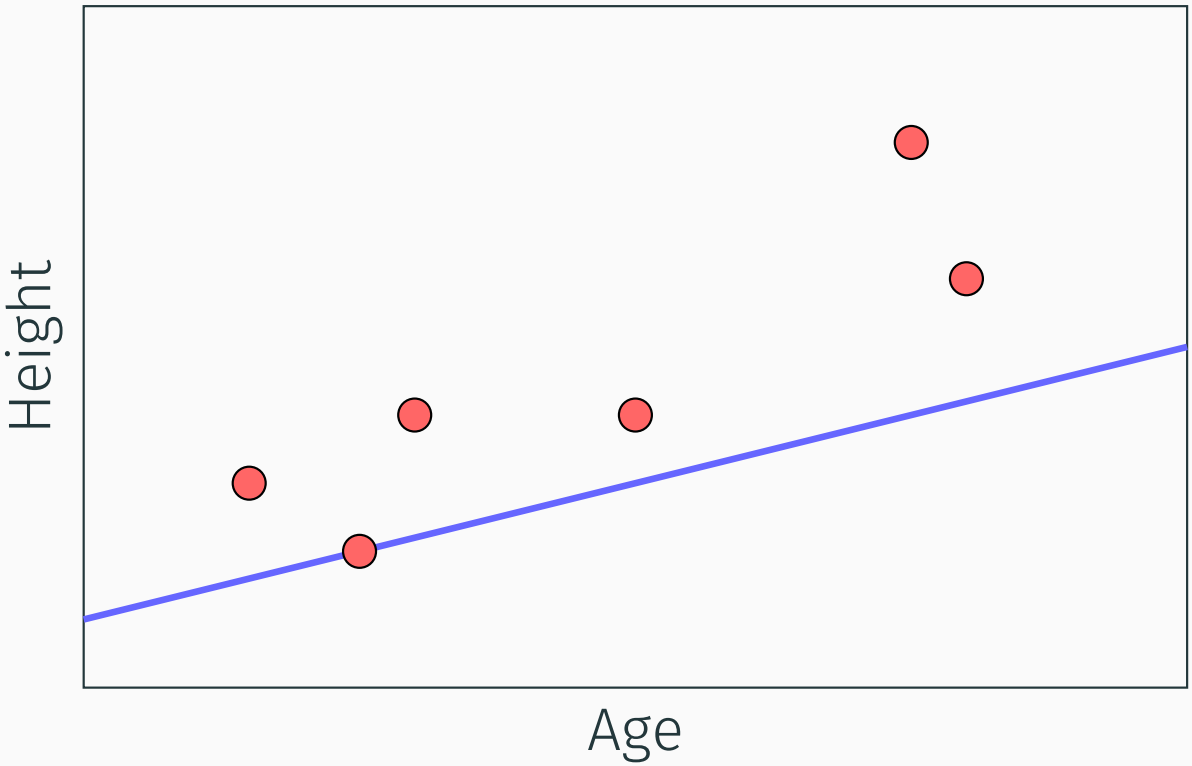

$$\hat{y} = 0.4x + 0.1$$

Training a neural network: Loss functions

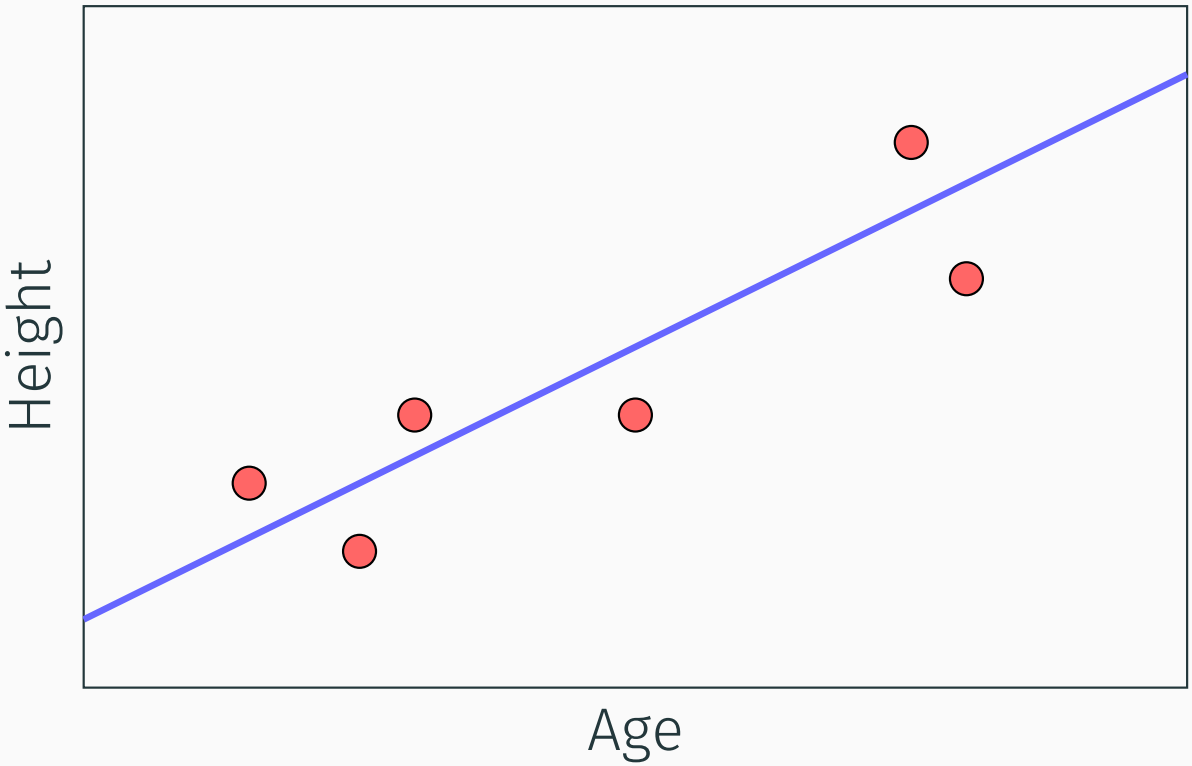

$$\hat{y} = 0.9x + 0.1$$

Training a neural network: Loss functions

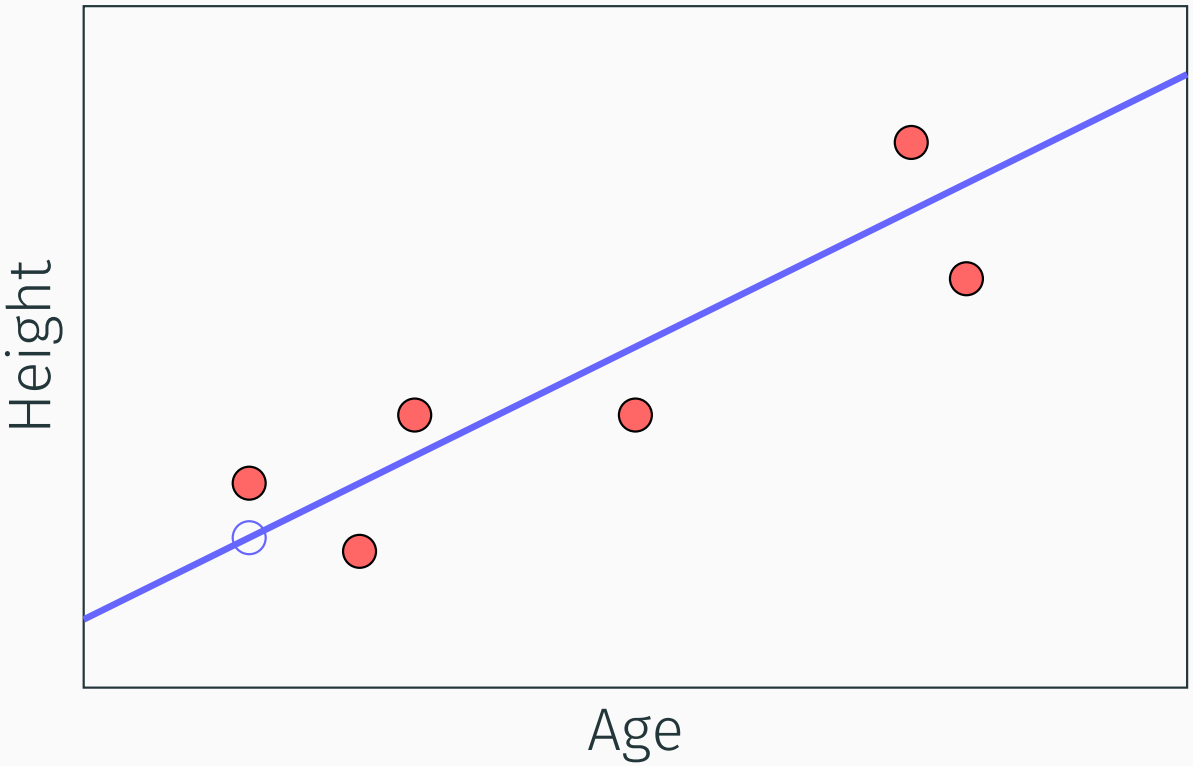

$$\hat{y} = 0.9x + 0.1$$

Training a neural network: Loss functions

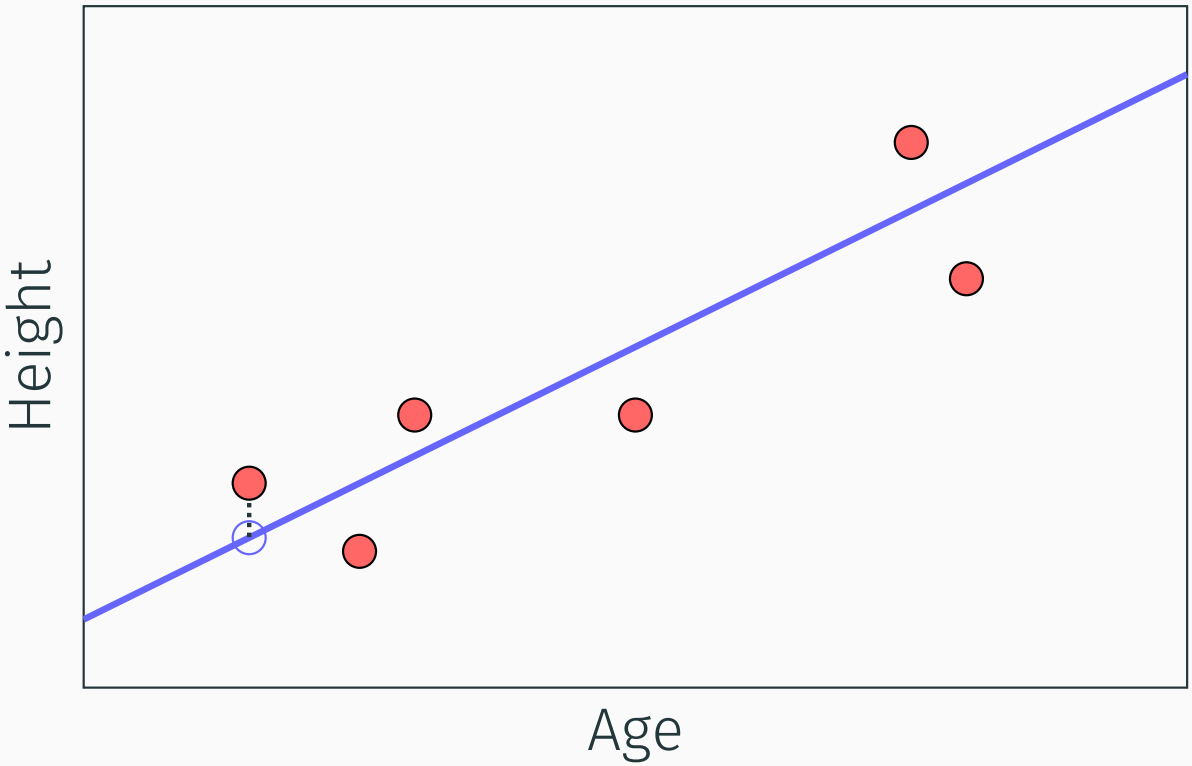

$$\hat{y} = 0.9x + 0.1$$

Training a neural network: Loss functions

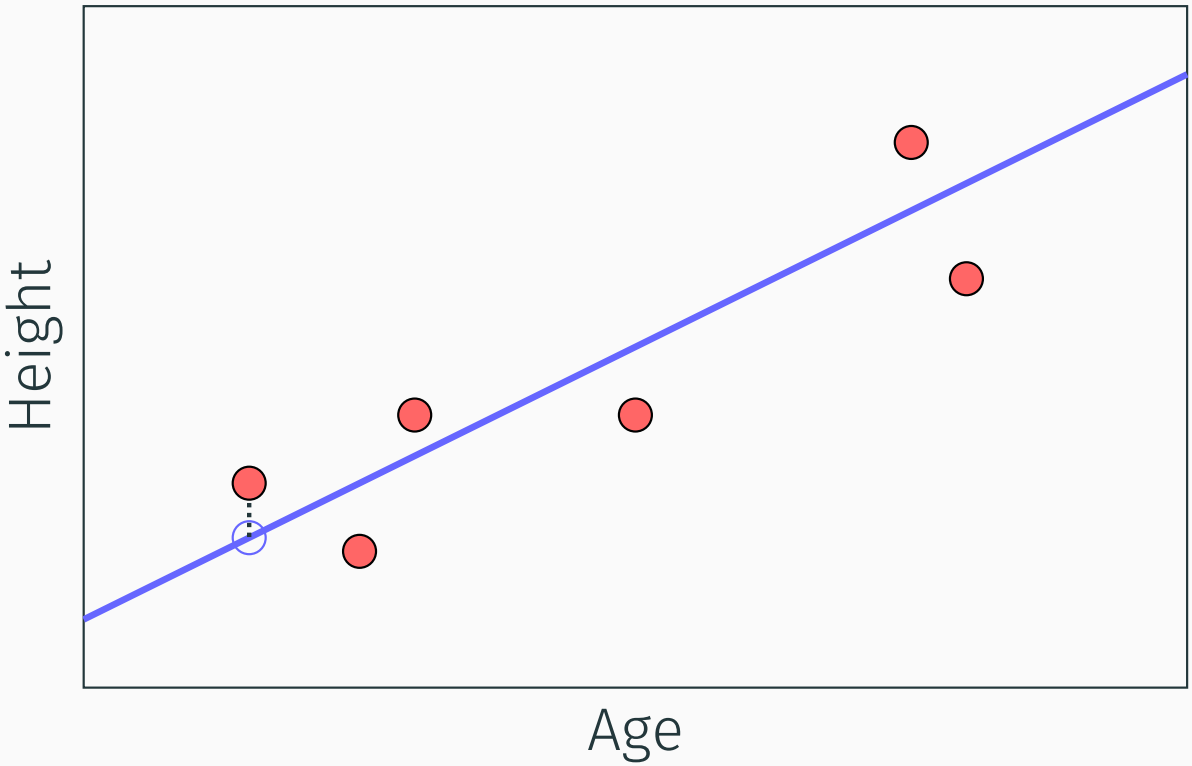

$$\hat{y} = 0.9x + 0.1$$

$$error_i = |\hat{y} - y|$$

Training a neural network: Loss functions

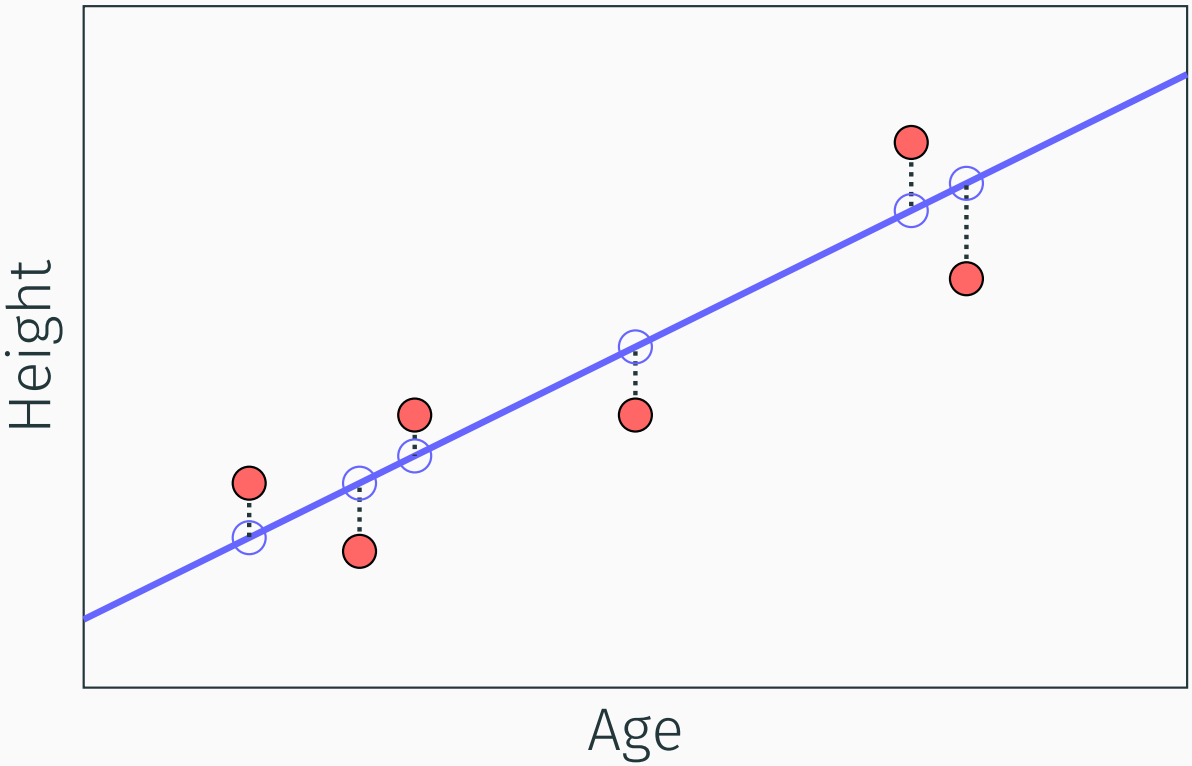

$$\hat{y} = 0.9x + 0.1$$

$$error_i = |\hat{y} - y|$$

$$mse = \frac{1}{n}\sum_{i<n}(\hat{y}_i - y_i)^2$$

Training a neural network: Gradient descent

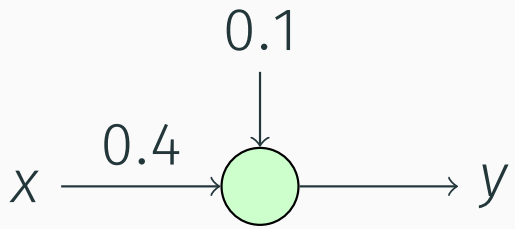

$$\hat{y} = 0.4x + 0.1 \qquad loss = (y - \hat{y})^2$$

Training a neural network: Gradient descent

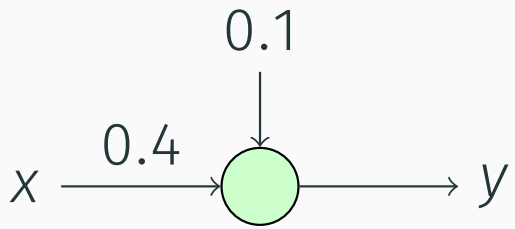

$$0.16 = 0.4 * 0.15 + 0.1 \qquad 0.019 = (0.3 - 0.16)^2$$

Training a neural network: Gradient descent

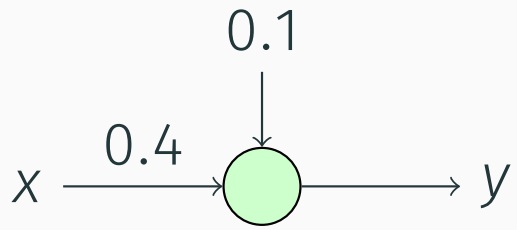

$$(0.3 - (0.4 * 0.15 + 0.1))^2 = 0.019$$

Training a neural network: Gradient descent

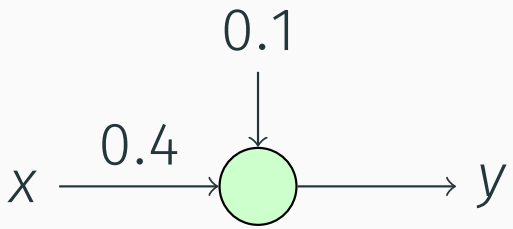

$$(\underset{\substack{\uparrow \\ y}}{0.3} - (\underbrace{\overset{\substack{w \\ \downarrow}}{0.4} * \overset{\substack{x \\ \downarrow}}{0.15} + \overset{\substack{b \\ \downarrow}}{0.1}}_{\substack{\uparrow \\ \hat{y}}}))^2 = \overset{\substack{loss \\ \downarrow}}{0.019}$$

Training a neural network: Gradient descent

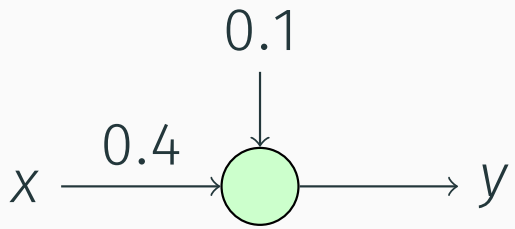

$$(0.3 - (0.4 * 0.15 + 0.1))^2 = 0.019$$

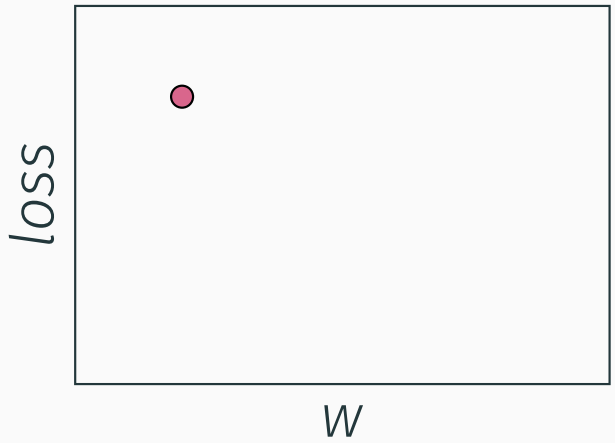

Training a neural network: Gradient descent

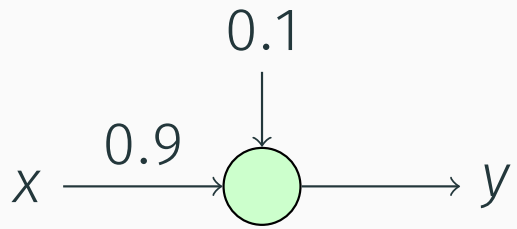

$$(0.3 - (0.9 * 0.15 + 0.1))^2 = 0.004$$

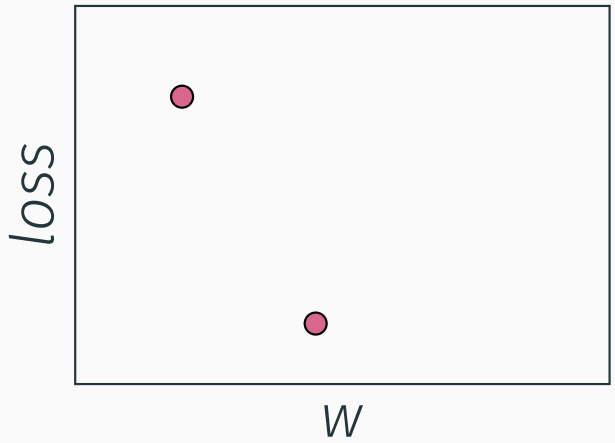

Training a neural network: Gradient descent

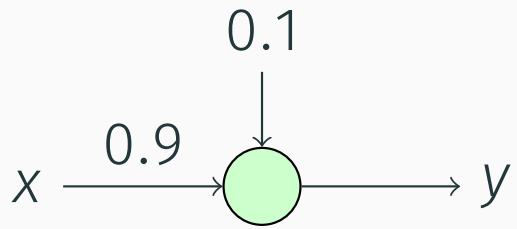

$$(0.3 - (0.9 * 0.15 + 0.1))^2 = 0.004$$

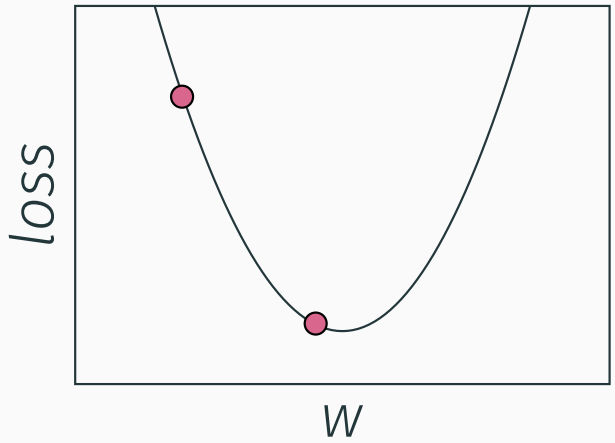

Training a neural network: Gradient descent

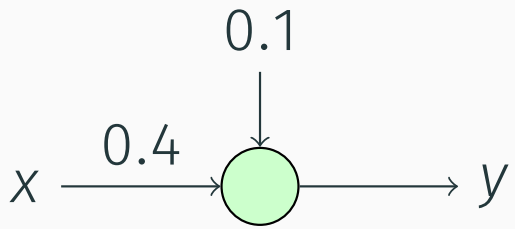

$$(0.3 - (0.4 * 0.15 + 0.1))^2 = 0.019$$

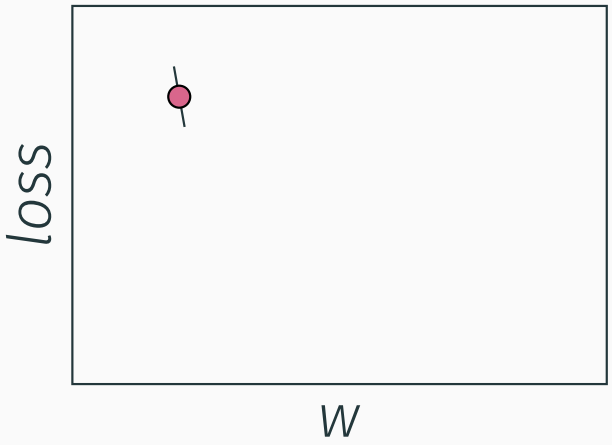

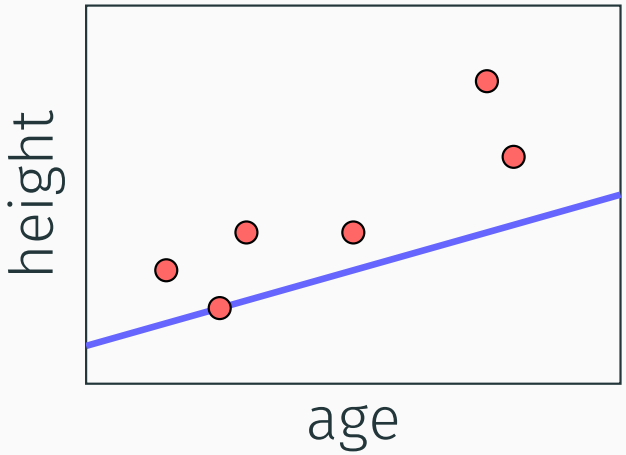

Training a neural network: Gradient descent

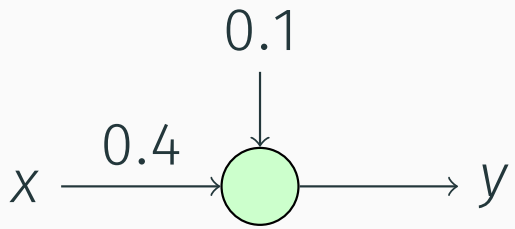

$$(0.3 - (0.4 * 0.15 + 0.1))^2 = 0.019$$

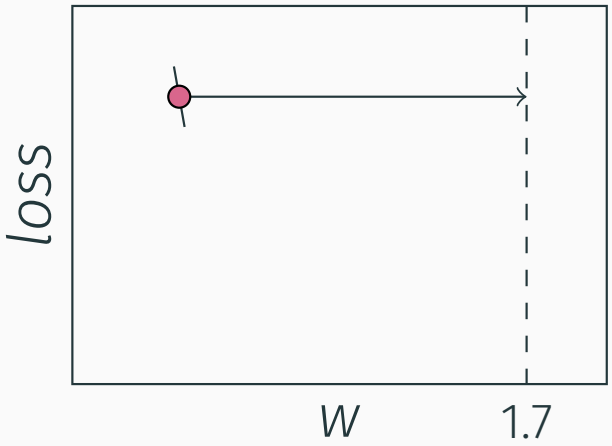

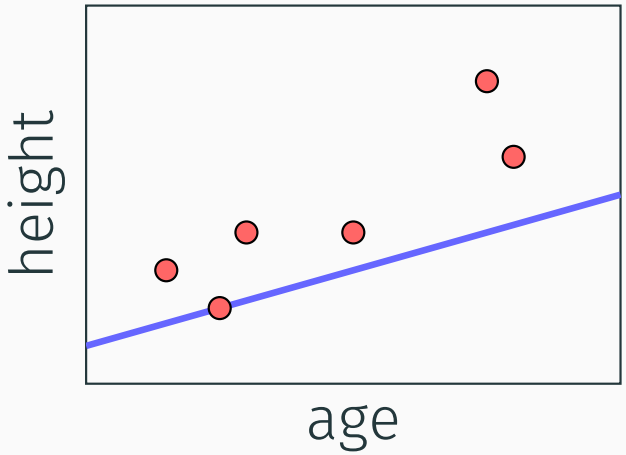

Training a neural network: Gradient descent

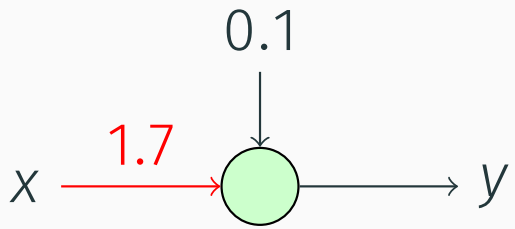

$$(0.3 - (1.7 * 0.15 + 0.1))^2 = 0.003$$

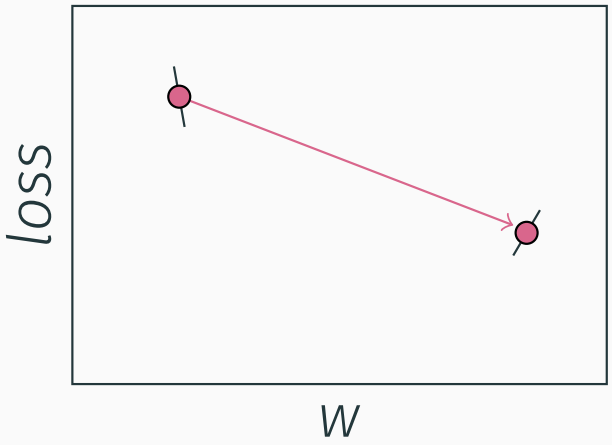

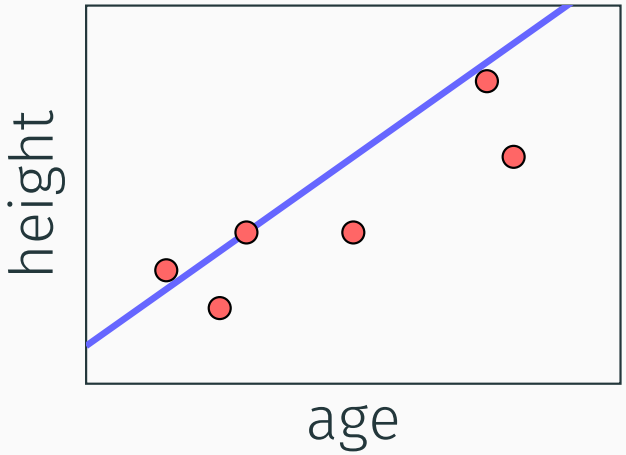

Training a neural network: Gradient descent

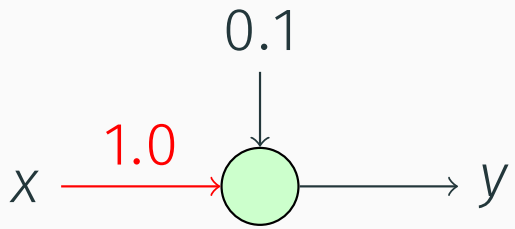

$$(0.3 - (1.0 * 0.15 + 0.1))^2 = 0.002$$

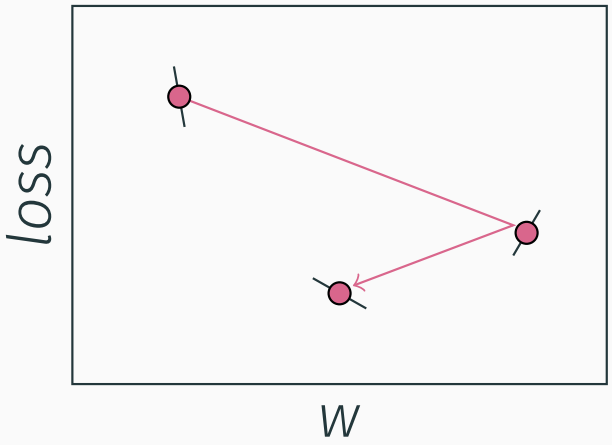

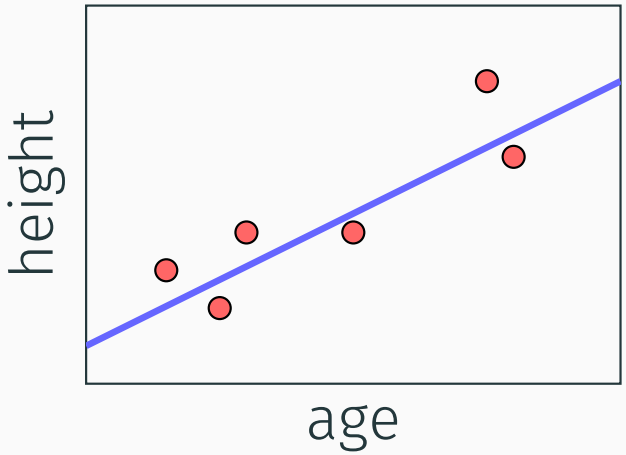

Training a neural network: Gradient descent

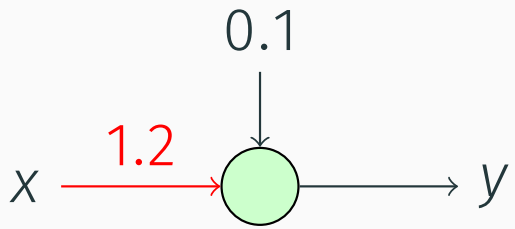

$$(0.3 - (1.2 * 0.15 + 0.1))^2 = 0.000$$

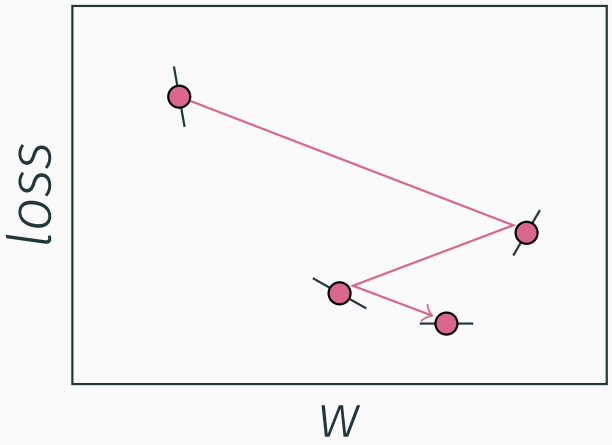

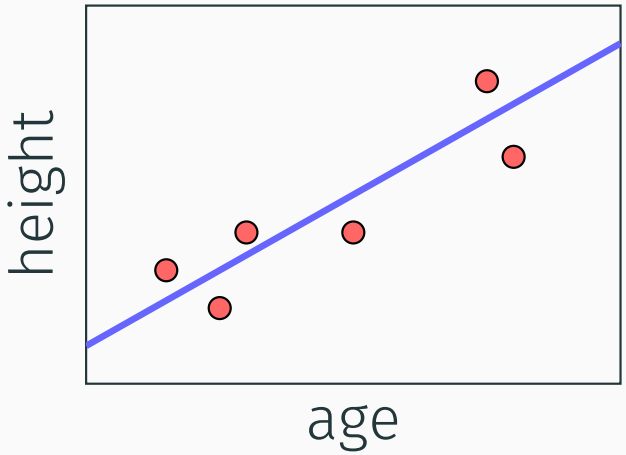

Training a neural network: Summary

- The loss function is a precise formalization of what we want the model to learn
- Calculating the gradient allows us to see how we can update the parameters of the model to decrease the loss
- Using gradient descent we can update the parameters sequentially until we have the perfect model with zero loss

Convolutional neural networks: Architecture

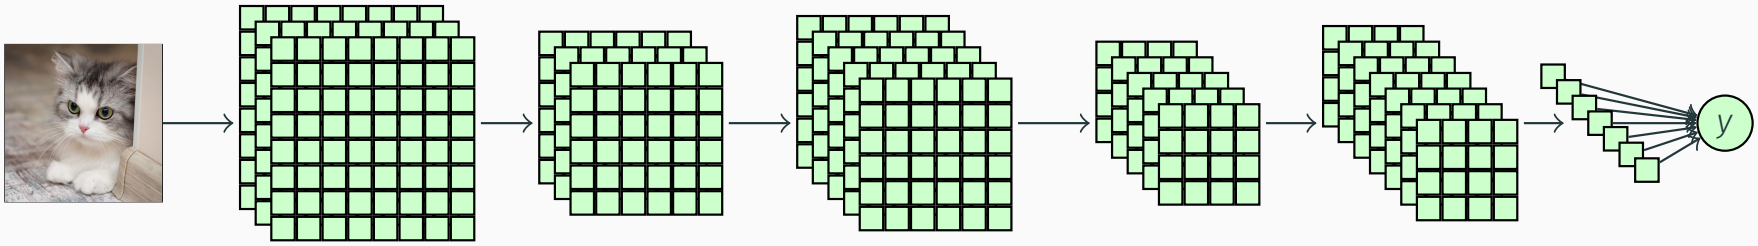

Convolutional neural networks: Images

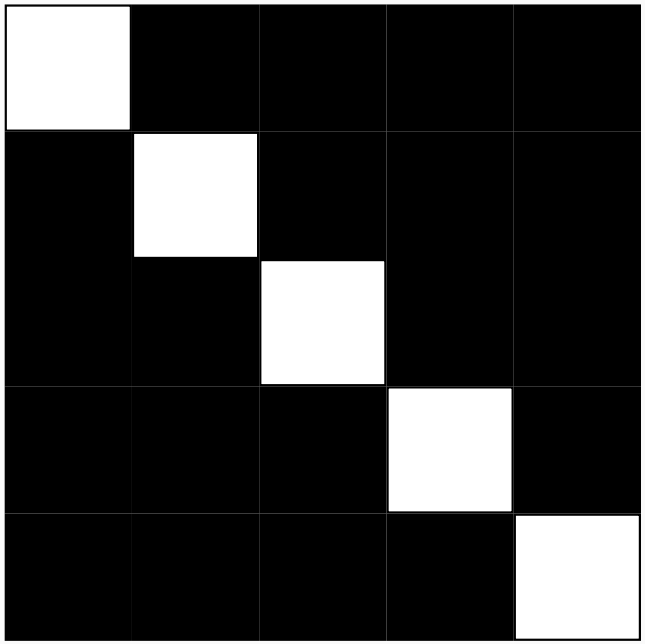

1	0	0	0	0
0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	1

Convolutional neural networks: Images

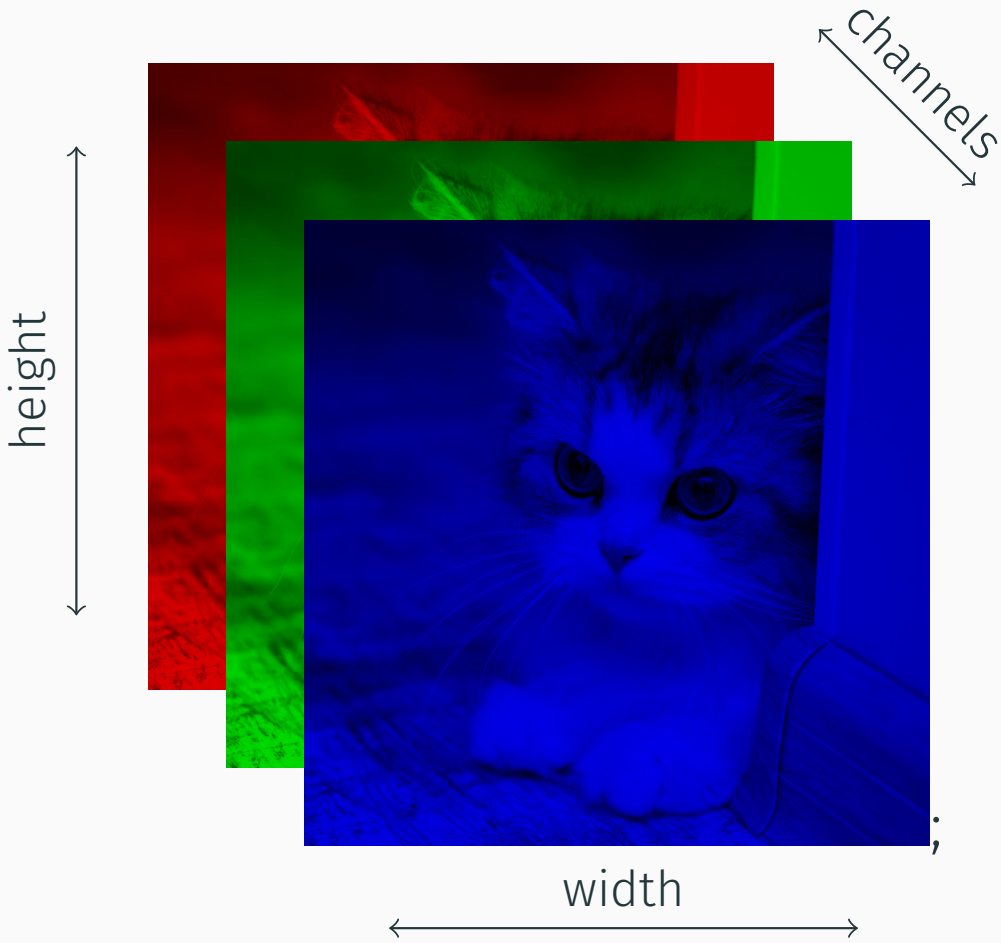

Convolutional neural networks: Convolution

0	0	1
0	1	0
1	0	0

$*$

1	0	0
0	1	0
0	0	1

Convolutional neural networks: Convolution

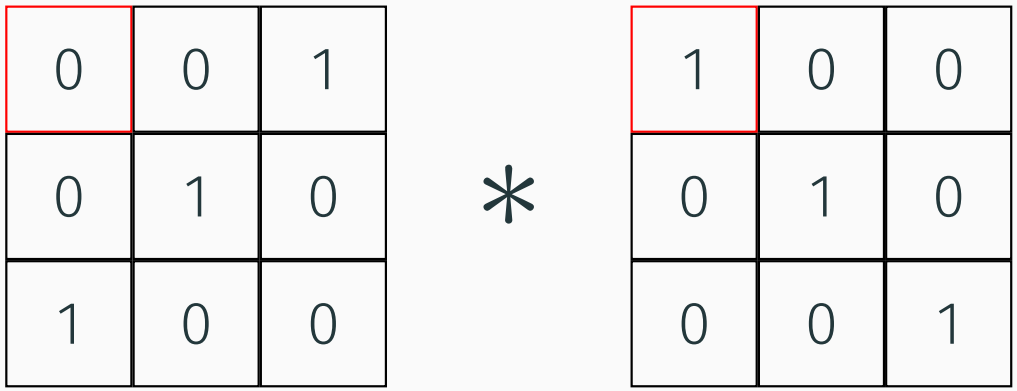

0*1

Convolutional neural networks: Convolution

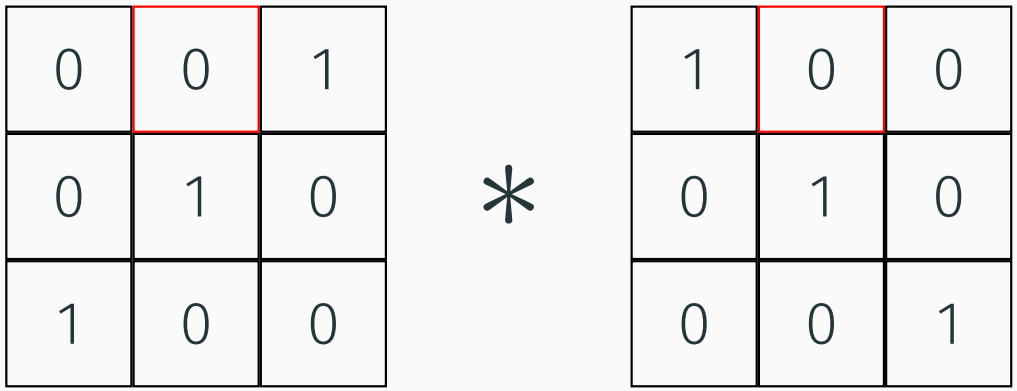

0*1+0*0

Convolutional neural networks: Convolution

0	0	1
0	1	0
1	0	0

$*$

1	0	0
0	1	0
0	0	1

0*1+0*0+1*0

Convolutional neural networks: Convolution

0	0	1
0	1	0
1	0	0

$*$

1	0	0
0	1	0
0	0	1

0*1+0*0+1*0+0*0

Convolutional neural networks: Convolution

0	0	1
0	1	0
1	0	0

$*$

1	0	0
0	1	0
0	0	1

0*1+0*0+1*0+0*0+1*1+0*0+1*0+0*0+0*1=1

Convolutional neural networks: Convolution

1	0	0
0	1	0
0	0	1

$*$

1	0	0
0	1	0
0	0	1

1*1+0*0+0*0+0*0+1*1+0*0+1*0+0*0+1*1=3

Convolutional neural networks: Convolution

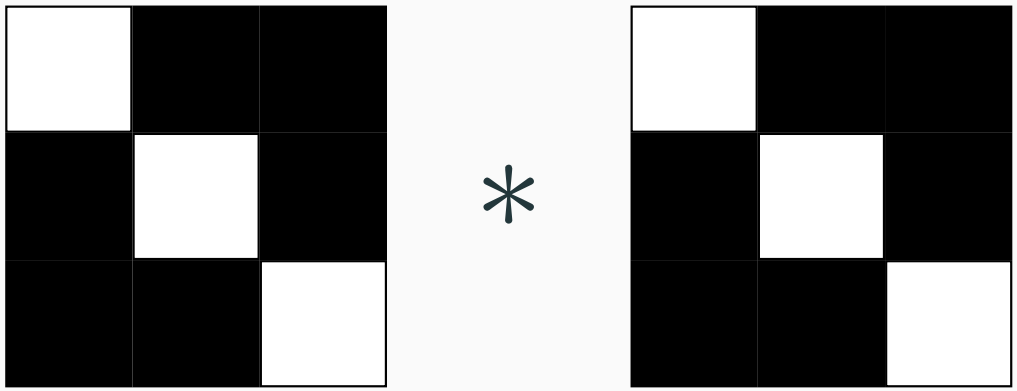

1*1+0*0+0*0+0*0+1*1+0*0+1*0+0*0+1*1=3

Convolutional neural networks: Convolutional layers

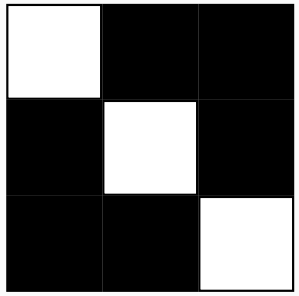
$*$
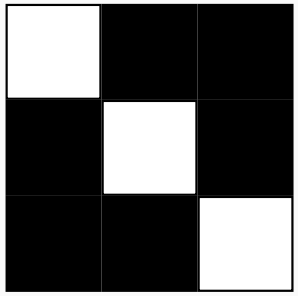
$= 3$

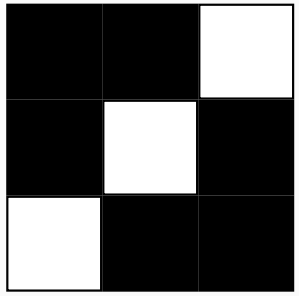
$*$
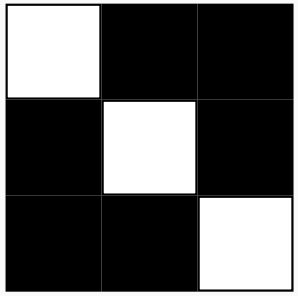
$= 1$

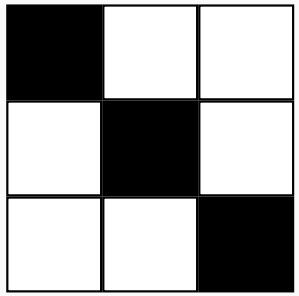
$*$
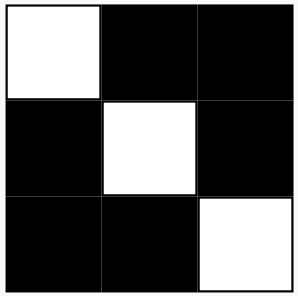
$= 0$

Convolutional neural networks: Convolutional layers

Image

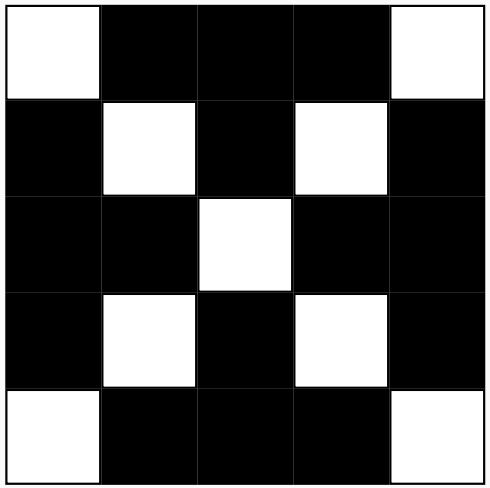

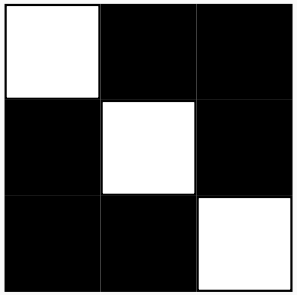

Pattern 1

Convolutional neural networks: Convolutional layers

Image

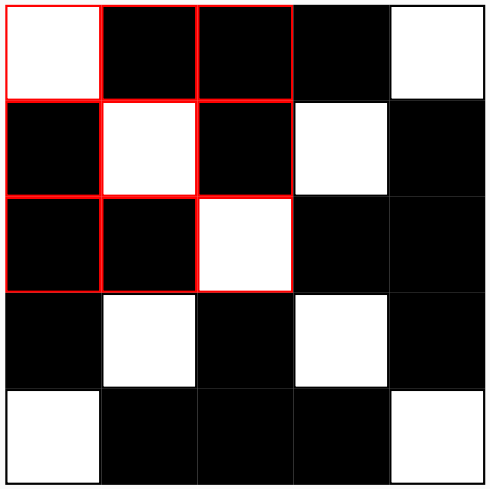

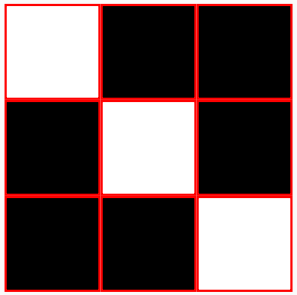

Pattern 1

3

Convolutional neural networks: Convolutional layers

Image

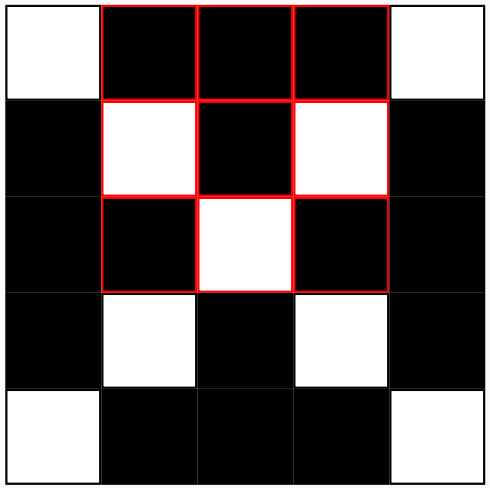

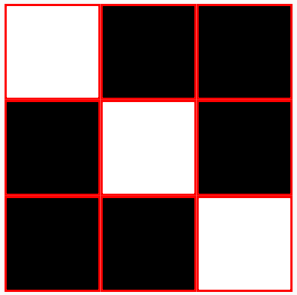

Pattern 1

3	0	

Convolutional neural networks: Convolutional layers

Image

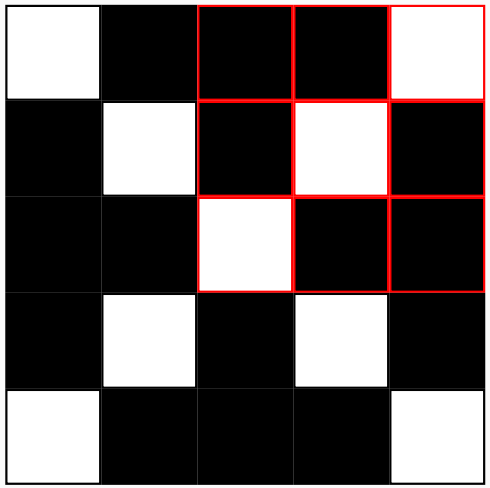

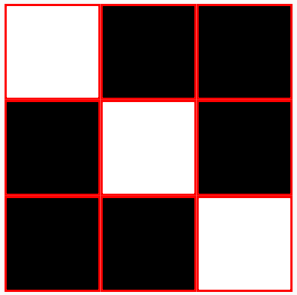

Pattern 1

3	0	1

Convolutional neural networks: Convolutional layers

Image

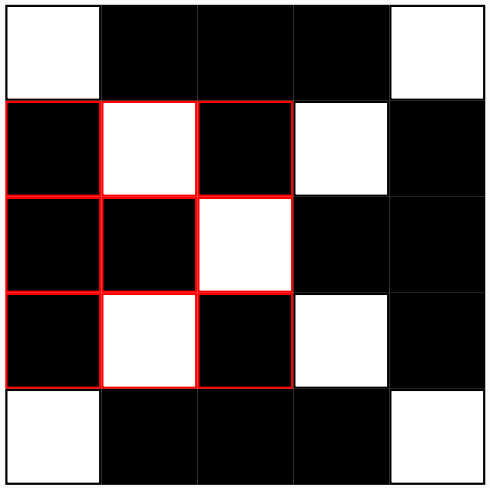

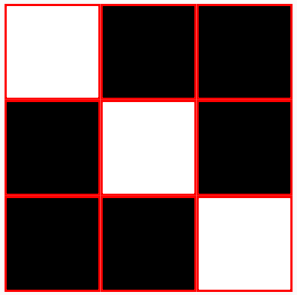

Pattern 1

3	0	1
0		

Convolutional neural networks: Convolutional layers

Image

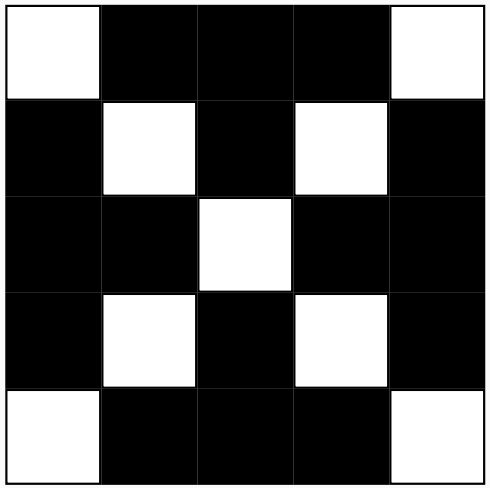

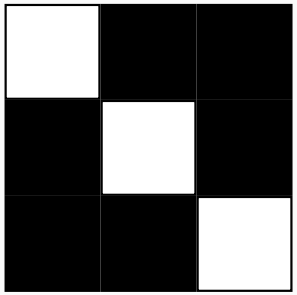

Pattern 1

3	0	1
0	3	0
1	0	3

Convolutional neural networks: Convolutional layers

Image

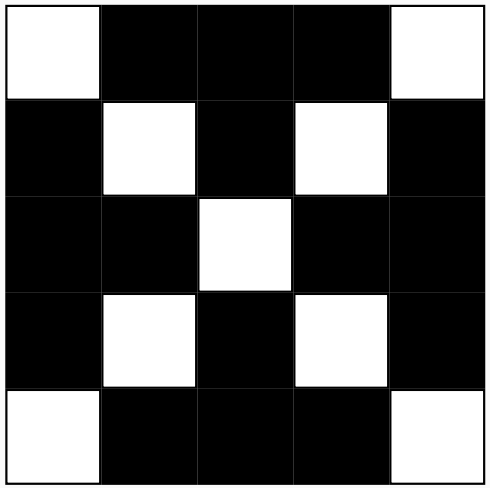

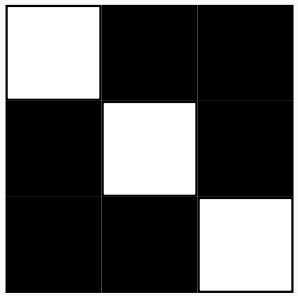

Pattern 1

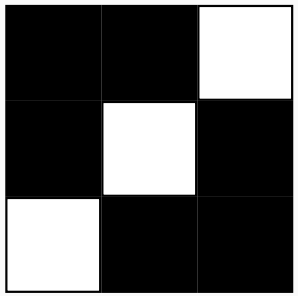

Pattern 2

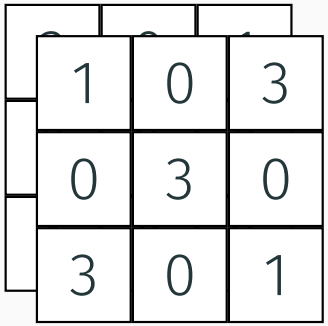

Convolutional neural networks: Convolutional layers

Image

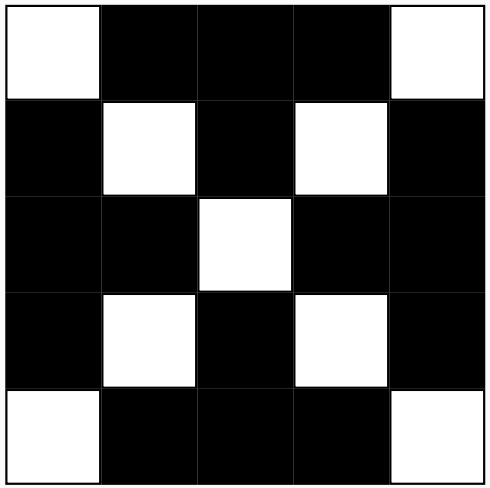

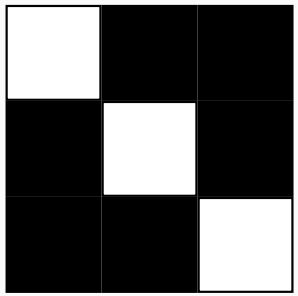

Pattern 1

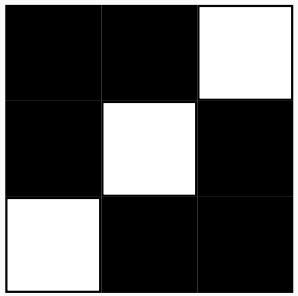

Pattern 2

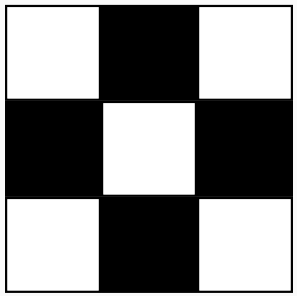

Pattern 3

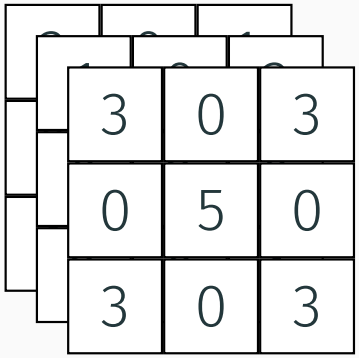

Convolutional neural networks: Convolutional layers

Image

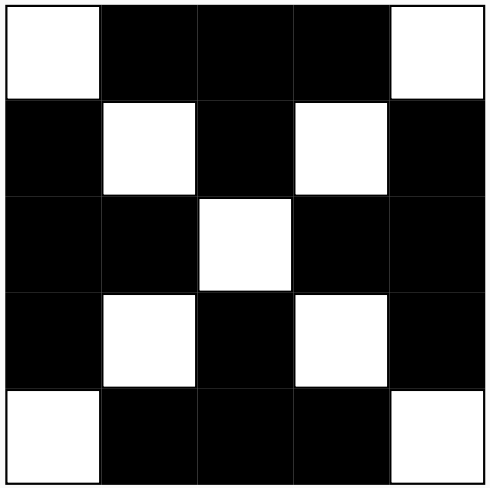

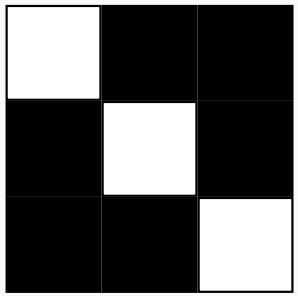

Pattern 1

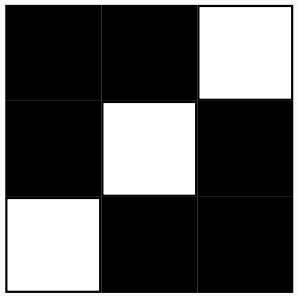

Pattern 2

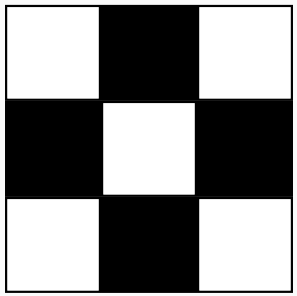

Pattern 3

Feature maps

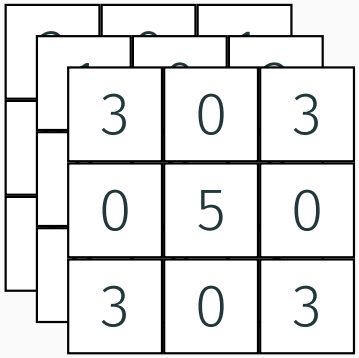

Convolutional neural networks: Convolutional layers

Image

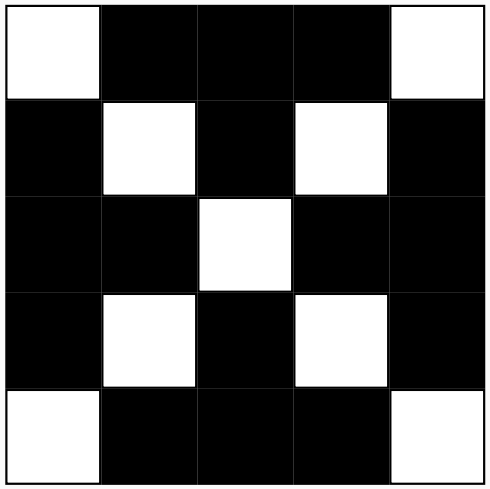

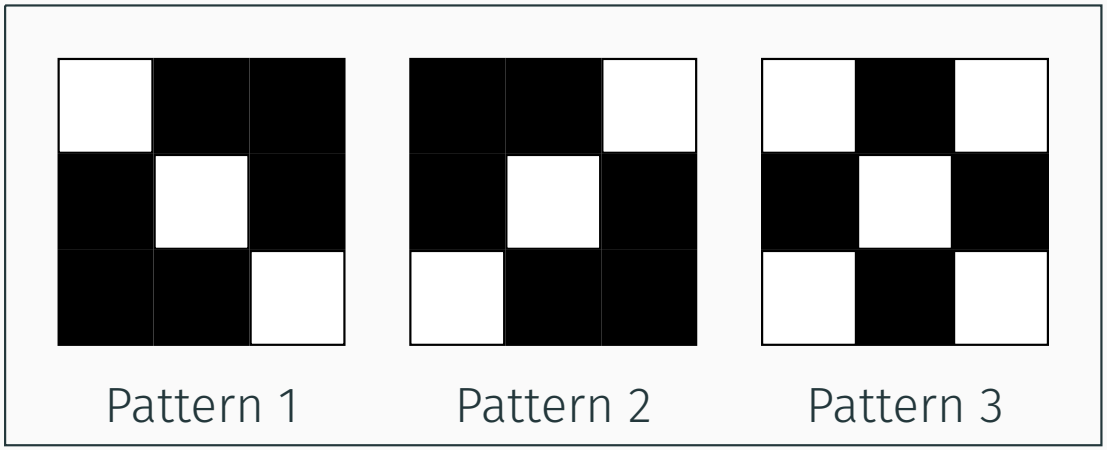

Weights

Feature maps

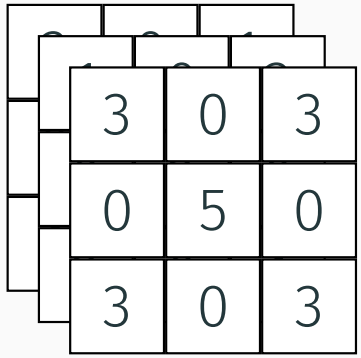

Convolutional neural networks: Convolutional layers

Image

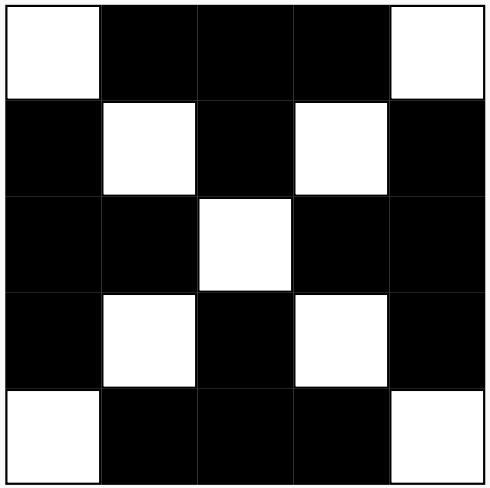

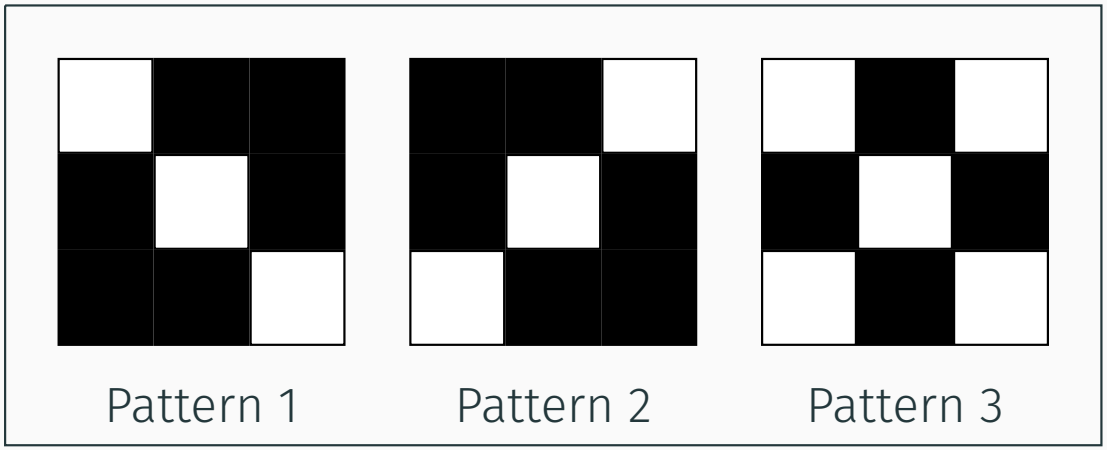

Weights

Feature maps

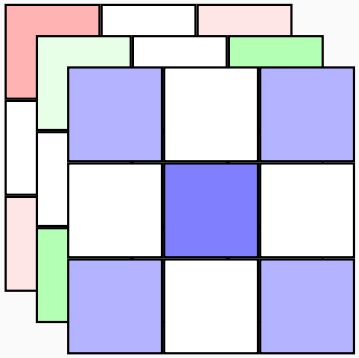

Convolutional neural networks: Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Convolutional neural networks: Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

5

Convolutional neural networks: Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

5	7

Convolutional neural networks: Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

5	7
13	

Convolutional neural networks: Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

5	7
13	15

Convolutional neural networks: Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

5	7
13	15

Convolutional neural networks: Architecture

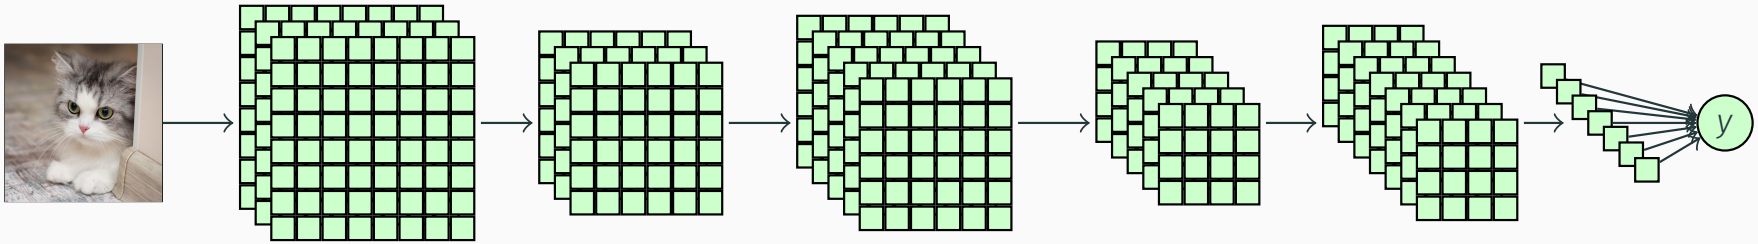

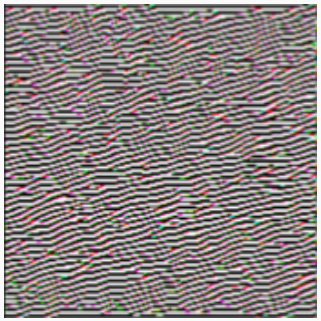

Convolutional neural networks: Architecture

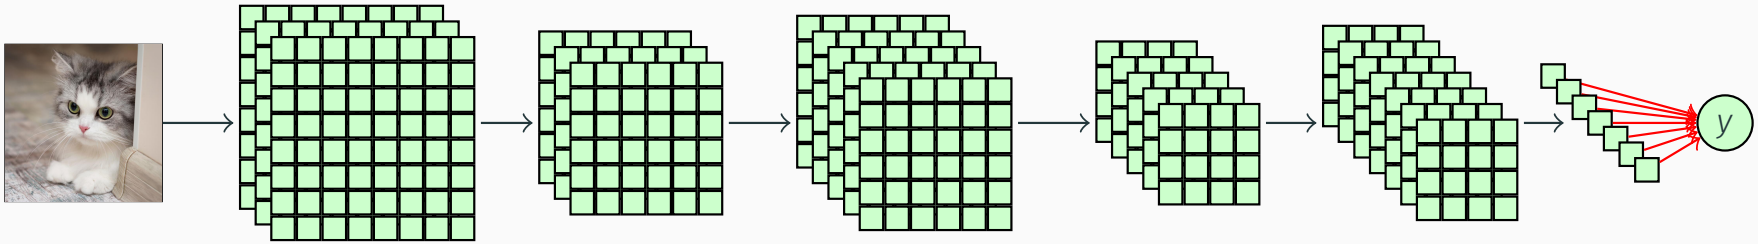

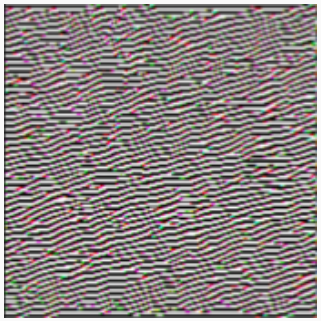

Convolutional neural networks: Summary

- Images in computers are stored as matrices of numbers
- The convolution operation is a pattern matcher
- A convolutional layer is a battery of pattern matchers where the patterns are learned during training
- The pooling operation reduces spatial dimensions and distills information
- A convolutional neural network consists of alternating convolution and pooling, which means looking for more and more abstract patterns spanning larger and larger region of the images
- The final layer of the network makes a prediction based on the patterns it has found