

Introduction to Machine Learning

Image recognition in Python and Tensorflow

Esten H. Leonardsen and Martin Hovin

19.01.23



Esten H. Leonardsen
(UiO and Biometrical AS)

Interests:

- Talking about esoteric theory
- Making deep learning tutorials



Martin Hovin
(Biometrical AS)

Interests:

- Installing tensorflow
- Debugging Estens code

Theory session:

- What is a statistical learning model?
- What is a loss function?
- How do we train a statistical learning model?
- How does a (deep) neural network work?
- What operations does a convolutional neural network use?
- What is transfer learning?
- What is overfitting?
- How do we combat it?

Practical session:

1. Set up a Python-environment containing Tensorflow
2. Use a pretrained convolutional neural network to predict
3. Fit a flower classifier using transfer learning
4. Improve the flower classifier



[illegible]

Summary

- What is a statistical learning model?
A formula representing the relationship between inputs and outputs
- What is a loss function?
A function quantifying how good a model is
- How do we train a statistical learning model?
Gradual updates of parameters using gradient descent
- How does a (deep) neural network work?
Sequentially applying (non-linear) artificial neurons to inputs
- What operations does a convolutional neural network use?
Alternating convolutions and pooling
- What is transfer learning?
Retraining an already trained model for a new problem
- What is overfitting?
When a model learns patterns in the training data that does not hold generally
- How do we combat it?
Rigorous testing, regularization and data augmentation