

PSY9511: Seminar 3

Variable selection and regularization

Esten H. Leonardsen

07.09.23

Outline

1. Introduction

- Python
- Coding tips: Separation of concerns

2. Variable selection

- Best subset selection
- Forward stepwise selection
- Backward stepwise selection

3. Regularization

- Ridge regression
- Lasso
- Elastic net

4. Dimensionality reduction

- Principal component regression
- Partial least squares

Python: Imports

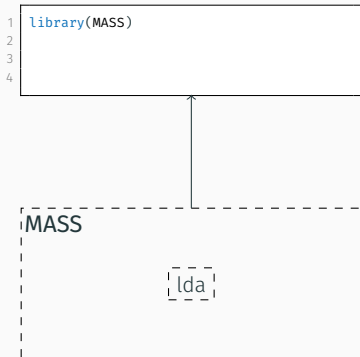
```
1  
2  
3  
4
```



```
MASS  
  
lda
```



Python: Imports



Python: Imports

```
1 library(MASS)
2
3 lda_fit <- lda(display ~ age + fb,
4               data = display)
```

MASS

lda

Python: Imports

```
1 library(MASS)
2
3 lda_fit <- lda(display ~ age + fb,
4               data = display)
```

MASS

lda

```
In[1]: from sklearn import *

lda = discriminant_analysis.
       LinearDiscriminantAnalysis()
lda.fit(display[['age', 'fb']],
        display['display'])
```

sklearn

discriminant_analysis

LinearDiscriminantAnalysis

Python: Imports

```
1 library(MASS)
2
3 lda_fit <- lda(display ~ age + fb,
4               data = display)
```

MASS

lda

```
In[1]: import sklearn

lda = sklearn.discriminant_analysis.
        LinearDiscriminantAnalysis()
lda.fit(display[['age', 'fb']],
        display['display'])
```

sklearn

discriminant_analysis

LinearDiscriminantAnalysis

Python: Imports

```
1 library(MASS)
2
3 lda_fit <- lda(display ~ age + fb,
4               data = display)
```

MASS

lda

```
In[1]: from sklearn.discriminant_analysis \
import LinearDiscriminantAnalysis

lda = LinearDiscriminantAnalysis()
lda.fit(display[['age', 'fb']],
        display['display'])
```

sklearn

discriminant_analysis

LinearDiscriminantAnalysis

Python: pandas

```
1 path <- '/Users/esten/Downloads/Auto.csv'  
2 df <- read.csv(path)  
3 head(df, 10)
```

	mpg	cylinders	displacement	horsepower
1	18	8	307.0	130
2	15	8	350.0	165
3	18	8	318.0	150
4	16	8	304.0	150
5	17	8	302.0	140
6	15	8	429.0	198
7	14	8	454.0	220
8	14	8	440.0	215
9	14	8	455.0	225
10	15	8	390.0	190

```
In[1]: import pandas as pd
```

```
path = '/Users/esten/Downloads/Auto.csv'  
df = pd.read_csv(path)  
df.head(10)
```

```
Out[1]:
```

	mpg	cylinders	displacement	horsepower
0	18	8	307.0	130
1	15	8	350.0	165
2	18	8	318.0	150
3	16	8	304.0	150
4	17	8	302.0	140
5	15	8	429.0	198
6	14	8	454.0	220
7	14	8	440.0	215
8	14	8	455.0	225
9	15	8	390.0	190

Python: numpy

```
In[1]: import numpy as np
```

```
In[2]: np.random.seed(42)
```

```
In[3]: np.arange(0, 10, 1)
```

```
Out[1]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In[4]: np.isnan([0, 1, np.nan, 3])
```

```
Out[2]: array([False, False,  True, False])
```

```
In[5]: np.amin([1, 0, 3, 2])
```

```
Out[3]: 0
```

```
Out[6]: np.argmin([1, 0, 3, 2])
```

```
Out[4]: 1
```

```
In[7]: np.nanmin([1, 0, 3, np.nan])
```

```
Out[5]: 0
```

Python: statsmodels

```
1 path <- '/Users/esten/Downloads/Auto.csv'
2 data <- read.csv(path)
3
4 model <- lm(mpg ~ cylinders + displacement +
5             horsepower + weight +
6             acceleration + year,
7             data=data)
8 summary(model)
```

Coefficients:

	Estimate	Std. Error	Pr(> t)
(Intercept)	-1.454e+01	4.764e+00	0.00244 **
cylinders	-3.299e-01	3.321e-01	0.32122
displacement	7.678e-03	7.358e-03	0.29733
horsepower	-3.914e-04	1.384e-02	0.97745
weight	-6.795e-03	6.700e-04	< 2e-16 ***
acceleration	8.527e-02	1.020e-01	0.40383
year	7.534e-01	5.262e-02	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*'
0.05

```
In[1]: import statsmodels.formula.api as smf

path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

model = smf.ols(
    formula='mpg ~ cylinders + displacement +
            horsepower + weight +
            acceleration + year',
    data=df
)
fit = model.fit()
print(fit.summary())
```

```
Out[1]:
```

	coef	std err	P> t	[0.025	0.975]
Intercept	-14.5353	4.764	0.002	-23.90	-5.16
cylinders	-0.3299	0.332	0.321	-0.98	0.32
displacement	0.0077	0.007	0.297	-0.00	0.02
horsepower	-0.0004	0.014	0.977	-0.02	0.02
weight	-0.0068	0.001	0.000	-0.00	-0.00
acceleration	0.0853	0.102	0.404	-0.11	0.28
year	0.7534	0.053	0.000	0.65	0.85

Python: scikit-learn

```
In[1]: from sklearn.linear_model import LinearRegression

path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

model = LinearRegression()
model.fit(df[predictors], df[target])
model.summary()
```

Python: scikit-learn

```
In[1]: from sklearn.linear_model import LinearRegression

path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

model = LinearRegression()
model.fit(df[predictors], df[target])
model.summary()
```

```
Out[1]: -----
AttributeError                                Traceback (most recent call last)
Cell In[52], line 13
     11 model = LinearRegression()
     12 model.fit(df[predictors], df[target])
--> 13 model.summary()

AttributeError: 'LinearRegression' object has no attribute 'summary'
```

Python: scikit-learn

```
In[1]: from sklearn.linear_model import LinearRegression

path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

model = LinearRegression()
model.fit(df[predictors], df[target])

# Print model coefficients
print(f'Intercept: {model.intercept_}')
print(f'Coefficients: {model.coef_}')

# Print model residuals
predictions = model.predict(df[predictors])
residuals = df[target] - predictions
print(f'Residuals: {residuals.values[:5]}...')
```

```
Out[1]: Intercept: -14.53525048050604
Coefficients: [-3.29859089e-01  7.67843024e-03 -3.91355574e-04 -6.79461791e-03
               8.52732469e-02  7.53367180e-01]
Residuals: [2.91708096  0.92742531  2.46368456  0.46552549  1.71359255]...
```

APIs of scikit-learn objects

To have a uniform API, we try to have a common basic API for all the objects. In addition, to avoid the proliferation of framework code, we try to adopt simple conventions and limit to a minimum the number of methods an object must implement.

Elements of the scikit-learn API are described more definitively in the [Glossary of Common Terms and API Elements](#).

Different objects

The main objects in scikit-learn are (one class can implement multiple interfaces):

Estimator: The base object, implements a `fit` method to learn from data, either:

```
estimator = estimator.fit(data, targets)
```

or:

```
estimator = estimator.fit(data)
```

Predictor: For supervised learning, or some unsupervised problems, implements:

```
prediction = predictor.predict(data)
```

Classification algorithms usually also offer a way to quantify certainty of a prediction, either using `decision_function` or `predict_proba`:

```
probability = predictor.predict_proba(data)
```

Transformer: For filtering or modifying the data, in a supervised or unsupervised way, implements:

```
new_data = transformer.transform(data)
```

When fitting and transforming can be performed much more efficiently together than separately, implements:

```
new_data = transformer.fit_transform(data)
```

Model: A model that can give a [goodness of fit](#) measure or a likelihood of unseen data, implements (higher is better):

```
score = model.score(data)
```

<https://scikit-learn.org/stable/developers/develop.html>

Python: scikit-learn

```
In[1]: from sklearn.linear_model import LinearRegression

path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

model = LinearRegression()
model.fit(df[predictors], df[target])

# Print model coefficients
print(f'Intercept: {model.intercept_}')
print(f'Coefficients: {model.coef_}')

# Print model residuals
predictions = model.predict(df[predictors])
residuals = df[target] - predictions
print(f'Residuals: {residuals.values[:5]}...')
```

```
Out[1]: Intercept: -14.53525048050604
Coefficients: [-3.29859089e-01  7.67843024e-03 -3.91355574e-04 -6.79461791e-03
               8.52732469e-02  7.53367180e-01]
Residuals: [2.91708096  0.92742531  2.46368456  0.46552549  1.71359255]...
```


Python: scikit-learn

```
In[1]: from sklearn.svm import SVR

path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

model = SVR(kernel='linear')
model.fit(df[predictors], df[target])

# Print model coefficients
print(f'Intercept: {model.intercept_}')
print(f'Coefficients: {model.coef_}')

# Print model residuals
predictions = model.predict(df[predictors])
residuals = df[target] - predictions
print(f'Residuals: {residuals.values[:5]}...')
```

```
Out[1]: Intercept: [-35.38646279]
Coefficients: [[-1.0526357  0.05910105 -0.03667206 -0.00831565  0.56218046
  0.96851648]]
Residuals: [3.0266171  0.62154228 3.10666275 1.34695011 3.07475274]...
```

Coding tips: Separation of concerns

```
In[1]: # Read and clean data
path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

# Split data
train = df.iloc[:int(len(df) * 0.8)]
validation = df.iloc[int(len(df) * 0.8):]

# Define input and output variables
predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

# Define necessary data structures for state
chosen_predictors = []
mses = []

while len(predictors) > 0:
    best_predictor = {'mse': float('inf'), 'predictor': None}

    for predictor in set(predictors) - set(chosen_predictors):
        potential_predictors = chosen_predictors + [predictor]

        # Fit and evaluate model
        model = LinearRegression()
        model.fit(train[potential_predictors], train[target])
        predictions = model.predict(validation[potential_predictors])
        test_mse = np.mean((validation[target] - predictions) ** 2)

        # Compare model with previous best
        if test_mse < best_predictor['mse']:
            best_predictor = {'mse': test_mse, 'predictor': predictor}

    # Update state
    chosen_predictors.append(best_predictor['predictor'])
    mses.append(best_predictor['mse'])
    predictors = [p for p in predictors if p != best_predictor['predictor']]
```

Coding tips: Separation of concerns

```
In[1]: # Read and clean data
path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

# Split data
train = df.iloc[:int(len(df) * 0.8)]
validation = df.iloc[int(len(df) * 0.8):]

# Define input and output variables
predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

# Define necessary data structures for state
chosen_predictors = []
mses = []

while len(predictors) > 0:
    best_predictor = {'mse': float('inf'), 'predictor': None}

    for predictor in set(predictors) - set(chosen_predictors):
        potential_predictors = chosen_predictors + [predictor]

        # Fit and evaluate model
        model = LinearRegression()
        model.fit(train[potential_predictors], train[target])
        predictions = model.predict(validation[potential_predictors])
        test_mse = np.mean((validation[target] - predictions) ** 2)

        # Compare model with previous best
        if test_mse < best_predictor['mse']:
            best_predictor = {'mse': test_mse, 'predictor': predictor}

    # Update state
    chosen_predictors.append(best_predictor['predictor'])
    mses.append(best_predictor['mse'])
    predictors = [p for p in predictors if p != best_predictor['predictor']]
```

Setup

Selection

Modelling

Housekeeping

Coding tips: Separation of concerns

```
In[1]: # Read and clean data
path = '/Users/esten/Downloads/Auto.csv'
df = pd.read_csv(path)

# Split data
train = df.iloc[:int(len(df) * 0.8)]
validation = df.iloc[int(len(df) * 0.8):]

# Define input and output variables
predictors = ['cylinders', 'displacement', 'horsepower',
              'weight', 'acceleration', 'year']
target = 'mpg'

# Define necessary data structures for state
chosen_predictors = []
mses = []

def fit_and_evaluate_model(model: LinearRegression, train: pd.DataFrame,
                           validation: pd.DataFrame, variables: List[str],
                           target: str):
    """ Fit a given model on a training dataset using a given set of variables
    and return MSE from a validation dataset. """
    model = LinearRegression()
    model.fit(train[potential_predictors], train[target])
    predictions = model.predict(validation[potential_predictors])

    return np.mean((validation[target] - predictions) ** 2)

while len(predictors) > 0:
    best_predictor = {'mse': float('inf'), 'predictor': None}

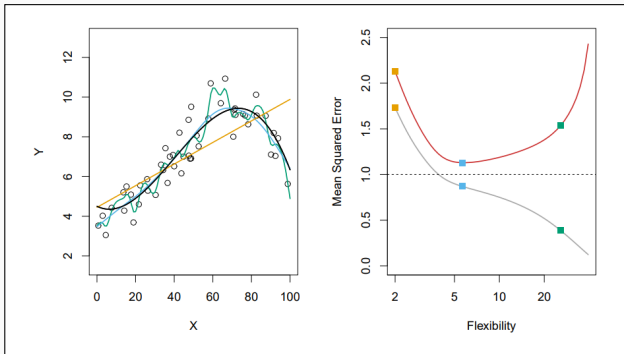
    for predictor in set(predictors) - set(chosen_predictors):
        potential_predictors = chosen_predictors + [predictor]
        test_mse = fit_and_evaluate_model(LinearRegression(), train, validation,
                                          variables=potential_predictors,
                                          target=target)

        # Compare model with previous best
        if test_mse < best_predictor['mse']:
            best_predictor = {'mse': test_mse, 'predictor': predictor}

    # Update state
    chosen_predictors.append(best_predictor['predictor'])
    mses.append(best_predictor['mse'])
    predictors = [p for p in predictors if p != best_predictor['predictor']]
```

Modelling

Regularization: Motivation



Regularization: Out-of-sample testing

Regularization: Methods

1. Variable selection
 - a. Best subset selection
 - b. Forward stepwise selection
 - c. Backward stepwise selection
2. Shrinkage
 - a. LASSO
 - b. Ridge Regression
 - c. Elastic net
3. Dimensionality reduction
 - a. Principal Component Regression
 - b. Partial Least Squares

Variable selection: Outline

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Variable selection: Outline

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Motivation

1. Simplify interpretation
2. Reduce model complexity (overfitting)

Variable selection: Outline

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Variable selection: Best subset selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution (best subset selection)

Train models on all subsets p and select the best one.

Variable selection: Best subset selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution (best subset selection)

Train models on all subsets p and select the best one.

Variable selection: Best subset selection