

PSY9511: Seminar 3

Regularization and variable selection

Esten H. Leonardsen

07.09.23

1. Assignment 1
2. Assignment 2
3. Regularization
 - Variable selection
 - Shrinkage (+ live coding 🤖)
 - Dimensionality reduction

Assignment 1



UNIVERSITETET
I OSLO

- Create a vector of 100 standard normally distributed numbers and visualize them with a histogram.
- Show rows 5, 8, 9, and 10 of the Auto dataset.
- Show the last three columns of the Auto dataset.
- Show all cars with five cylinders in the Auto dataset.

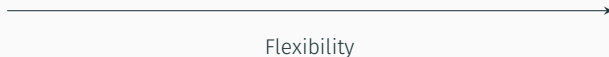
- Create a vector of 100 standard normally distributed numbers and visualize them with a histogram.
- Show rows 5, 8, 9, and 10 of the Auto dataset.
- Show the last three columns of the Auto dataset.
- Show all cars with five cylinders in the Auto dataset.

<http://localhost:8889/notebooks/notebooks%2FAssignment%201.ipynb>

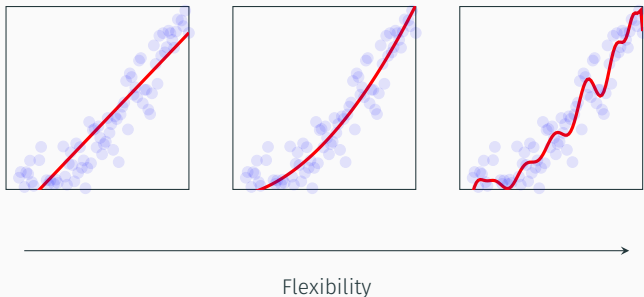


Assignment 2: Bias-variance trade-off

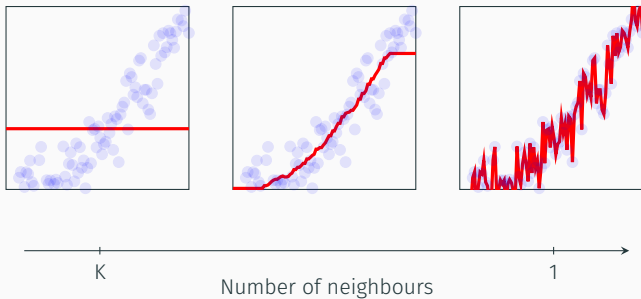




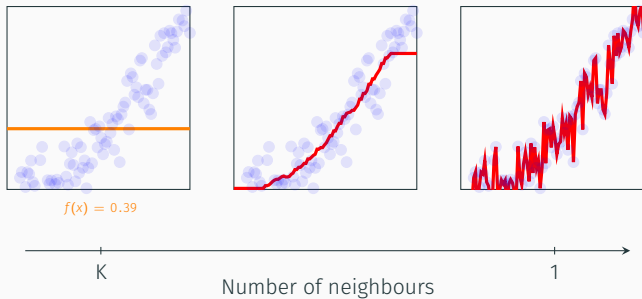
Assignment 2: Bias-variance trade-off



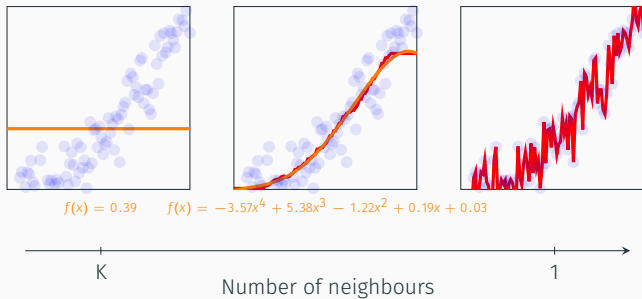
Assignment 2: Bias-variance trade-off



Assignment 2: Bias-variance trade-off



Assignment 2: Bias-variance trade-off



Assignment 2: Bias-variance trade-off

$$f(x) = 0.39 \quad f(x) = -3.57x^4 + 5.38x^3 - 1.22x^2 + 0.19x + 0.03$$



Assignment 2: Bias-variance trade-off

$$f(x) = 0.39 \quad f(x) = -3.57x^4 + 5.38x^3 - 1.22x^2 + 0.19x + 0.03$$

↑
1



Assignment 2: Bias-variance trade-off

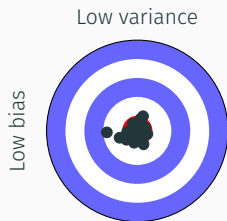
$$\begin{array}{cccccc} f(x) = 0.39 & f(x) = -3.57x^4 + 5.38x^3 - 1.22x^2 + 0.19x + 0.03 \\ \uparrow & \uparrow & \uparrow & \uparrow & \uparrow & \uparrow \\ 1 & 1 & 2 & 3 & 4 & 5 \end{array}$$

Model flexibility: Denotes the complexity of the approximated function $\hat{y} = \hat{f}(x)$.

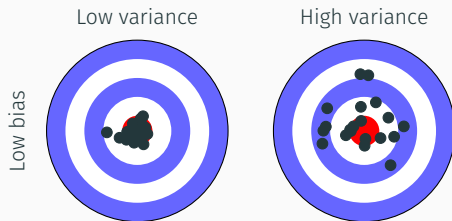
- Informally: Wigglyness of the line
- Formally: Number of parameters in the function (degrees of freedom)



Assignment 2: Bias-variance trade-off



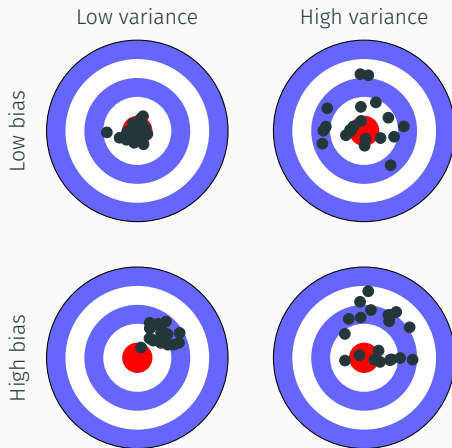
Assignment 2: Bias-variance trade-off



Assignment 2: Bias-variance trade-off



Assignment 2: Bias-variance trade-off



Assignment 2



UNIVERSITETET
I OSLO

Assignment 2: Data splitting



Assignment 2: Random seeds



Assignment 2: Log-odds vs probability vs class



Assignment 2: Eye test



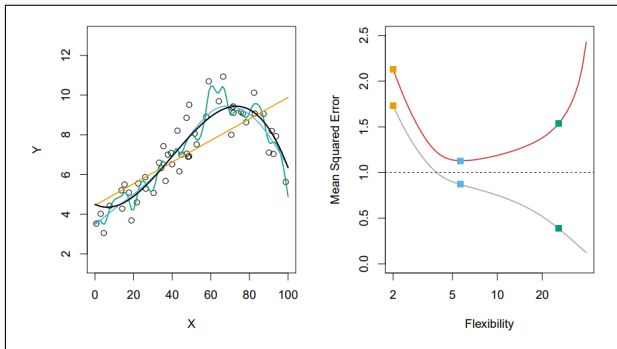
Regularization



UNIVERSITETET
I OSLO

$$y \sim \beta_0 + \beta_1 * x_1 + \beta_2 * x_2 + \beta_3 * x_3$$

Regularization: Motivation



$$y \sim \beta_0 + \beta_1 * x_1 + \beta_2 * x_2 + \beta_3 * x_3$$

Regularization: Out-of-sample testing

```
In[1]: import pandas as pd

df = pd.read_csv('/Users/esten/Downloads/Auto.csv')
train = df.iloc[:int(len(df) * 0.8)]
validation = df.iloc[int(len(df) * 0.8):]

print(f'Using {len(train)} samples for training')
print(f'Using {len(validation)} samples for validation')
```

```
Out[1]: Using 317 samples for training
Using 80 samples for validation
```



1. Variable selection
 - a. Best subset selection
 - b. Forward stepwise selection
 - c. Backward stepwise selection
2. Shrinkage
 - a. LASSO
 - b. Ridge Regression
3. Dimensionality reduction
 - a. Principal Component Regression
 - b. Partial Least Squares



Variable selection



UNIVERSITETET
I OSLO

The number of predictors we are using in our model directly impacts model complexity.

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Problem

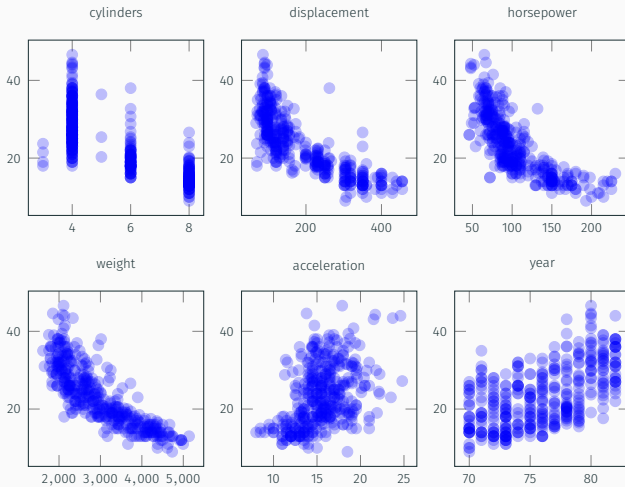
We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Motivation

1. Reduce model complexity (overfitting)
2. Simplify interpretation

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .



Variable selection: Best subset selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Train models on all subsets p and select the best one.



Variable selection: Best subset selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

```
In[1]: import numpy as np

from itertools import chain, combinations
from sklearn.linear_model import LinearRegression

subsets = list(chain.from_iterable(combinations(predictors, r) \
                                   for r in range(len(predictors)+1)))

best = {'mse': float('inf'), 'subset': None}

for subset in subsets:
    if len(subset) == 0:
        continue

    model = LinearRegression()
    model.fit(train[list(subset)], train[target])
    predictions = model.predict(validation[list(subset)])
    mse = np.mean((predictions - validation[target]) ** 2)

    if mse < best['mse']:
        best = {'mse': mse, 'subset': subset}

print(f'MSE: {best["mse"]:.2f}, predictors: {best["subset"]}')
```

```
Out[1]: MSE: 29.68, predictors: ('cylinders', 'displacement', 'horsepower', 'weight', 'year')
```



Variable selection: Best subset selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Train models on all subsets p and select the best one.

+ Positives

Guaranteed to find the optimal solution.

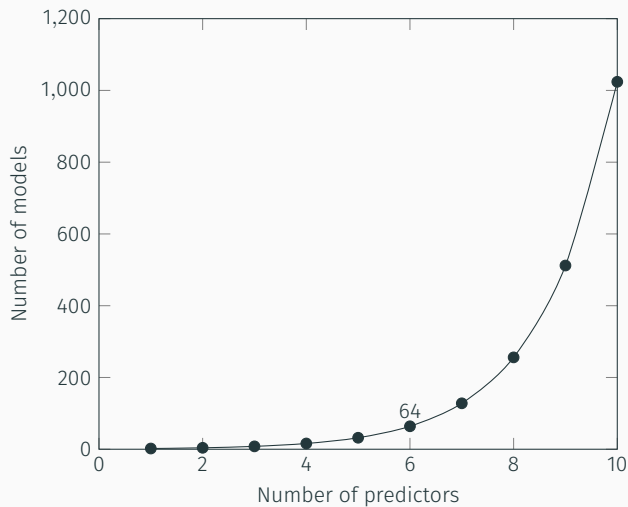
Simple implementation

- Drawbacks

Need to train many ($2^{|P|}$) models.



Variable selection: Best subset selection



Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.

$y \sim 1$ $mse = 146.47$



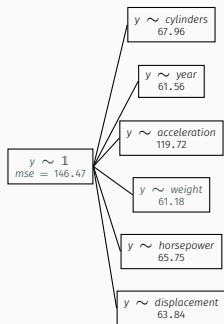
Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



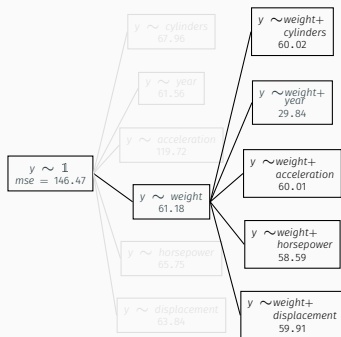
Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



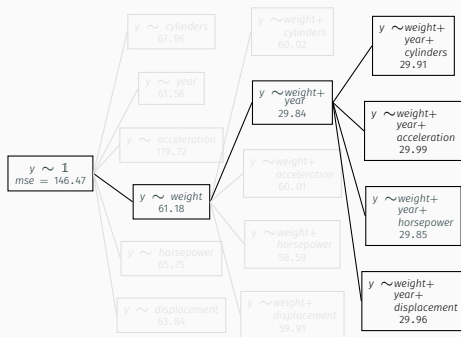
Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



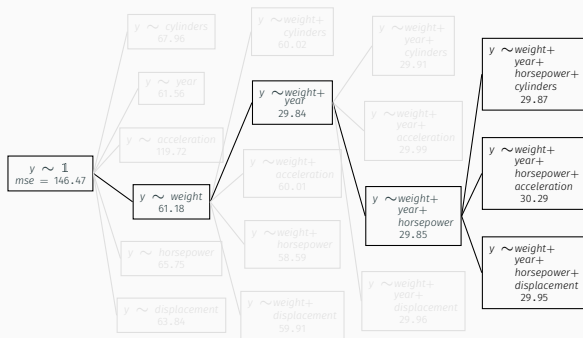
Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



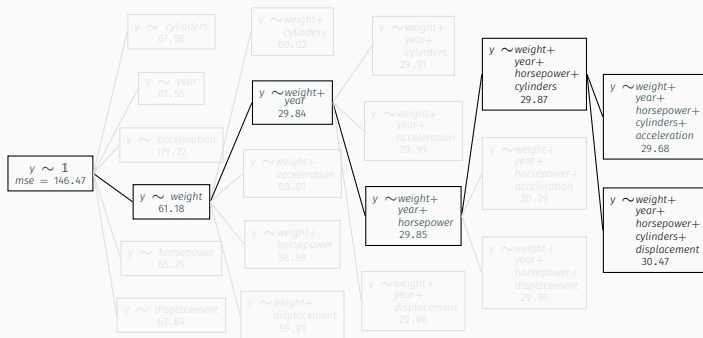
Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



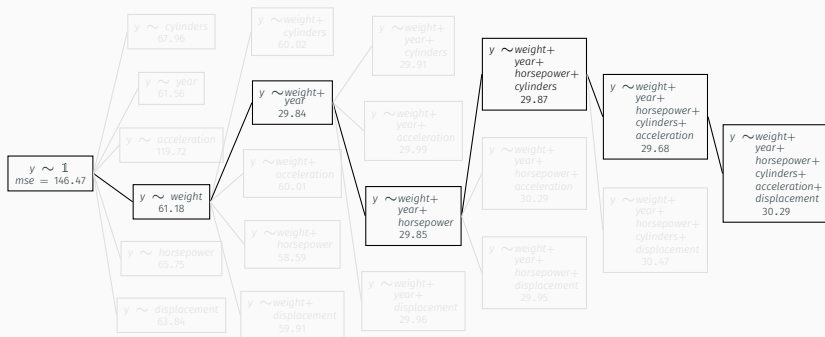
Variable selection: Forward stepwise selection

Problem

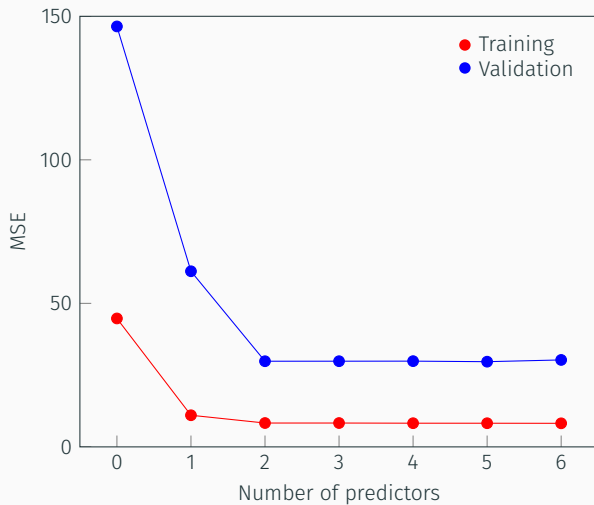
We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

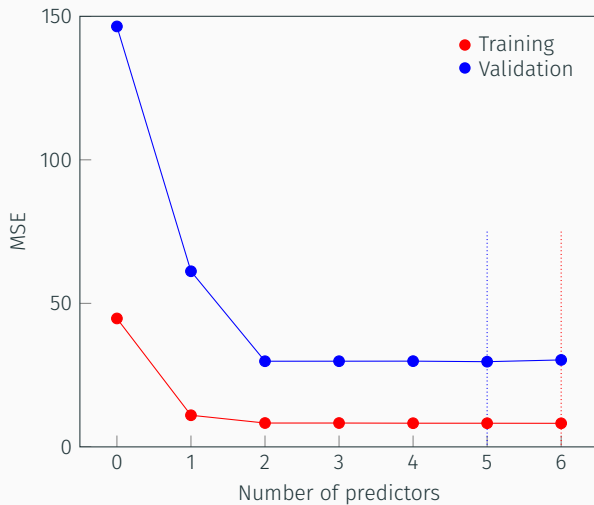
Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



Variable selection: Forward stepwise selection



Variable selection: Forward stepwise selection



Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

```
In[1]: def fit_and_evaluate(train: pd.DataFrame, validation: pd.DataFrame,
    predictors: List[str], target: str):
    model = LinearRegression()
    model.fit(train[predictors], train[target])

    train_predictions = model.predict(train[predictors])
    validation_predictions = model.predict(validation[predictors])

    return np.mean((train_predictions - train[target]) ** 2), \
           np.mean((validation_predictions - validation[target]) ** 2)

predictors = ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year']
target = 'mpg'

train['intercept'] = 1
validation['intercept'] = 1
train_mse, validation_mse = fit_and_evaluate(train, validation,
                                           predictors=['intercept'],
                                           target=target)
print(f'[]: {validation_mse:.2f} ({train_mse:.2f})')

chosen_predictors = []

while len(chosen_predictors) < len(predictors):
    best_predictor = {'train_mse': None, 'validation_mse': float('inf'),
                     'predictor': None}

    for predictor in set(predictors) - set(chosen_predictors):
        train_mse, validation_mse = fit_and_evaluate(train, validation,
                                                    predictors=chosen_predictors + [predictor],
                                                    target=target)

        if validation_mse < best_predictor['validation_mse']:
            best_predictor = {'train_mse': train_mse, 'validation_mse': validation_mse, 'predictor': predictor}

    chosen_predictors.append(best_predictor['predictor'])
```



Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.

+ Positives

Need to train fewer models.

- Drawbacks

Not guaranteed to find the optimal solution.



Variable selection: Backward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with all predictors. Iteratively remove the predictor that yields the best model until all you have none left.



Variable selection: Backward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with all predictors. Iteratively remove the predictor that yields the best model until all you have none left.

+ Positives

Need to train fewer models.

- Drawbacks

Not guaranteed to find the optimal solution.



Shrinkage



UNIVERSITETET
I OSLO

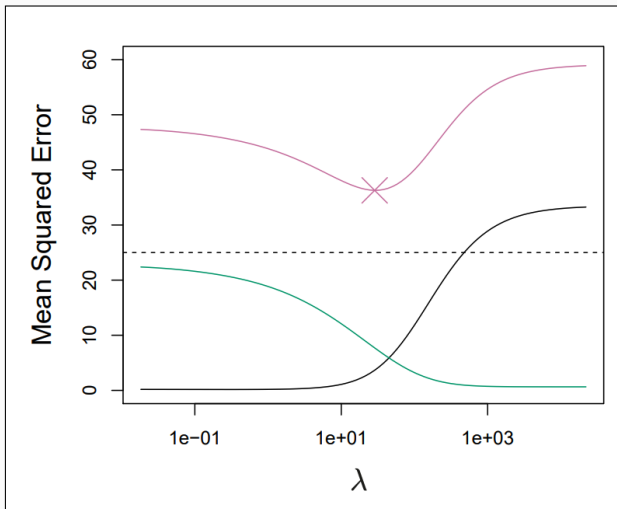
$$y \sim \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5 + \beta_6 x_6$$


```
Out[1]:
```

	coef	std err	P> t	[0.025	0.975]
Intercept	-14.5353	4.764	0.002	-23.90	-5.16
cylinders	-0.3299	0.332	0.321	-0.98	0.32
displacement	0.0077	0.007	0.297	-0.00	0.02
horsepower	-0.0004	0.014	0.977	-0.02	0.02
weight	-0.0068	0.001	0.000	-0.00	-0.00
acceleration	0.0853	0.102	0.404	-0.11	0.28
year	0.7534	0.053	0.000	0.65	0.85

$$y \sim \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5 + \beta_6 x_6$$

$$mse = bias^2 + variance + irreducible\ error$$



$$\text{salary} \sim \beta_0 + \beta_1 * \text{age}$$

$$\text{salary} \sim \beta_0 + \beta_1 * \text{age}$$

$$\text{salary} \sim 300000 + 10000 * \text{age}$$

$$\text{salary} \sim 600000 + 0 * \text{age}$$

$$loss_{mse} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2$$

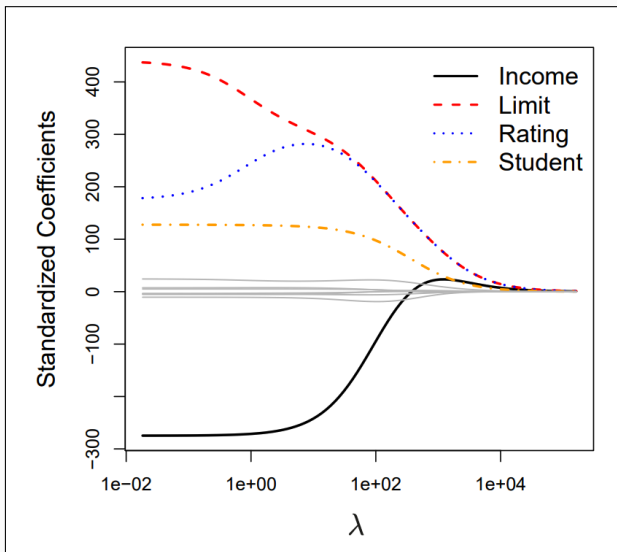
$$\text{loss}_{\text{ridge}} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p \beta_j^2$$

Shrinkage: Ridge regression

$$\text{loss}_{\text{ridge}} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p \beta_j^2$$

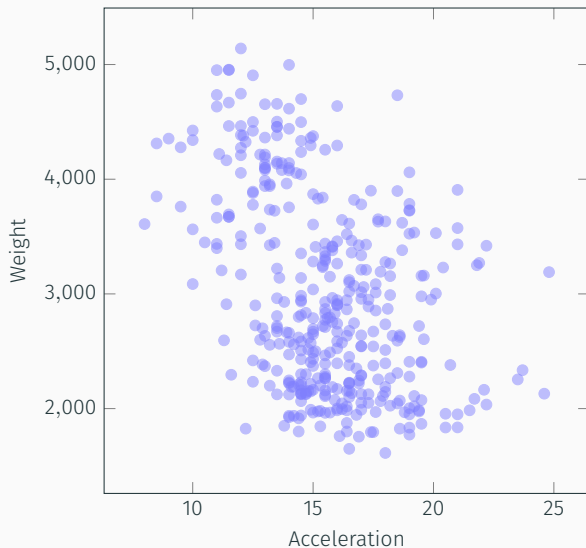
$$\Downarrow \\ \lambda \rightarrow \infty \Rightarrow \beta \rightarrow 0$$

Shrinkage



$$loss_{ridge} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p \beta_j^2$$

Shrinkage: Feature standardization



z-score standardization

z-score standardization

$$x = \frac{x - \mu_x}{\sigma_x^2}$$

z-score standardization

$$X = \frac{x - \mu_X}{\sigma_X^2}$$

```
In[1]: for col in predictors:
        print(f'{col}: {np.mean(df[col]):.2f} ({np.std(df[col]):.2f})')

        # z-score standardization
        for col in predictors:
            df[col] = (df[col] - np.mean(df[col])) / np.std(df[col])

        for col in predictors:
            print(f'{col} after: {np.mean(df[col]):.2f} ({np.std(df[col]):.2f})')
```

z-score standardization

$$x = \frac{x - \mu_x}{\sigma_x^2}$$

```
In[1]: for col in predictors:
        print(f'{col}: {np.mean(df[col]):.2f} ({np.std(df[col]):.2f})')

        # z-score standardization
        for col in predictors:
            df[col] = (df[col] - np.mean(df[col])) / np.std(df[col])

        for col in predictors:
            print(f'{col} after: {np.mean(df[col]):.2f} ({np.std(df[col]):.2f})')
```

```
Out[1]: cylinders: 5.47 (1.70)
displacement: 194.41 (104.51)
horsepower: 104.47 (38.44)
weight: 2977.58 (848.32)
acceleration: 15.54 (2.76)
year: 75.98 (3.68)
cylinders after: -0.00 (1.00)
displacement after: -0.00 (1.00)
horsepower after: -0.00 (1.00)
weight after: -0.00 (1.00)
acceleration after: 0.00 (1.00)
year after: -0.00 (1.00)
```



<http://localhost:8888/notebooks/notebooks/Live%20coding.ipynb>

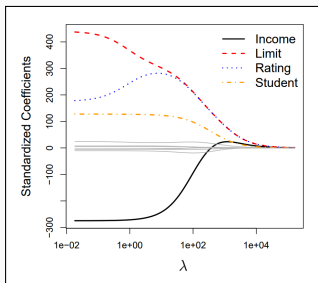
$$loss_{ridge} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p \beta_j^2$$

Regularization through shrinking the model covariates towards zero.

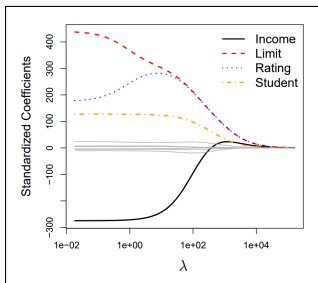
$$loss_{ridge} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p \beta_j^2$$

$$loss_{lasso} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p |\beta_j|$$

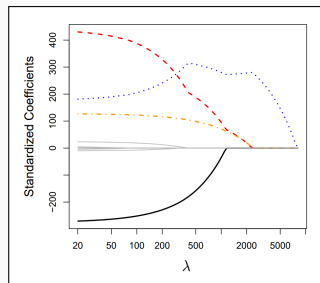
Ridge



Ridge



LASSO

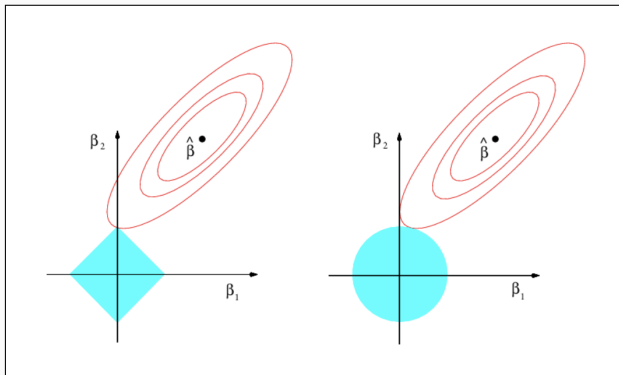


Predictor	Ridge	LASSO
Intercept	23.44	23.44
Weight	-5.59	-4.78
Year	2.75	2.00
Horsepower	-0.07	-0.09
Cylinders	-0.54	0
Acceleration	0.19	0
Displacement	0.66	0

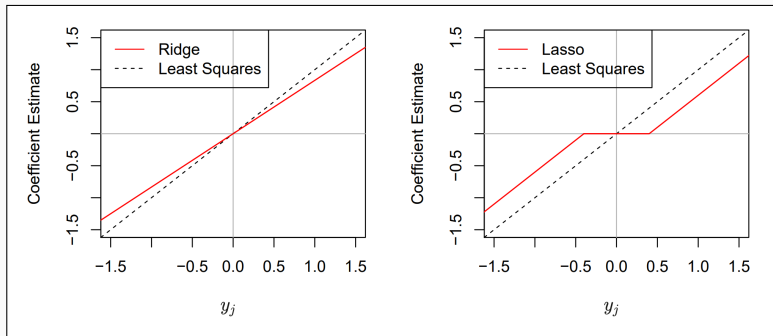
Predictor	Ridge	LASSO
Intercept	23.44	23.44
Weight	-5.59	-4.78
Year	2.75	2.00
Horsepower	-0.07	-0.09
Cylinders	-0.54	0
Acceleration	0.19	0
Displacement	0.66	0

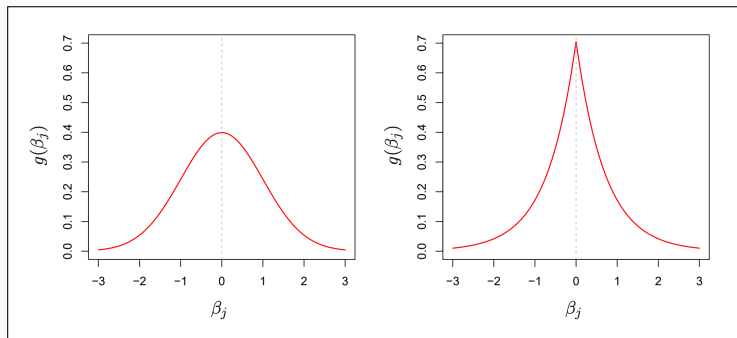
A coefficient of 0 does not mean the predictor has no predictive value for the outcome!

Shrinkage: LASSO



Whiteboard! 🎉





$$loss_{mse} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2$$

Fits the **best** model
to the data.

$$loss_{mse} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2$$

Fits the **best** model to the data.

$$loss_{ridge} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p \beta_j^2$$

Fits the **best** model to the data while **shrinking** coefficients towards zero.

$$loss_{mse} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2$$

Fits the **best** model to the data.

$$loss_{ridge} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p \beta_j^2$$

Fits the **best** model to the data while **shrinking** coefficients towards zero.

$$loss_{lasso} = \sum_{i=0}^n \left(y_i - \sum_{j=0}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=0}^p |\beta_j|$$

Fits the **best** model to the data while **shrinking** coefficients towards zero such that some variables are effectively **removed**.