

# PSY9511: Seminar 7

Deep learning for computer vision tasks

---

Esten H. Leonardsen

07.11.24



UNIVERSITY  
OF OSLO

# Outline

1. Exercise 4
2. Deep learning with artificial neural networks
3. Convolutional neural networks for computer vision



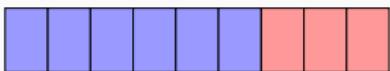
# Weekly exercises

- The weekly exercises are **mandatory**
- The deadlines are **strict**

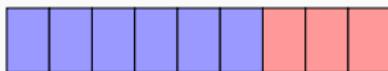
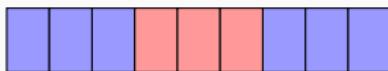
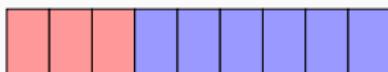


# Validation procedures

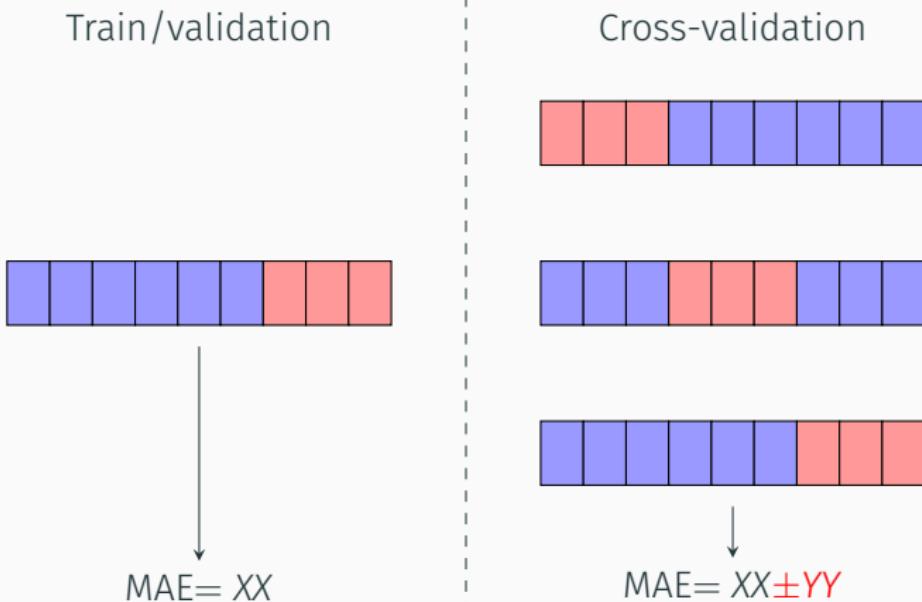
Train/validation



Cross-validation



# Validation procedures



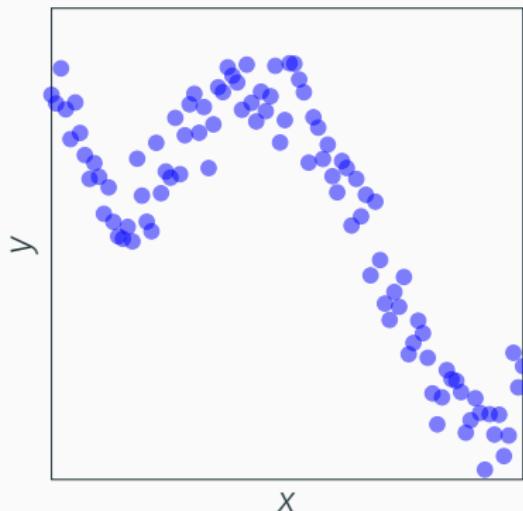
# Deep learning

---

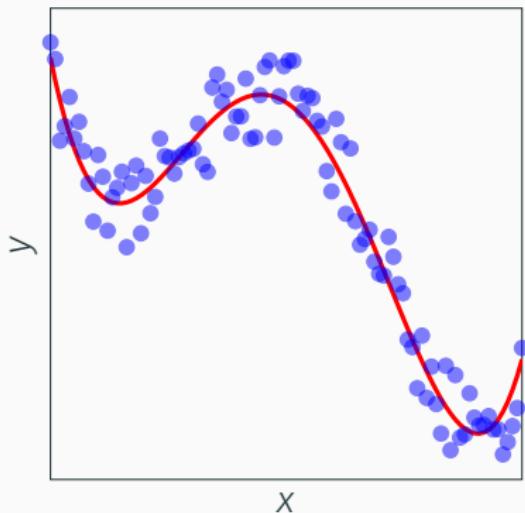


UNIVERSITY  
OF OSLO

# Deep learning: Motivation



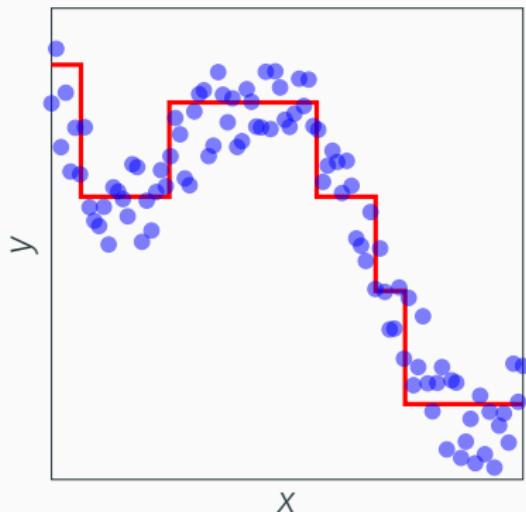
# Deep learning: Motivation



$$\hat{y} = s(x)$$



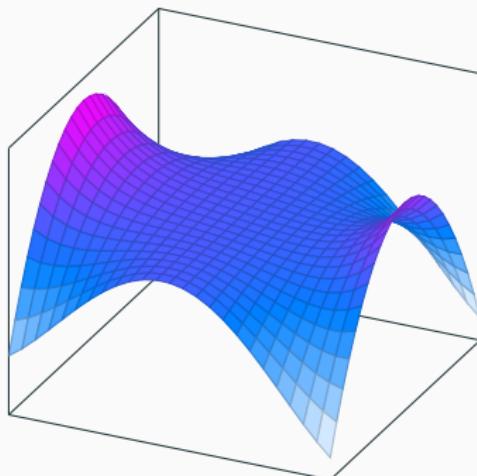
# Deep learning: Motivation



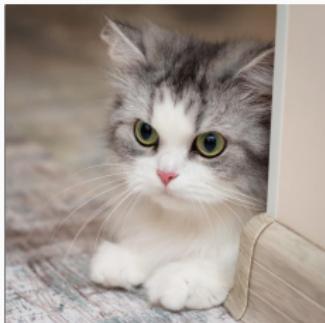
$$\hat{y} = \begin{cases} 4 & \dots \\ 3 & 0.2 \leq x < 0.6 \\ 1.5 & \dots \end{cases}$$



# Deep learning: Motivation

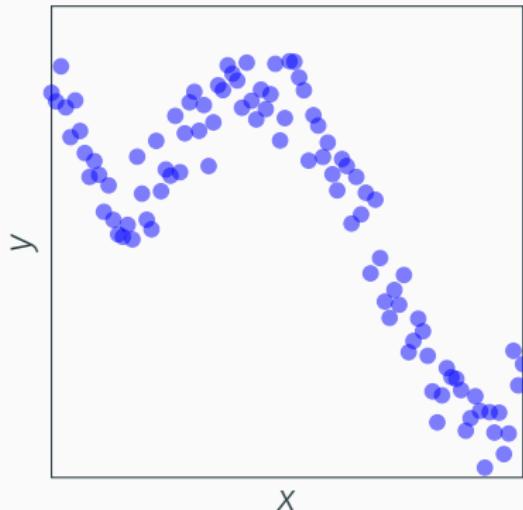


# Deep learning: Motivation

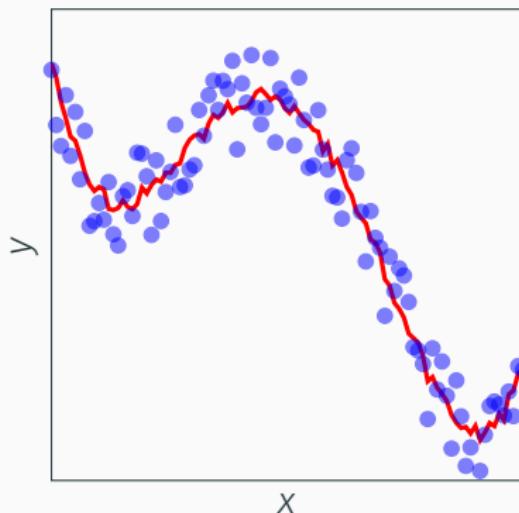


The cat wagged  
its tail

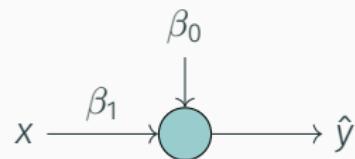
# Deep learning: Motivation



# Deep learning: Motivation



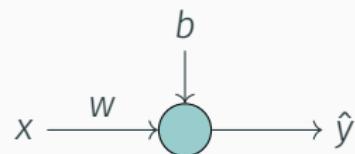
# Deep learning: Artificial neural networks



$$\hat{y} = \beta_0 + \beta_1 x$$



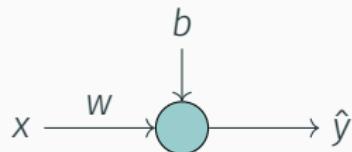
# Deep learning: Artificial neural networks



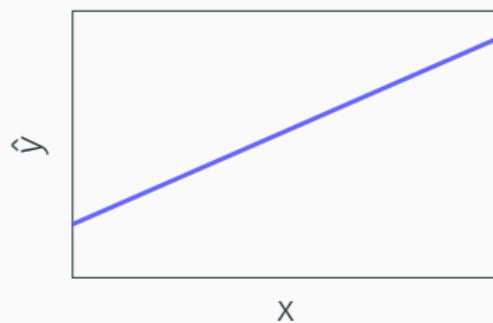
$$\hat{y} = wx + b$$



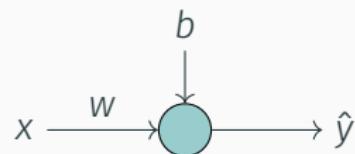
# Deep learning: Artificial neural networks



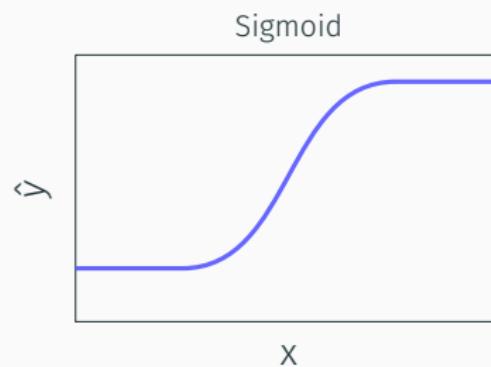
$$\hat{y} = wx + b$$



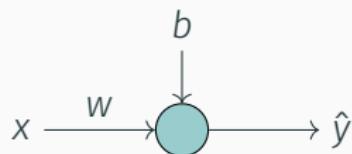
# Deep learning: Artificial neural networks



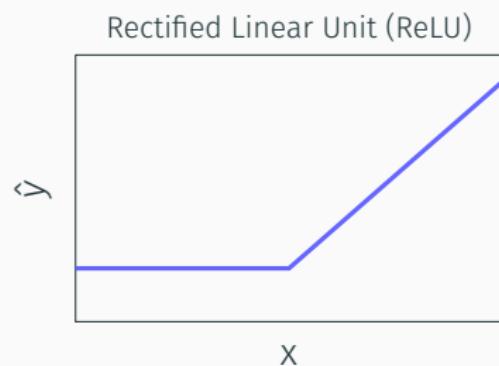
$$\hat{y} = \frac{e^{wx+b}}{1 + e^{wx+b}}$$



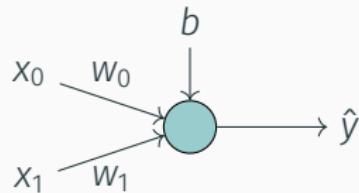
# Deep learning: Artificial neural networks



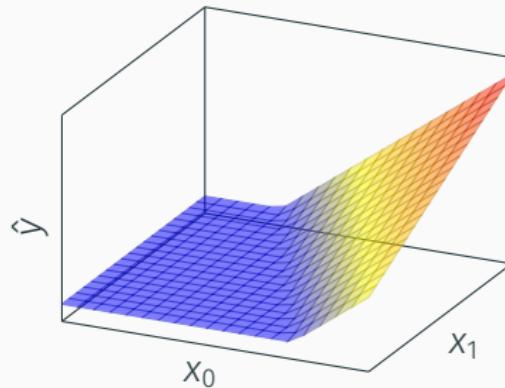
$$\hat{y} = \max(0, wx + b)$$



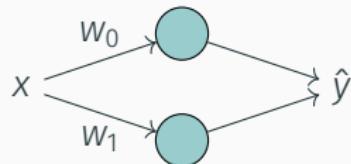
# Deep learning: Artificial neural networks



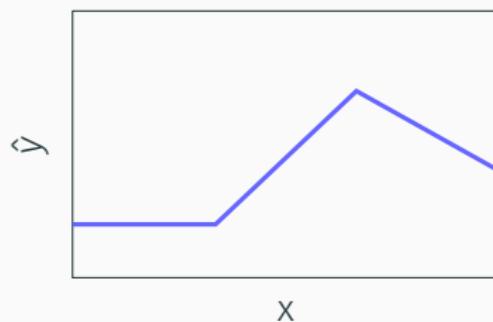
$$\hat{y} = \max(0, w_0x_0 + w_1x_1 + b)$$



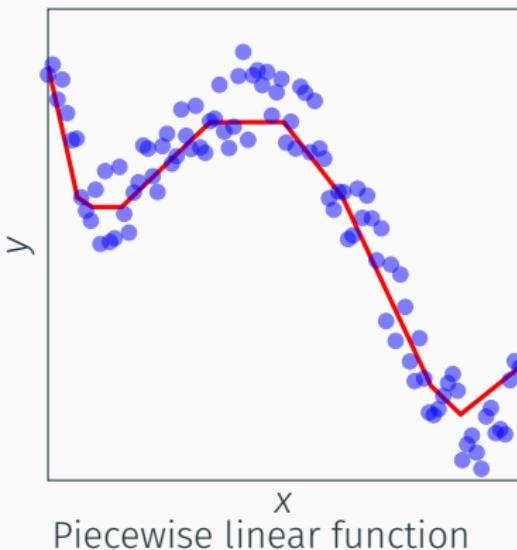
# Deep learning: Artificial neural networks



$$\hat{y} = \max(0, w_0x + b_0) + \max(0, w_1x + b_1)$$



# Deep learning: Artificial neural networks



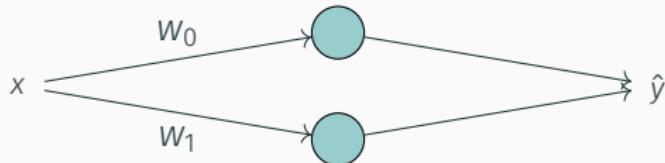
# Deep learning: Artificial neural networks

**Universal approximation theorem:**

*"Any relationship that can be described with a polynomial function can be approximated by a neural network with a single hidden layer."*



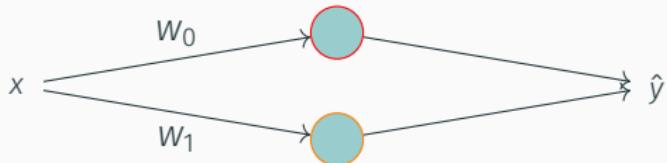
# Deep learning: Artificial neural networks



$$\hat{y} = \max(0, w_0x + b_0) + \max(0, w_1x + b_1)$$



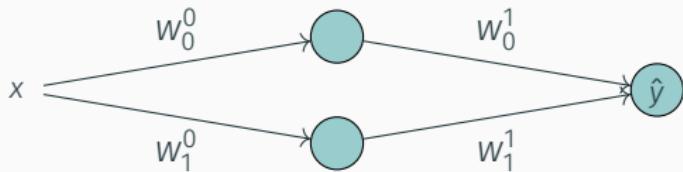
# Deep learning: Artificial neural networks



$$\hat{y} = \boxed{\max(0, w_0x + b_0)} + \boxed{\max(0, w_1x + b_1)}$$



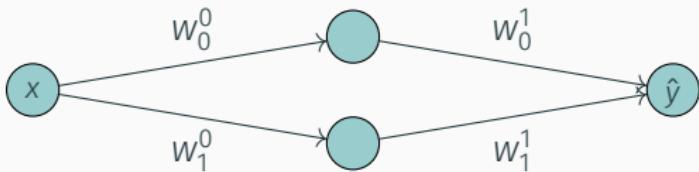
# Deep learning: Artificial neural networks



$$\hat{y} = \max(0, w_0^1 * \max(0, w_0^0 * x + b_0^0) + w_1^1 * \max(0, w_1^0 * x + b_1^0) + b_0^1)$$



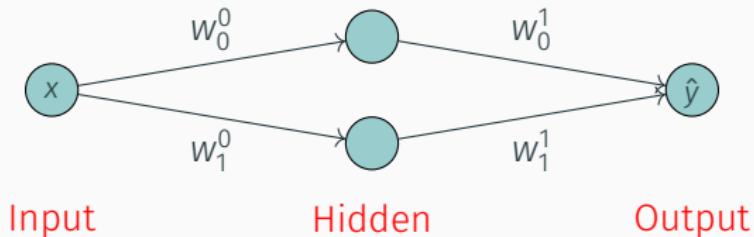
# Deep learning: Artificial neural networks



$$\hat{y} = \max(0, w_0^1 * \max(0, w_0^0 * x + b_0^0) + w_1^1 * \max(0, w_1^0 * x + b_1^0) + b_0^1)$$



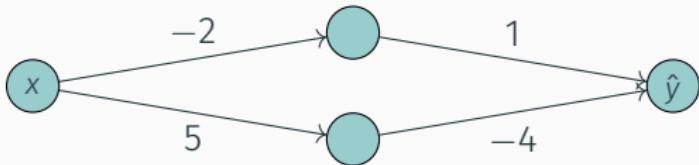
# Deep learning: Artificial neural networks



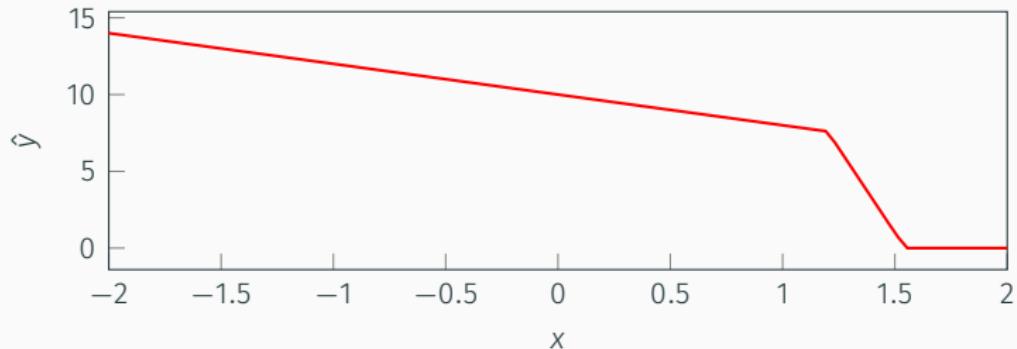
$$\hat{y} = \max(0, w_0^1 * \max(0, w_0^0 * x + b_0^0) + w_1^1 * \max(0, w_1^0 * x + b_1^0) + b_0^1)$$



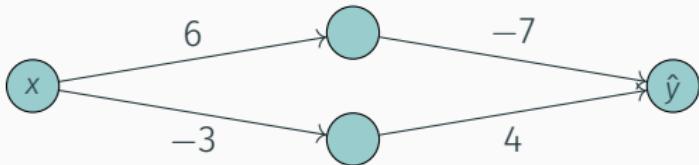
# Deep learning: Artificial neural networks



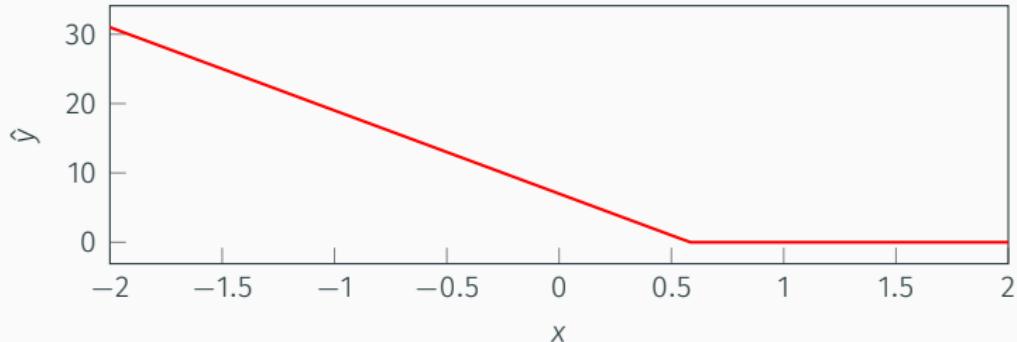
$$\hat{y} = \max(0, 1 * \max(0, (-2) * x + 3) + (-4) * \max(0, 5 * x + (-6)) + 7)$$



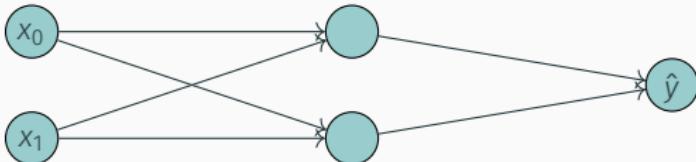
# Deep learning: Artificial neural networks



$$\hat{y} = \max(0, -7 * \max(0, 6 * x + (-5)) + 4 * \max(0, (-3) * x + 2) - 1)$$



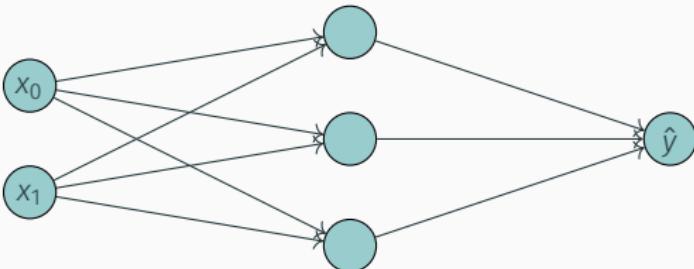
# Deep learning: Artificial neural networks



$$\begin{aligned}\hat{y} = & \max(0, w_{0,0}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\ & w_{1,0}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\ & b_1)\end{aligned}$$



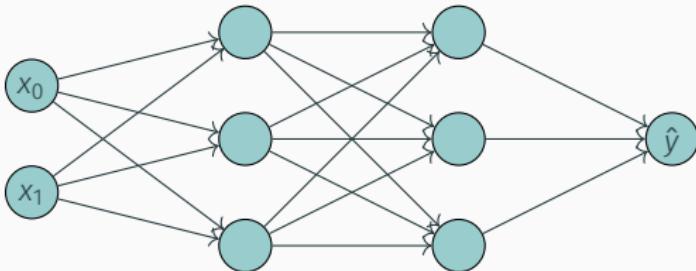
# Deep learning: Artificial neural networks



$$\begin{aligned}\hat{y} = & \max(0, w_{0,0}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\& w_{1,0}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\& w_{2,0}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\& b_1)\end{aligned}$$



# Deep learning: Artificial neural networks



$$\begin{aligned}\hat{y} = & \max(0, w_{0,0}^2 * \max(0, w_{0,0}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\ & w_{1,0}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\ & w_{2,0}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\ & b_{1,0}) + \\ & w_{1,0}^2 * \max(0, w_{0,1}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\ & w_{1,1}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\ & w_{2,1}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\ & b_{1,1}) + \\ & w_{2,0}^2 * \max(0, w_{0,2}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\ & w_{1,2}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\ & w_{2,2}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\ & b_{1,2}) + \\ & b_2)\end{aligned}$$



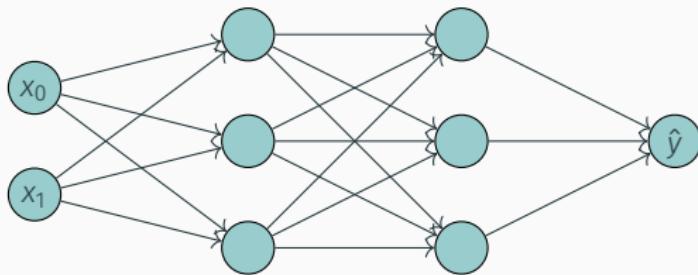
# Deep learning: Artificial neural networks

Artificial neural networks: Combines artificial neurons, simple computational units that compute a non-linear function of their inputs, in a computational graph

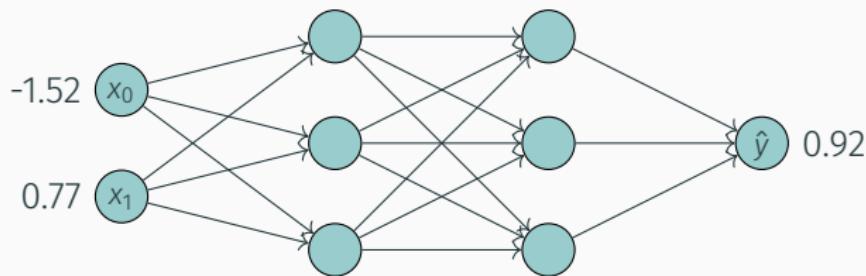
- Can approximate arbitrarily complex polynomial functions (given enough neurons)
- Organized in layers. We can expand a model in width (e.g. more neurons per layer) or depth (e.g. more layers)



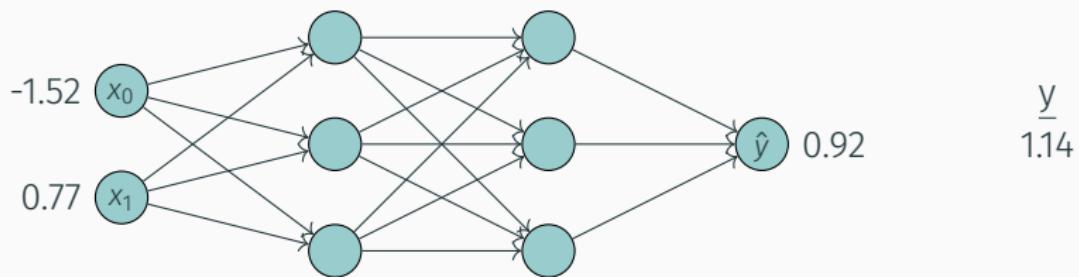
# Deep learning: Loss functions



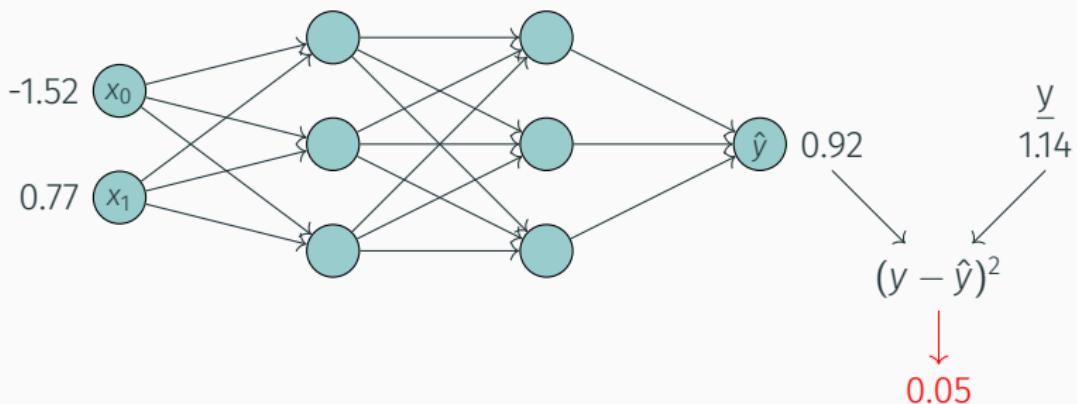
# Deep learning: Loss functions



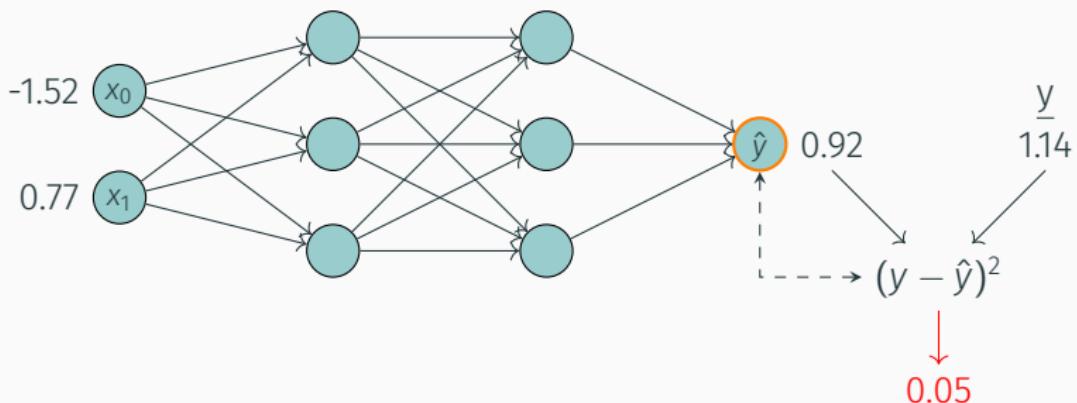
# Deep learning: Loss functions



# Deep learning: Loss functions



# Deep learning: Loss functions



# Deep learning: Loss functions

```
In[1]: from tensorflow.keras.layers import Input, Dense  
from tensorflow.keras import Model  
  
inputs = Input(shape=(2,))  
hidden1 = Dense(units=3, activation='relu')(inputs)  
hidden2 = Dense(units=3, activation='relu')(hidden1)  
outputs = Dense(units=1, activation=?)(hidden2)  
  
model = Model(inputs, outputs)  
model.compile(loss=?)
```



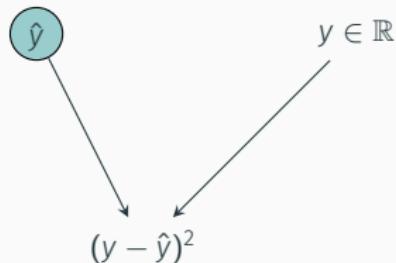
# Deep learning: Loss functions

```
In[1]: from tensorflow.keras.layers import Input, Dense  
from tensorflow.keras import Model  
  
inputs = Input(shape=(2,))  
hidden1 = Dense(units=3, activation='relu')(inputs)  
hidden2 = Dense(units=3, activation='relu')(hidden1)  
outputs = Dense(units=1, activation='?')(hidden2)  
  
model = Model(inputs, outputs)  
model.compile(loss=?)
```



# Deep learning: Loss functions

## Regression



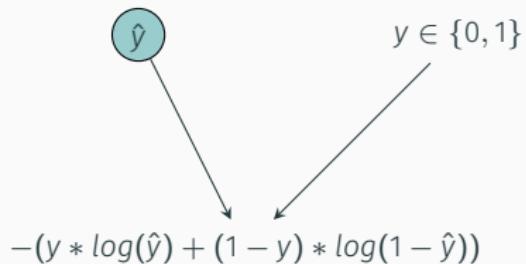
In[2]:

```
...
outputs = Dense(units=1, activation=None)(...)
...
model.compile(loss='mean_squared_error')
```



# Deep learning: Loss functions

## Binary classification



In[3]:

```
...  
outputs = Dense(units=1, activation='sigmoid')(...)  
...  
model.compile(loss='binary_crossentropy')
```



# Deep learning: Loss functions

## Multiclass classification

$\hat{y}$

$y \in \{cat, dog, bat\}$



# Deep learning: Loss functions

## Multiclass classification

$\hat{y}$

$x_0$	$y$
	cat
	dog
	bat
	dog



# Deep learning: Loss functions

## Multiclass classification

$\hat{y}$

$x_0$	cat	dog	bat
	1	0	1
	0	1	0
	0	0	1
	0	1	0



# Deep learning: Loss functions

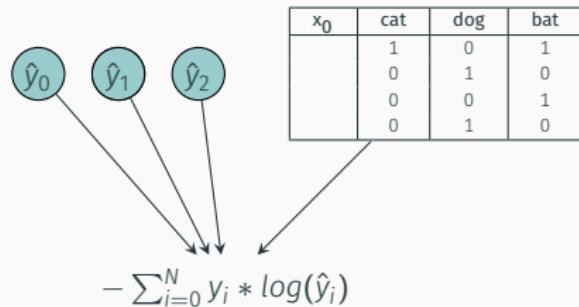
## Multiclass classification

$\hat{y}_0$     $\hat{y}_1$     $\hat{y}_2$

$x_0$	cat	dog	bat
	1	0	1
	0	1	0
	0	0	1
	0	1	0



# Deep learning: Loss functions



In[3]:

```
...
outputs = Dense(units=1, activation='softmax')(...)
...
model.compile(loss='categorical_crossentropy')
```



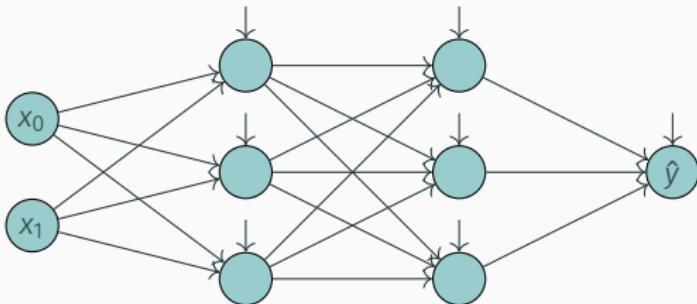
# Deep learning: Loss functions

Loss functions: Behaves for neural networks as any other statistical learning model. However, important to **configure the final layer correctly**

- Regression: Mean squared error
  - No activation
- Binary classification: Binary cross-entropy
  - Sigmoid activation
- Multiclass classification: Categorical cross-entropy
  - Softmax activation



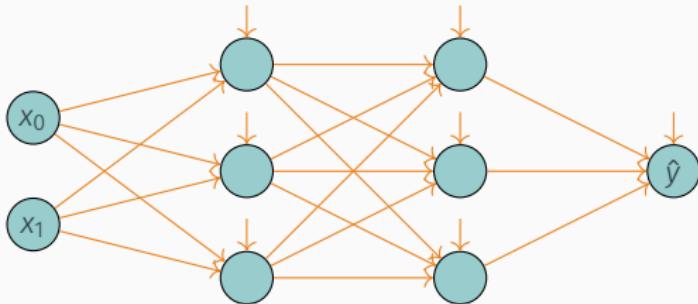
# Deep learning: Training



$$\begin{aligned}\hat{y} = & \max(0, w_{0,0}^2 * \max(0, w_{0,0}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\& w_{1,0}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\& w_{2,0}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\& b_{1,0}) + \\& w_{1,0}^2 * \max(0, w_{0,1}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\& w_{1,1}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\& w_{2,1}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\& b_{1,1}) + \\& w_{2,0}^2 * \max(0, w_{0,2}^1 * \max(0, w_{0,0}^0 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\& w_{1,2}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\& w_{2,2}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\& b_{1,2}) + \\& b_2)\end{aligned}$$



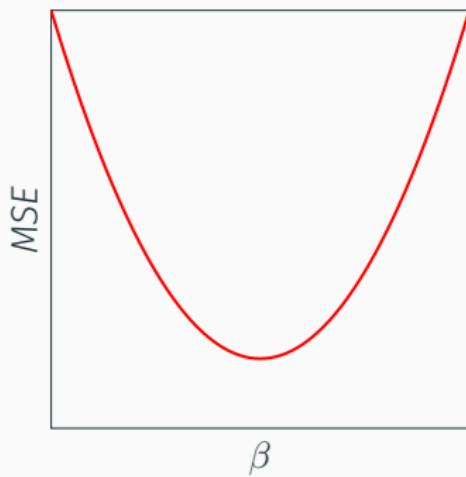
# Deep learning: Training



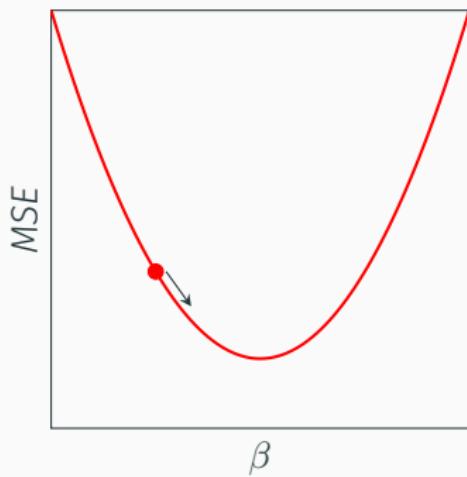
$$\begin{aligned}\hat{y} = & \max(0, w_{0,0}^2 * \max(0, w_{0,0}^1 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\& w_{1,0}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\& w_{2,0}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\& b_{1,0}) + \\& w_{1,0}^2 * \max(0, w_{0,1}^1 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\& w_{1,1}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\& w_{2,1}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\& b_{1,1}) + \\& w_{2,0}^2 * \max(0, w_{0,2}^1 * x_0 + w_{1,0}^0 * x_1 + b_{0,0}) + \\& w_{1,2}^1 * \max(0, w_{0,1}^0 * x_0 + w_{1,1}^0 * x_1 + b_{0,1}) + \\& w_{2,2}^1 * \max(0, w_{0,2}^0 * x_0 + w_{1,2}^0 * x_1 + b_{0,2}) + \\& b_{1,2}) + \\& b_2)\end{aligned}$$



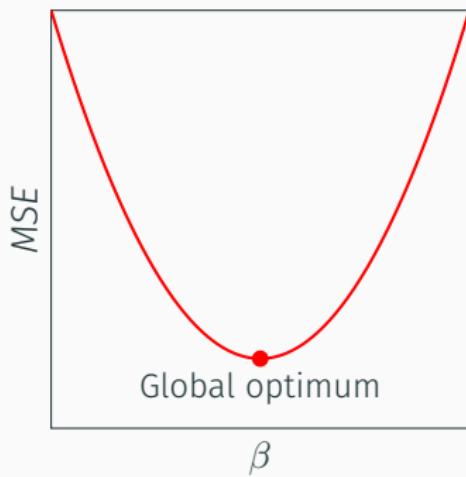
# Deep learning: Training



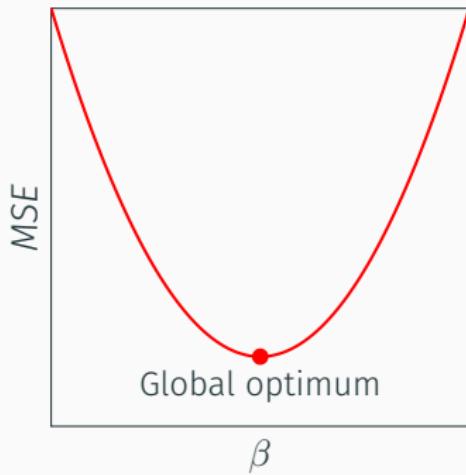
# Deep learning: Training



# Deep learning: Training



# Deep learning: Training



- $|\beta| \rightarrow 10^6 - 10^{12}$
- $\beta_x \implies \beta_y$



## Learning representations by back-propagating errors

David E. Rumelhart\*, Geoffrey E. Hinton†  
& Ronald J. Williams\*

\* Institute for Cognitive Science, C-015, University of California,  
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,  
Pittsburgh, Philadelphia 15213, USA

---

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure<sup>1</sup>.



## Learning representations by back-propagating errors

David E. Rumelhart\*, Geoffrey E. Hinton†  
& Ronald J. Williams\*

\* Institute for Cognitive Science, C-015, University of California,  
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,  
Pittsburgh, Philadelphia 15213, USA

---

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure<sup>1</sup>.

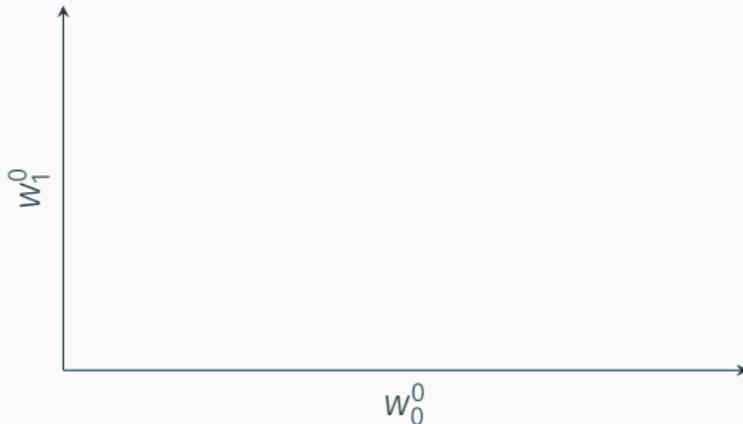


# Deep learning: Training

$$\Delta R(\theta^m) = \frac{\partial R(\theta)}{\partial \theta} \Big|_{\theta=\theta^m}$$



# Deep learning: Training

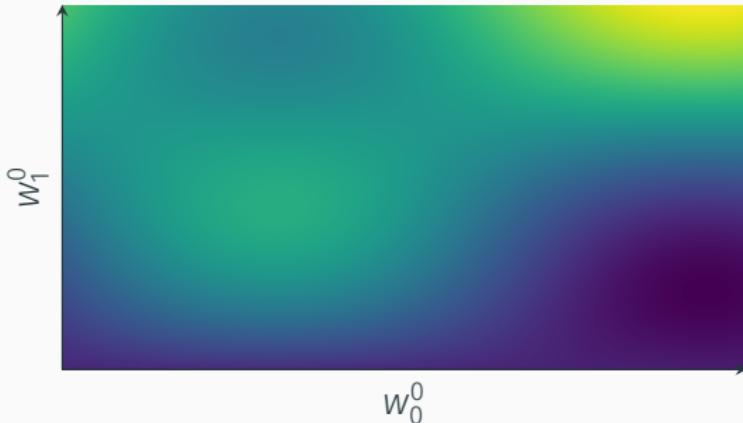


$$\Delta R(\theta^m) = \frac{\partial R(\theta)}{\partial \theta} \Big|_{\theta=\theta^m}$$

- $\theta$  are the parameters of the model



# Deep learning: Training

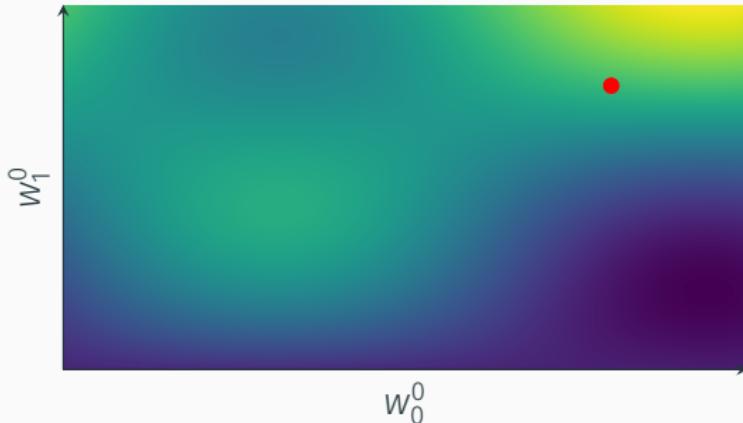


$$\Delta R(\theta^m) = \frac{\partial R(\theta)}{\partial \theta} \Big|_{\theta=\theta^m}$$

- $\theta$  are the parameters of the model
- $R(\theta)$  is the loss as a function of the parameters



# Deep learning: Training

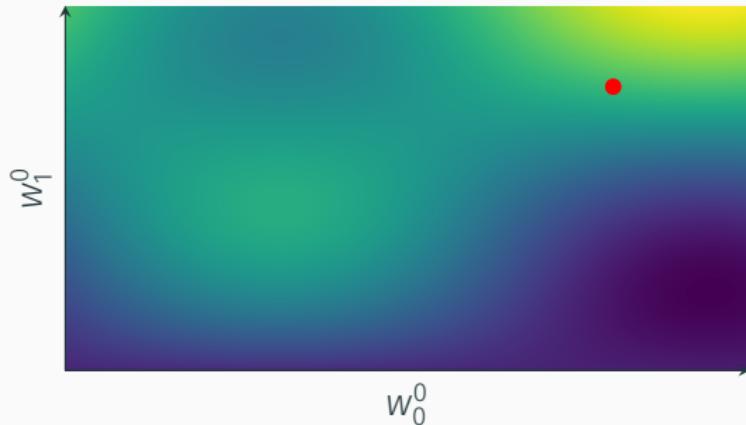


$$\Delta R(\theta^m) = \left. \frac{\partial R(\theta)}{\partial \theta} \right|_{\theta=\theta^m}$$

- $\theta$  are the parameters of the model
- $R(\theta)$  is the loss as a function of the parameters
- $\theta^m$  is a specific configuration of parameters



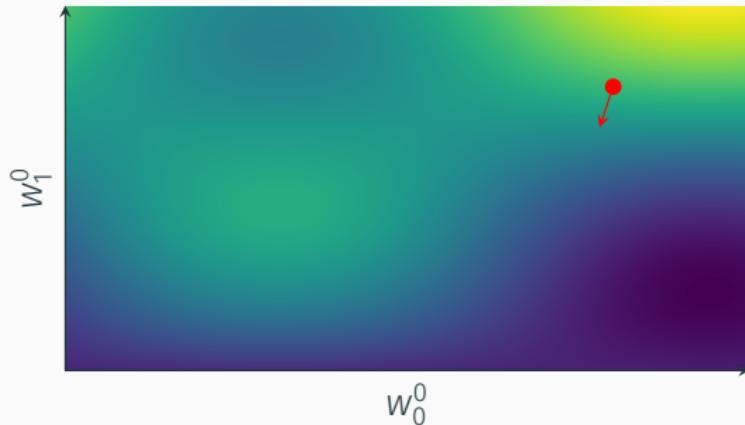
# Deep learning: Training



$$\frac{\partial R_i(\theta)}{\partial w_{jk}} = \frac{\partial R(\theta)}{\partial f_\theta(x_i)} \cdot \frac{\partial f_\theta(x_i)}{\partial g(z_{ik})} \cdot \frac{\partial g(z_{ik})}{\partial z_{ik}} \cdot \frac{z_{ik}}{w_{kj}}$$



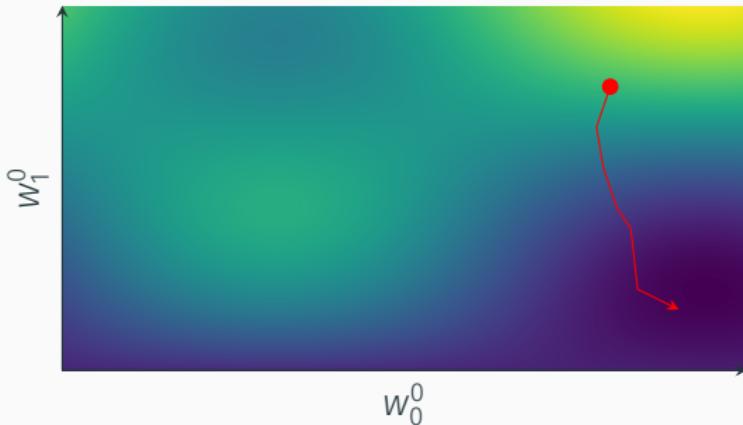
# Deep learning: Training



$$\frac{\partial R_i(\theta)}{\partial w_{jk}} = \frac{\partial R(\theta)}{\partial f_\theta(x_i)} \cdot \frac{\partial f_\theta(x_i)}{\partial g(z_{ik})} \cdot \frac{\partial g(z_{ik})}{\partial z_{ik}} \cdot \frac{z_{ik}}{w_{kj}}$$



# Deep learning: Training



$$\frac{\partial R_i(\theta)}{\partial w_{jk}} = \frac{\partial R(\theta)}{\partial f_\theta(x_i)} \cdot \frac{\partial f_\theta(x_i)}{\partial g(z_{ik})} \cdot \frac{\partial g(z_{ik})}{\partial z_{ik}} \cdot \frac{z_{ik}}{w_{kj}}$$



# Deep learning: Training

Backpropagation: Uses gradient descent to iteratively determine how the model weights should be updated to minimize the loss function

- **Gradient descent**: Calculate gradient based on all data points



# Deep learning: Training

Backpropagation: Uses gradient descent to iteratively determine how the model weights should be updated to minimize the loss function

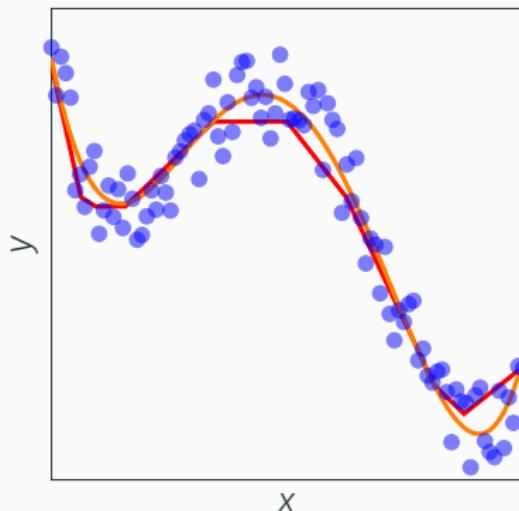
- **Gradient descent**: Calculate gradient based on all data points
- **Stochastic gradient descent**: Calculate gradient based on a batch of data points



# Deep learning: The conundrum

Splines: A smooth curve implemented via piecewise polynomial functions

Neural networks: A piecewise linear function implemented as a hierarchy of artificial neurons



# Deep learning: The conundrum

Splines: A smooth curve implemented via piecewise polynomial functions

- Requires us to carefully balance the complexity of the function

Neural networks: A piecewise linear function implemented as a hierarchy of artificial neurons



# Deep learning: The conundrum

Splines: A smooth curve implemented via piecewise polynomial functions

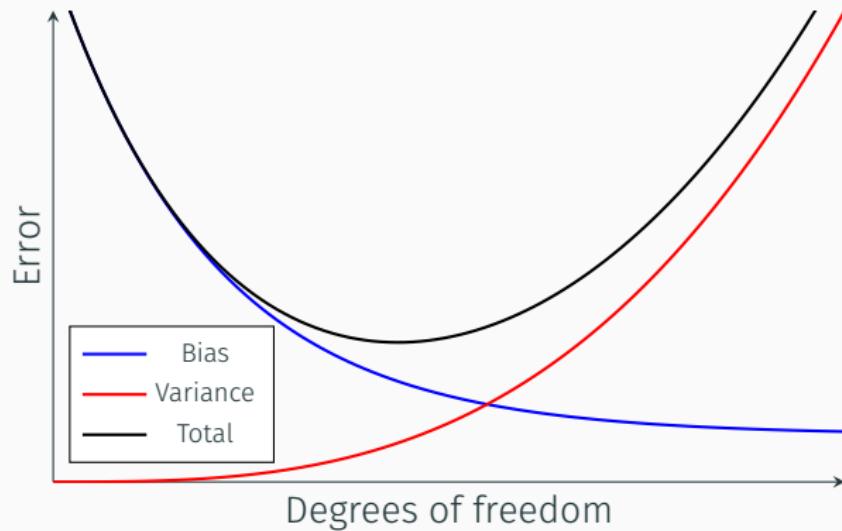
- Requires us to carefully balance the complexity of the function

Neural networks: A piecewise linear function implemented as a hierarchy of artificial neurons

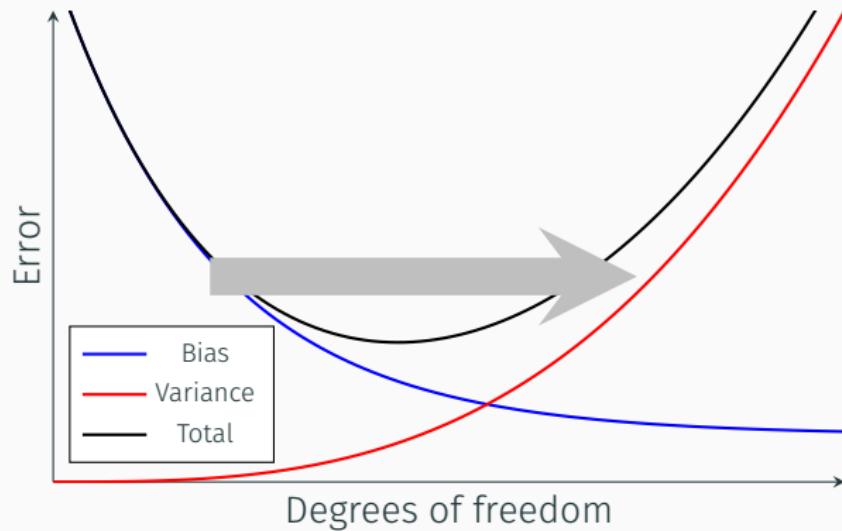
- Overparameterization 😵



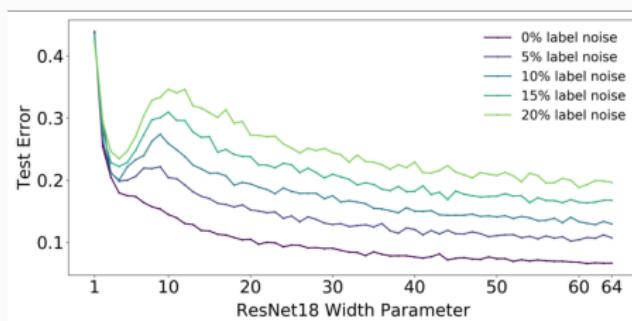
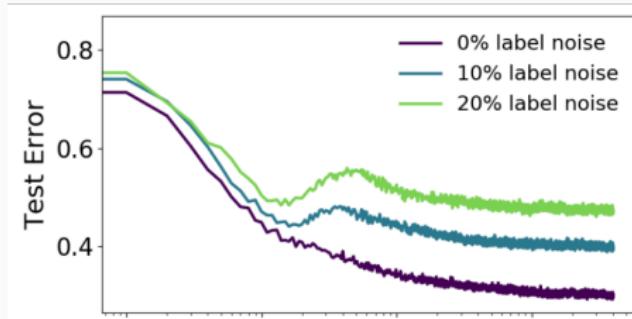
# Deep learning: The conundrum



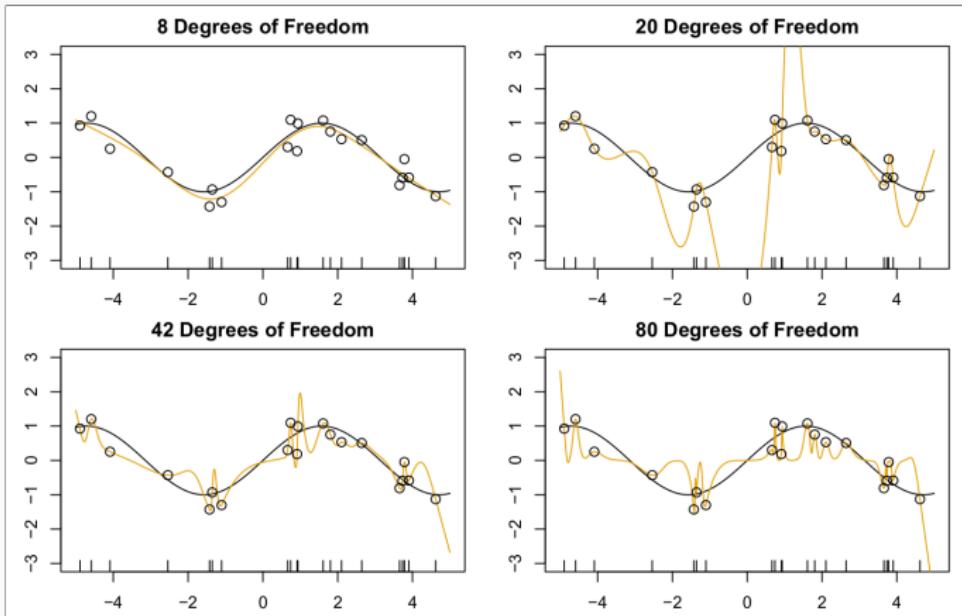
# Deep learning: The conundrum



# Deep learning: The conundrum



# Deep learning: The conundrum



# Deep learning: The conundrum

Overparameterization: Deep artificial neural networks generally have far more parameters than necessary (and often more than the number of data points)

- At face value, it is surprising that this does not yield severe overfitting
- However, it can be shown that neural networks, after perfectly fitting their training data, generally become more well-behaved and less wild



# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression

$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$



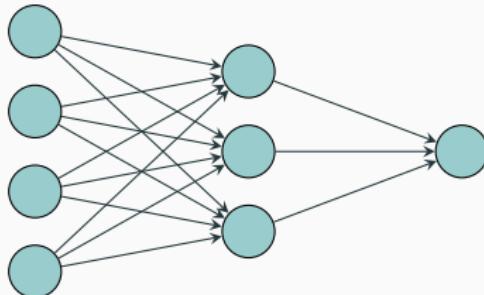
# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



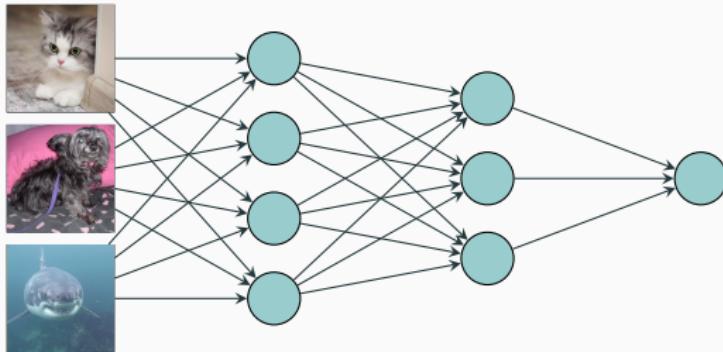
# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



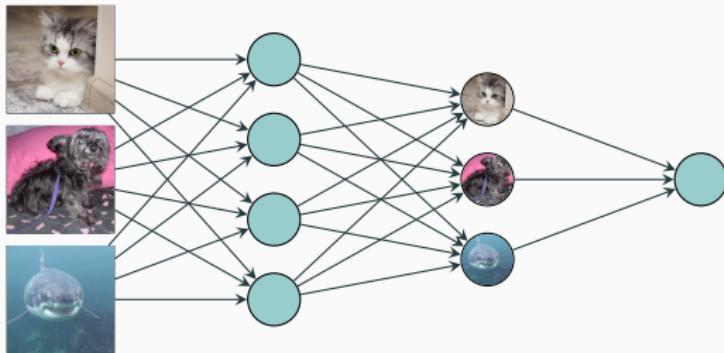
# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



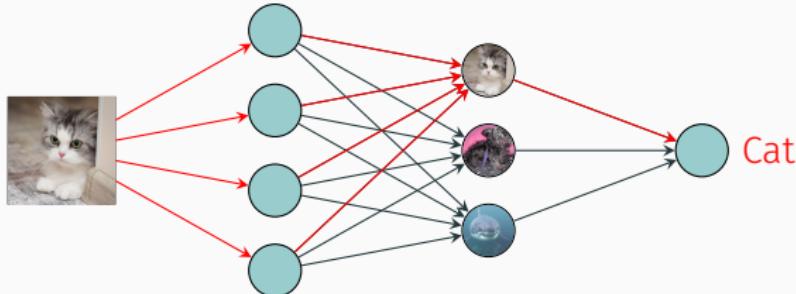
# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



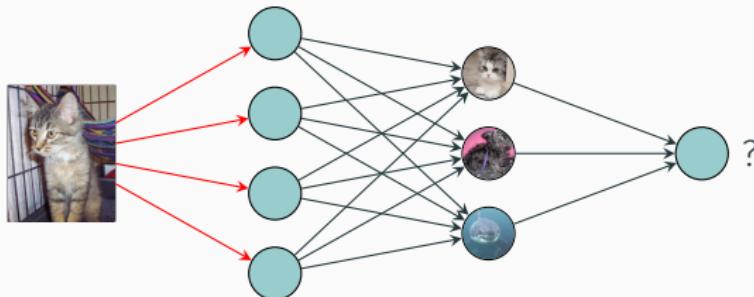
# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



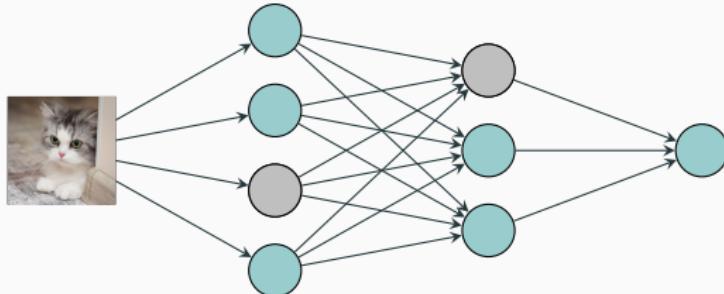
# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



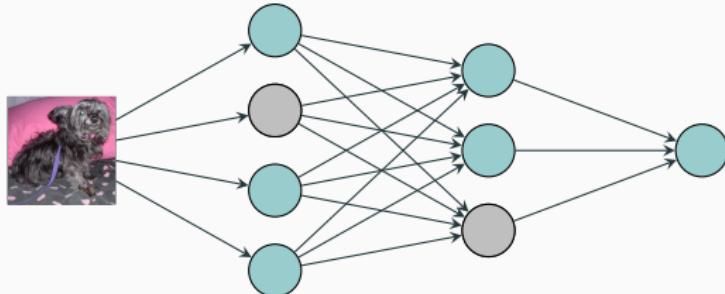
# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training



# Deep learning: Regularization

- Weight decay: Applies an  $\ell_2$ -penalty to the weights, similarly to ridge regression  
$$R(\theta; \lambda) = R(\theta) + \lambda \sum \theta^2$$
- Dropout: Randomly kills a fraction of the neurons during training
- Data augmentation: Attempts to generate more data



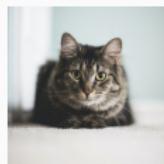
# Computer vision

---



UNIVERSITY  
OF OSLO

# Computer vision: Background



Cat



# Computer vision: Background



Sunflower



Ladybug



Cat



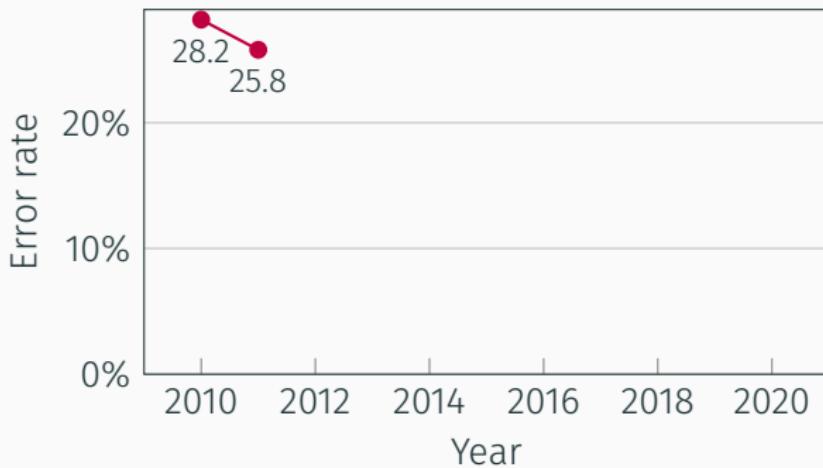
Airplane



Shark

ImageNet: ~14m images, ~22k categories

# Computer vision: Background



# Computer vision: Background



# Computer vision: Background



# Computer vision: Background



# Computer vision: Image data



# Computer vision: Image data



# Computer vision: Image data



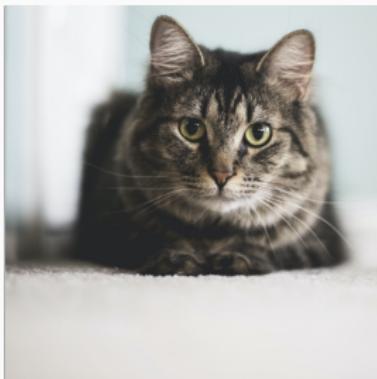
# Computer vision: Image data



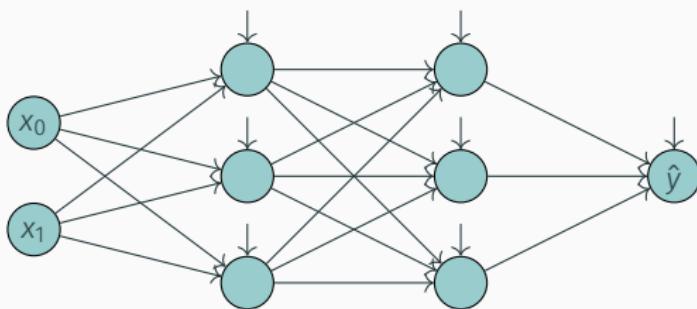
# Computer vision: Image data



# Computer vision: Image data



# Convolutional neural networks: Inputs



# Convolutional neural networks: Inputs

$x_0$

$x_1$

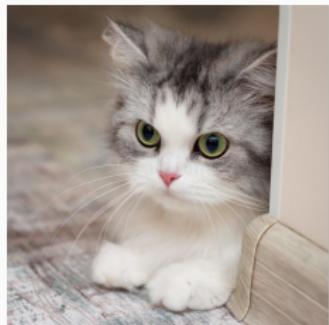


# Convolutional neural networks: Inputs

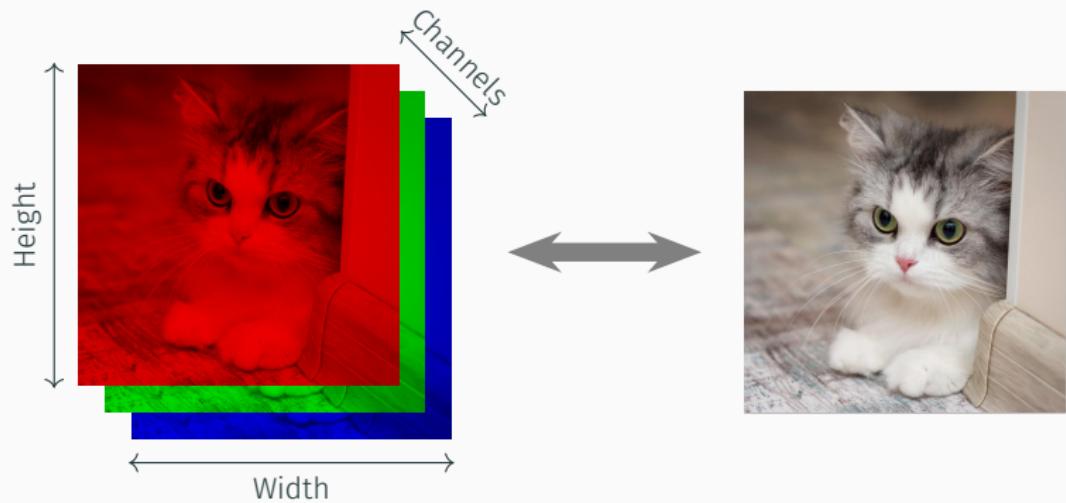
$x_0$

$x_1$

?



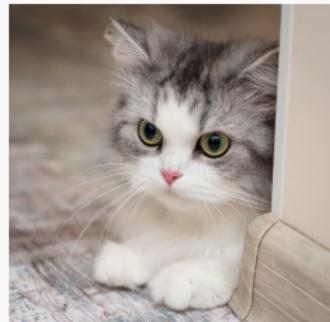
# Convolutional neural networks: Inputs



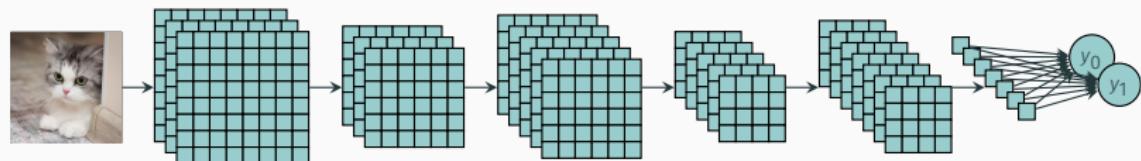
# Convolutional neural networks: Inputs

$x_{000}$	$x_{010}$	$x_{020}$	$x_{030}$	$x_{040}$		
$x_{100}$	$x_{110}$	$x_{120}$	$x_{130}$	$x_{140}$		
$x_{200}$	$x_{210}$	$x_{220}$	$x_{230}$	$x_{240}$		
$x_{300}$	$x_{310}$	$x_{320}$	$x_{330}$	$x_{340}$		
$x_{400}$	$x_{410}$	$x_{420}$	$x_{430}$	$x_{440}$		

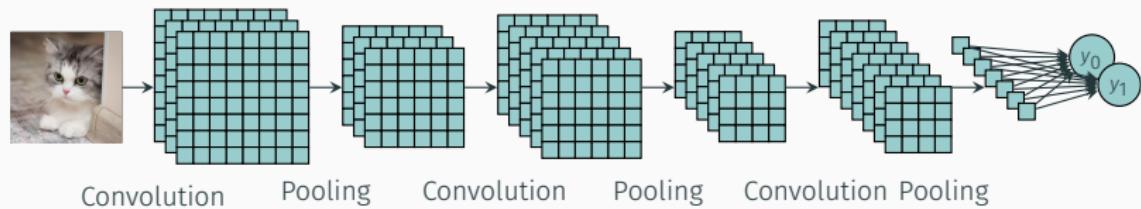
The diagram illustrates a convolution operation on a 5x5 input grid. The input grid contains 25 elements labeled  $x_{ij0}$  where  $i, j \in \{0, 1, 2, 3, 4\}$ . A 2x2 stride convolution is applied, indicated by arrows pointing from the input grid to a 3x3 output grid. The output grid has 9 elements labeled  $x_{ij1}$  where  $i, j \in \{0, 1, 2\}$ . The arrows show that only every second element from the input grid is selected for the output grid, creating a downsampling effect.



# Convolutional neural networks: Architecture

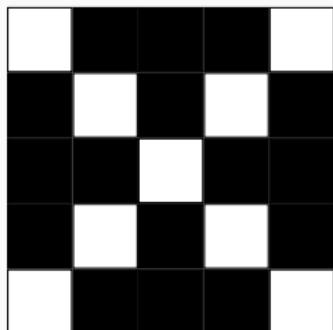


# Convolutional neural networks: Architecture



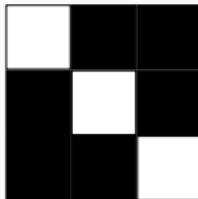
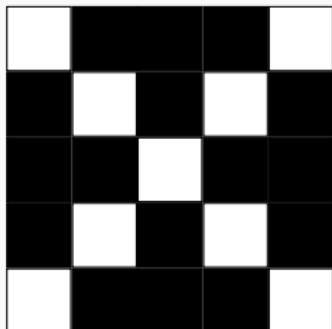
# Convolutional neural networks: Convolution

Image



# Convolutional neural networks: Convolution

Image

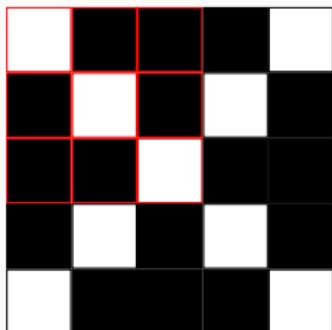


Pattern 1

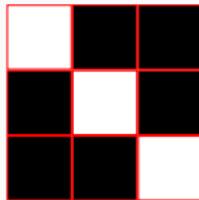


# Convolutional neural networks: Convolution

Image



3

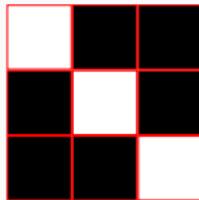
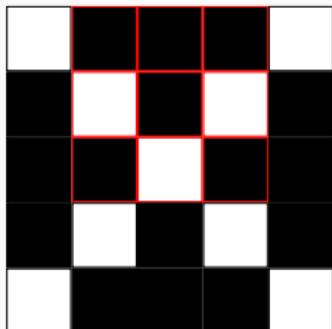


Pattern 1



# Convolutional neural networks: Convolution

Image

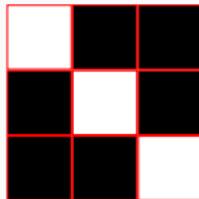
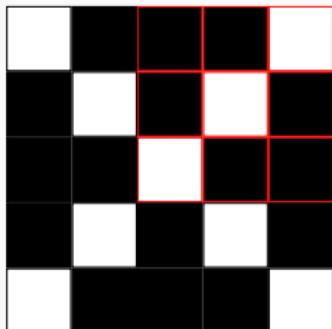


Pattern 1



# Convolutional neural networks: Convolution

Image

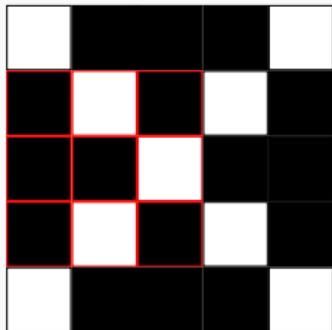


Pattern 1

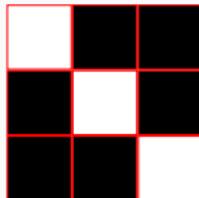


# Convolutional neural networks: Convolution

Image



3	0	1
0		

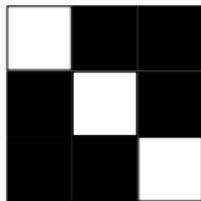
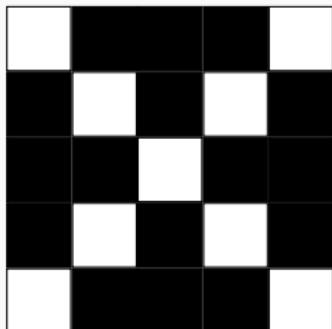


Pattern 1



# Convolutional neural networks: Convolution

Image



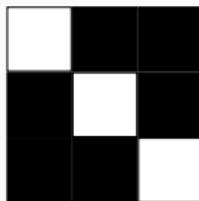
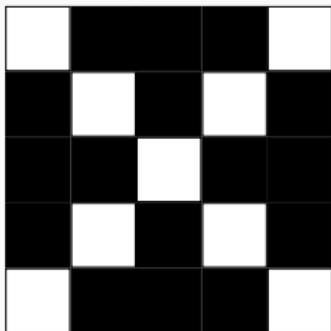
Pattern 1

3	0	1
0	3	0
1	0	3

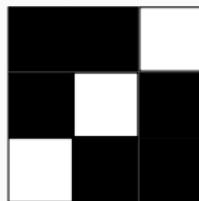


# Convolutional neural networks: Convolution

Image



Pattern 1



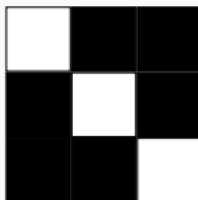
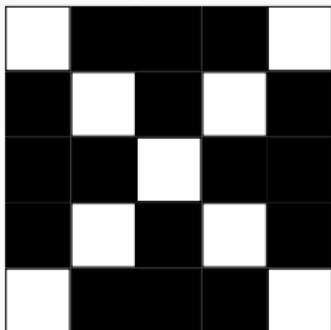
Pattern 2

1	0	3
0	3	0
3	0	1

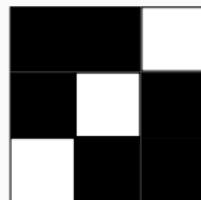


# Convolutional neural networks: Convolution

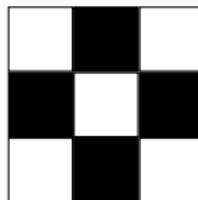
Image



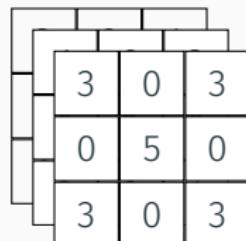
Pattern 1



Pattern 2

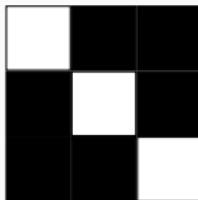
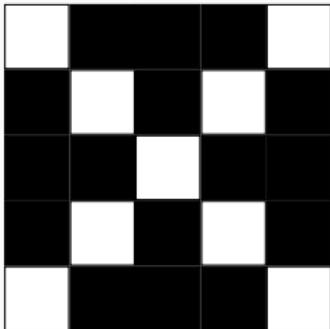


Pattern 3

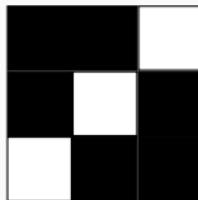


# Convolutional neural networks: Convolution

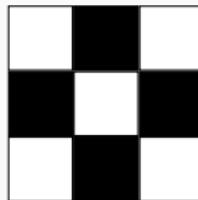
Image



Pattern 1

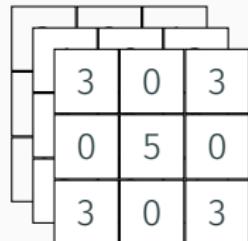


Pattern 2



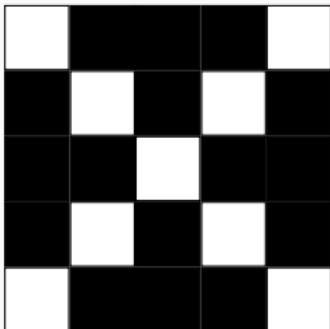
Pattern 3

Feature maps

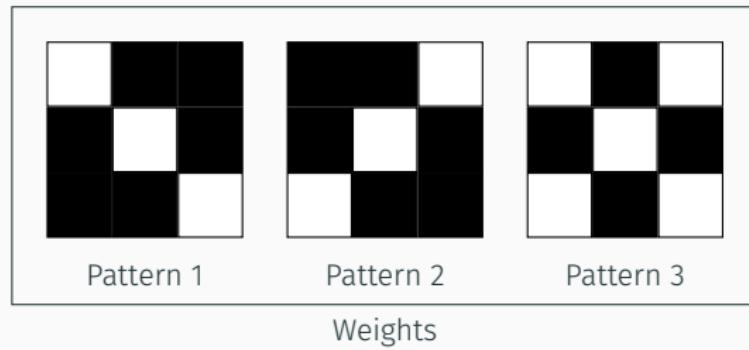
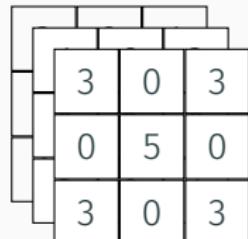


# Convolutional neural networks: Convolution

Image

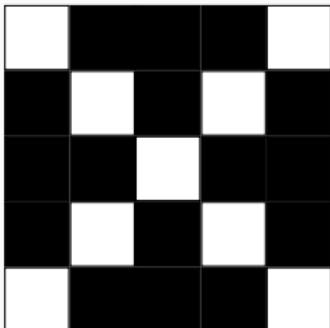


Feature maps

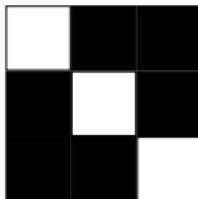
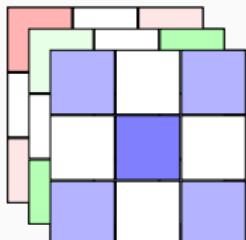


# Convolutional neural networks: Convolution

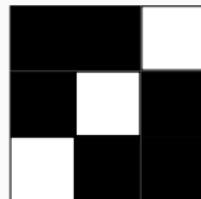
Image



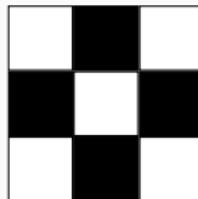
Feature maps



Pattern 1



Pattern 2



Pattern 3

Weights



# Convolutional neural networks: (Max-)Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	1
12	13	14	15



# Convolutional neural networks: (Max-)Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	1
12	13	14	15



# Convolutional neural networks: (Max-)Pooling

Feature map

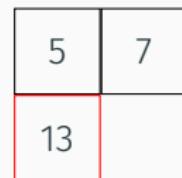
0	1	2	3
4	5	6	7
8	9	10	1
12	13	14	15



# Convolutional neural networks: (Max-)Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	1
12	13	14	15



# Convolutional neural networks: (Max-)Pooling

Feature map

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

5	7
13	15



# Convolutional neural networks: (Max-)Pooling

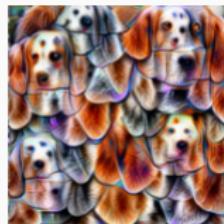
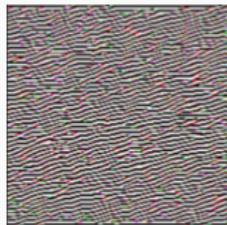
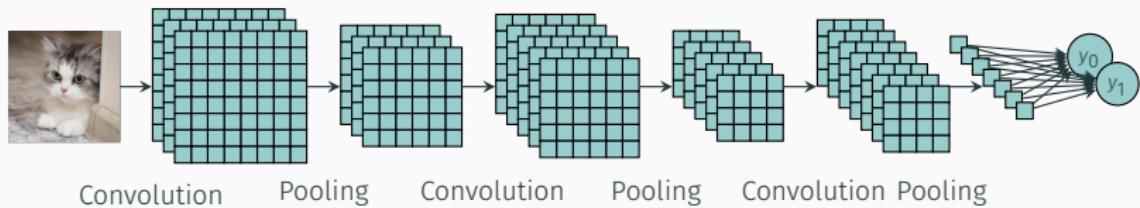
Feature map

0	1	2	3
4	5	6	7
8	9	10	1
12	13	14	15

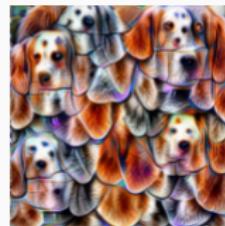
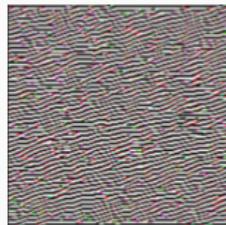
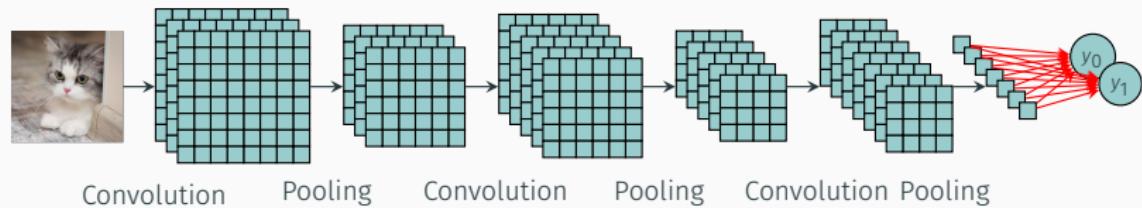
5	7
13	15



# Convolutional neural networks: Overview



# Convolutional neural networks: Overview

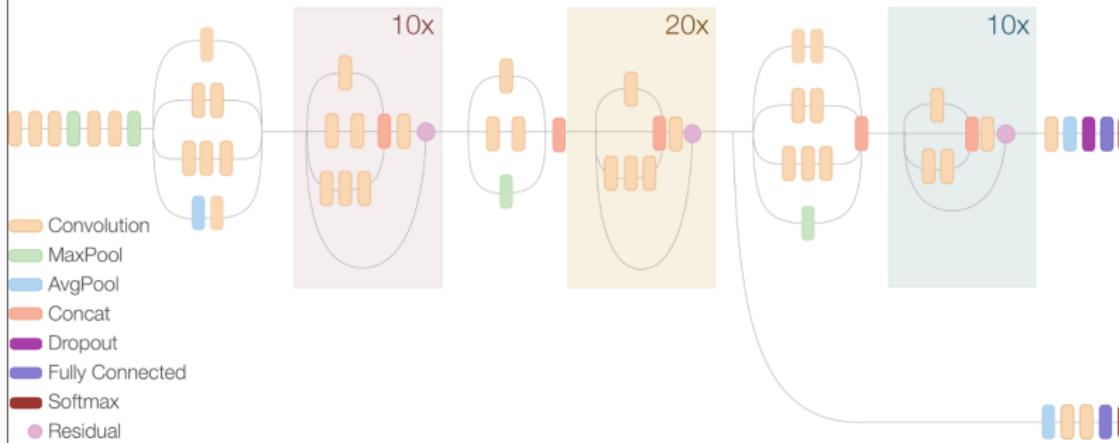


# Convolutional neural networks: Overview

## Inception Resnet V2 Network



## Compressed View



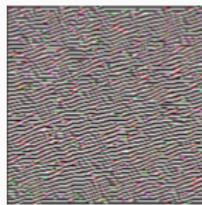
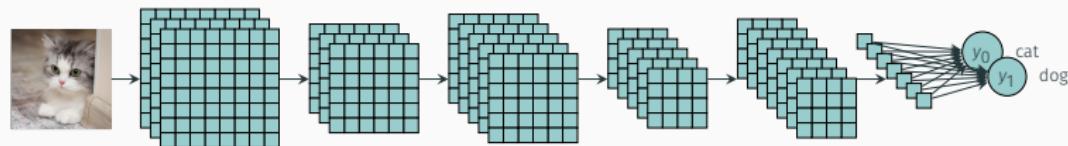
# Convolutional neural networks: Overview

Convolutional neural networks (CNNs): Artificial neural networks tailored specifically for image data.

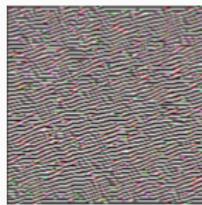
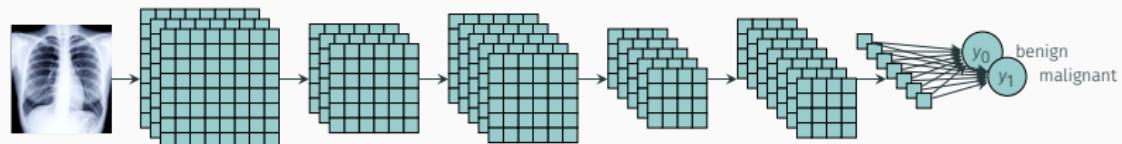
- Takes raw pixel data as input, e.g. arrays of size  $H \times W \times C$
- Mainly consists of convolutions and pooling operations, which lets us recognize larger and more abstract patterns the deeper we get in the model
- The trainable parameters are the weights of the convolutional kernels, e.g. the patterns the model is looking for
- New architectures extend beyond this basic formula, by employing residual connections, attention mechanisms etc.



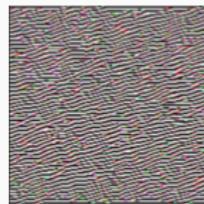
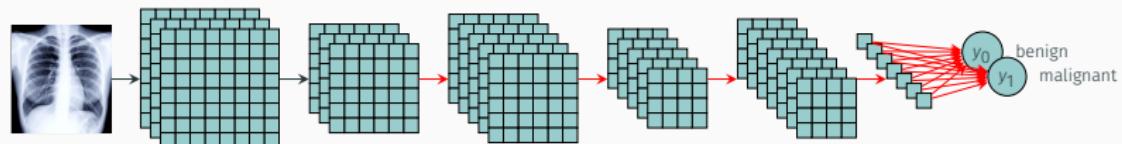
# Computer vision: Transfer learning



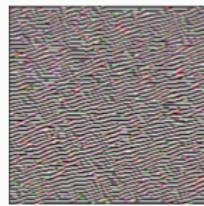
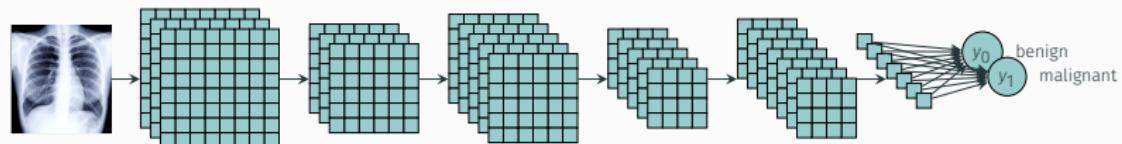
# Computer vision: Transfer learning



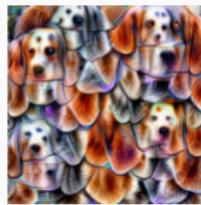
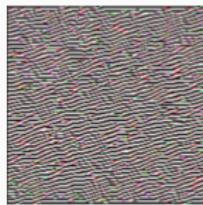
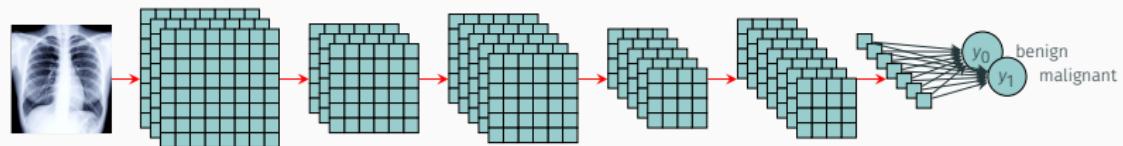
# Computer vision: Transfer learning



# Computer vision: Transfer learning



# Computer vision: Transfer learning



# Computer vision: Transfer learning

Transfer learning: Utilizing pretrained models to solve new tasks.

- Allows us to solve tasks where we don't have enough data to train a model from scratch
- Common to either freeze the weights of the convolutional part of pre-trained model and train a new classifier on top of them, or to finetune the entire model

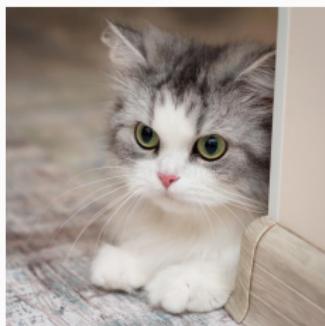


# Computer vision: Transfer learning

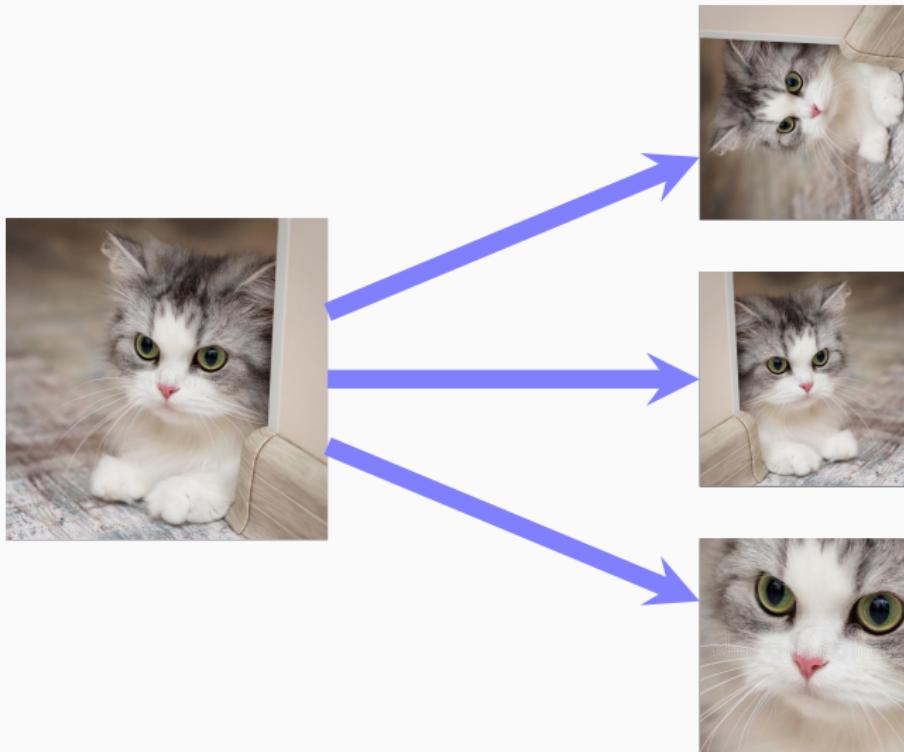
<https://keras.io/api/applications/>



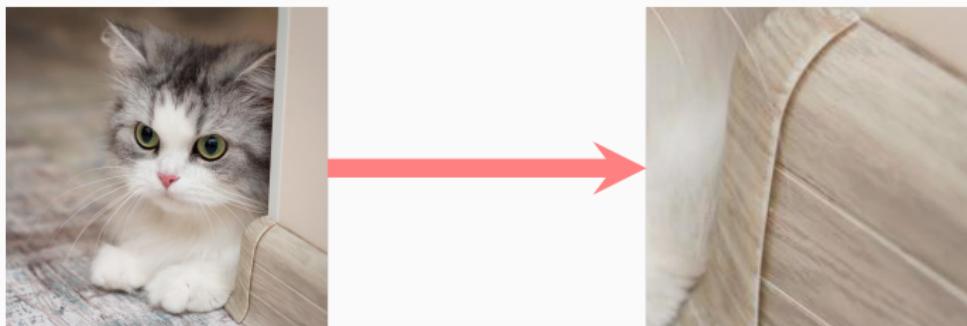
# Computer vision: Data augmentation



# Computer vision: Data augmentation



# Computer vision: Data augmentation



# Computer vision: Data augmentation

<https://albumentations.ai/>



# Computer vision: Tutorial

<https://github.com/estenhl/flowers>

