

PSY9511: Seminar 3

Regularization and variable selection

Esten H. Leonardsen

11.09.24



UNIVERSITETET
I OSLO

Outline

1. Recap
2. Assignment 2
3. Regularization
 - Variable selection
 - Shrinkage



Assignment 2



UNIVERSITETET
I OSLO

Assignment 2

1. Download the Auto.csv dataset from the ISLP website.
2. Read the Auto.csv-dataset into memory.
3. In the horsepower-column, some values are missing. These are encoded with '?'. Remove these rows from the dataset.
4. Create a new column 'muscle'. This column should contain a 1 for all muscle cars (e.g. cars that have above average horsepower) and 0 for the rest.
5. Split the dataset into a training set and a test set, by randomly drawing 80% of the rows for the former and 20% of the rows for the latter.

<http://localhost:8890/notebooks/notebooks/Assignment%202.ipynb>



Assignment 2

1. Fit a simple linear regression model using horsepower as the predictor and mpg as the outcome using the training data.
2. Create a scatter plot with horsepower on the x-axis and mpg on the y-axis using the testing data. Plot the regression line found by the model in the plot.
3. Use the model to generate predictions for the training set. Calculate and report the mean absolute error (MAE) of these predictions.
4. Use the model to generate predictions for the test set. Calculate and report the MAE of the predictions.

<http://localhost:8890/notebooks/notebooks/Assignment%202.ipynb>



Assignment 2

Reflection: Is the training or testing MAE is lower? Does this match your expectation? What would be the general pattern we expect here (e.g. one is lower than the other, they are the same, etc.), and why do we expect that?



Assignment 2

1. Fit a multivariate linear regression model using horsepower, weight, displacement, and year as predictors and mpg as the outcome.
2. Print the intercept and coefficients of the model .
3. Use the model to generate predictions for the training set. Calculate and report the MAE of these predictions.
4. Use the model to generate predictions for the test set. Calculate and report the MAE of the predictions.

<http://localhost:8890/notebooks/notebooks/Assignment%202.ipynb>



Assignment 2

Is the training MAE lower or higher than in the simple linear regression model? Does it have to be this way, or could it have been otherwise? What about the testing MAE?



Assignment 2

1. Fit a logistic regression model using weight, displacement and year as predictors and our newly created muscle-column as the outcome.
2. Use the model to generate predictions for the training set. Calculate and report the accuracy of these predictions.
3. Use the model to generate predictions for the testing set. Calculate and report the accuracy of these predictions.


<http://localhost:8890/notebooks/notebooks/Assignment%202.ipynb>



Assignment 2: Random seeds

```
In[1]: import pandas as pd


df = pd.DataFrame(...)
train = df.sample(frac=0.8)
test = df.drop(train.index)
```



Assignment 2: Random seeds

```
In[1]: import pandas as pd


df = pd.DataFrame(...)
train = df.sample(frac=0.8)
test = df.drop(train.index)
```



Assignment 2: Random seeds


```
In[1]: import pandas as pd

df = pd.DataFrame(...)
train = df.sample(frac=0.8)
test = df.drop(train.index)
```



Assignment 2: Random seeds

```
In[1]: import pandas as pd
import numpy as np
np.random.seed(42)
df = pd.DataFrame(...)
train = df.sample(frac=0.8)
test = df.drop(train.index)
```



Assignment 2: Log-odds vs probability vs class

```
model <- glm(muscle ~ year + weight, data=df, family='binomial')  
predict(model, df)
```



Assignment 2: Log-odds vs probability vs class

```
model <- glm(muscle ~ year + weight, data=df, family='binomial')  
predict(model, df)
```

1	2	3	4	5	6	7	8	9	10	11
1.2460	1.9245	1.0019	0.9911	1.0485	4.2506	4.1465	4.5522	2.4889	1.4578	1.6223



Assignment 2: Log-odds vs probability vs class

```
model <- glm(muscle ~ year + weight, data=df, family='binomial')  
predict(model, df)
```

1	2	3	4	5	6	7	8	9	10	11
1.2460	1.9245	1.0019	0.9911	1.0485	4.2506	4.1465	4.5522	2.4889	1.4578	1.6223

```
predict(model, df, type="response")
```

1	2	3	4	5	6	7	8	9	10	11
0.7766	0.8726	0.7314	0.7293	0.7405	0.9853	0.9844	0.9895	0.9233	0.8112	0.8352



Assignment 2: Log-odds vs probability vs class

```
In[1]: model = LogisticRegression()  
model.fit(df[['year', 'weight']], df['muscle'])  
model.predict(df[['year', 'weight']])
```

```
Out[1]: array([0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0])
```

```
In[1]: model.predict_proba(df[['year', 'weight']])
```

```
Out[1]: array([[0.14, 0.86], [0.08, 0.92], [0.17, 0.83], [0.18, 0.82]])
```



Assignment 2: Eye test

```
model <- glm(muscle ~ year + weight, data=df, family='binomial')  
predict(model, df)
```

1	2	3	4	5	6	7	8	9	10	11
1.2460	1.9245	1.0019	0.9911	1.0485	4.2506	4.1465	4.5522	2.4889	1.4578	1.6223



Assignment 2: Eye test

```
model <- glm(muscle ~ year + weight, data=df, family='binomial')
predict(model, df)
```

1	2	3	4	5	6	7	8	9	10	11
1.2460	1.9245	1.0019	0.9911	1.0485	4.2506	4.1465	4.5522	2.4889	1.4578	1.6223

```
model <- lm(mpg ~ horsepower, df)
summary(model)
```

```
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 19.59412  0.96187 20.371 < 2e-16 ***
horsepower102 0.40588  4.08087  0.099  0.920840
horsepower103 0.70588  4.08087  0.173  0.862789
horsepower105 0.90588  1.49529  0.606  0.545091
```



Regularization



UNIVERSITETET
I OSLO


Regularization: Motivation

$$y \sim \beta_0 + \beta_1 * x_1 + \beta_2 * x_2$$




Regularization: Motivation

$$y \sim \beta_0 + \beta_1 * x_1 + \beta_2 * x_2$$



Regularization: Motivation

$$y \sim \beta_0 + \beta_1 * x_1 + \beta_2 * x_2 + \beta_3 * x_3 + \beta_4 * x_4$$



Regularization: Motivation

```
In[1]: import pandas as pd

df = pd.read_csv('/Users/esten/Downloads/Auto.csv')
train = df.iloc[:int(len(df) * 0.8)]
validation = df.iloc[int(len(df) * 0.8):]

print(f'Using len(train) samples for training')
print(f'Using len(validation) samples for validation')
```

```
Out[1]: Using 317 samples for training
Using 80 samples for validation
```



Regularization: Motivation

1. Variable selection
 - a. Best subset selection
 - b. Forward stepwise selection
 - c. Backward stepwise selection
2. Shrinkage
 - a. LASSO
 - b. Ridge Regression
3. Dimensionality reduction: Lecture 6 and self-study



Regularization: Motivation

1. Variable selection
 - a. Best subset selection
 - b. Forward stepwise selection
 - c. Backward stepwise selection
2. Shrinkage
 - a. LASSO
 - b. Ridge Regression
3. Dimensionality reduction: Lecture 6 and self-study



Variable selection



UNIVERSITETET
I OSLO

The number of predictors we are using in our model directly impacts model complexity.



Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Motivation

To reduce model complexity (and therefore risk of overfitting), avoid problems due to collinearity, and simplify subsequent interpretations.

Problem


We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Motivation


To reduce model complexity (and therefore risk of overfitting), **avoid problems due to collinearity**, and simplify subsequent interpretations.

Variable selection


cylinders




displacement




horsepower




weight



acceleration



year



Variable selection: Best subset selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Train models on all subsets p and select the best one.



Variable selection: Best subset selection

```
In[1]: import numpy as np

from itertools import chain, combinations
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

subsets = list(chain.from_iterable(combinations(predictors, r)
                                   for r in range(len(predictors)+1)))

best = {'mse': float('inf'), 'subset': None}

for subset in subsets:
    if len(subset) == 0:
        continue

    model = LinearRegression()
    model.fit(train[list(subset)], train[target])
    predictions = model.predict(validation[list(subset)])
    mse = mean_squared_error(validation[target], predictions)

    best = best if mse > best['mse'] else {'mse': mse, 'subset': subset}

print(f'MSE: {best["mse"]:.2f}, predictors: {best["subset"]}')
```

```
Out[1]: MSE: 29.68, predictors: ('cylinders', 'displacement', 'horsepower', 'weight', 'year')
```



Variable selection: Best subset selection

```
In[1]: import numpy as np

from itertools import chain, combinations
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

subsets = list(chain.from_iterable(combinations(predictors, r)
                                   for r in range(len(predictors)+1)))

best = {'mse': float('inf'), 'subset': None}

for subset in subsets:
    if len(subset) == 0:
        continue

    model = LinearRegression()
    model.fit(train[list(subset)], train[target])
    predictions = model.predict(validation[list(subset)])
    mse = mean_squared_error(validation[target], predictions)

    best = best if mse > best['mse'] else {'mse': mse, 'subset': subset}

print(f'MSE: {best["mse"]:.2f}, predictors: {best["subset"]}')
```

```
Out[1]: MSE: 29.68, predictors: ('cylinders', 'displacement', 'horsepower', 'weight', 'year')
```



Variable selection: Best subset selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Train models on all subsets p and select the best one.

+ Positives


- Guaranteed to find the optimal solution.
- Simple implementation

- Drawbacks

- Need to train many ($2^{|P|}$) models.



Variable selection: Best subset selection



Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.



Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the **best model** until all are included.




Variable selection: Forward stepwise selection


$$\begin{array}{l} y \sim \mathbf{1} \\ mse = 146.47 \end{array}$$




Variable selection: Forward stepwise selection




Variable selection: Forward stepwise selection




Variable selection: Forward stepwise selection




Variable selection: Forward stepwise selection




Variable selection: Forward stepwise selection




Variable selection: Forward stepwise selection



Variable selection: Forward stepwise selection



Variable selection: Forward stepwise selection



Variable selection: Forward stepwise selection

```
In[1]: def fit_and_evaluate(train: pd.DataFrame, validation: pd.DataFrame,
    predictors: List[str], target: str):
    model = LinearRegression()
    model.fit(train[predictors], train[target])

    train_predictions = model.predict(train[predictors])
    validation_predictions = model.predict(validation[predictors])

    return np.mean((train_predictions - train[target]) ** 2),
           np.mean((validation_predictions - validation[target]) ** 2)

predictors = ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year']
target = 'mpg'

train['intercept'] = 1
validation['intercept'] = 1
train_mse, validation_mse = fit_and_evaluate(train, validation, predictors=['intercept'], target=target)
print(f'[: {validation_mse:.2f}] ({train_mse:.2f})')

chosen_predictors = []

while len(chosen_predictors) < len(predictors):
    best_predictor = {'train_mse': None, 'validation_mse': float('inf'),
                    'predictor': None}

    for predictor in set(predictors) - set(chosen_predictors):
        train_mse, validation_mse = fit_and_evaluate(train, validation, predictors=chosen_predictors + [predictor], target=target)

        if validation_mse < best_predictor['validation_mse']:
            best_predictor = {'train_mse': train_mse, 'validation_mse': validation_mse, 'predictor': predictor}

    chosen_predictors.append(best_predictor['predictor'])

print(f'{chosen_predictors}: {best_predictor["validation_mse"]:.2f} ({best_predictor["train_mse"]:.2f})')
```



Variable selection: Forward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with no predictors. Iteratively add the predictor that yields the best model until all are included.

+ Positives

- Need to train fewer models.

- Drawbacks

- Not guaranteed to find the optimal solution.



Variable selection: Backward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with all predictors. Iteratively remove the predictor that yields the best model until all you have none left.

+ Positives

- Need to train fewer models.

- Drawbacks

- Not guaranteed to find the optimal solution.



Variable selection: Backward stepwise selection

Problem

We have a set of predictors $P = \{x_0, x_1, \dots\}$ and a target variable y , and we want to find the subset $p \subseteq P$ that yields the best (linear) model for predicting y .

Solution

Start with **all** predictors. Iteratively **remove** the predictor that yields the best model until all you have none left.

+ Positives

- Need to train fewer models.

- Drawbacks

- Not guaranteed to find the optimal solution.



Variable selection: Backward stepwise selection

Does forward and backward stepwise selection yield the same set of predictors?



Shrinkage



UNIVERSITETET
I OSLO

Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \beta_5 X_5 + \beta_6 X_6$$



Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \beta_5 X_5 + \beta_6 X_6$$



Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \beta_5 X_5 + \beta_6 X_6$$

$$\beta_n \rightarrow 0$$



Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \beta_5 X_5 + \beta_6 X_6$$

$$\beta_n \rightarrow 0$$

1. $\beta_1 = 0 \implies$ One less degree of freedom in our function



Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \beta_5 X_5 + \beta_6 X_6$$


$$\beta_n \rightarrow 0$$

1. $\beta_1 = 0 \implies$ One less degree of freedom in our function

$$mse = bias^2 + variance + irreducible\ error$$



Shrinkage: Outline



$$mse = bias^2 + variance + irreducible\ error$$



Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \beta_5 X_5 + \beta_6 X_6$$

$$\beta_n \rightarrow 0$$

1. $\beta_1 = 0 \implies$ One less degree of freedom in our function
2. A little more bias \implies A lot less variance



Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5 + \beta_6 x_6$$

$$\beta_n \rightarrow 0$$

1. $\beta_1 = 0 \implies$ One less degree of freedom in our function
2. A little more bias \implies A lot less variance

$$y \sim \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5 + \beta_6 x_6$$

↑
↓



Shrinkage: Outline

$$y \sim \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5 + \beta_6 x_6$$

$$\beta_n \rightarrow 0$$

1. $\beta_1 = 0 \implies$
One less degree of freedom in our function
2. A little more bias \implies A lot less variance
3. Parameters depend on each other \implies
Fewer degrees of freedom

